

Methods of Cloud Computing

Chapter 3: Management of Virtual Resources



Complex and Distributed Systems
Faculty IV
Technische Universität Berlin



Operating Systems and Middleware
Hasso-Plattner-Institut
Universität Potsdam

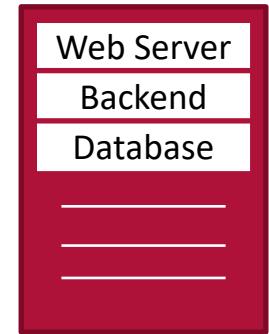
Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

“Iron Age”: Bare-Metal Servers



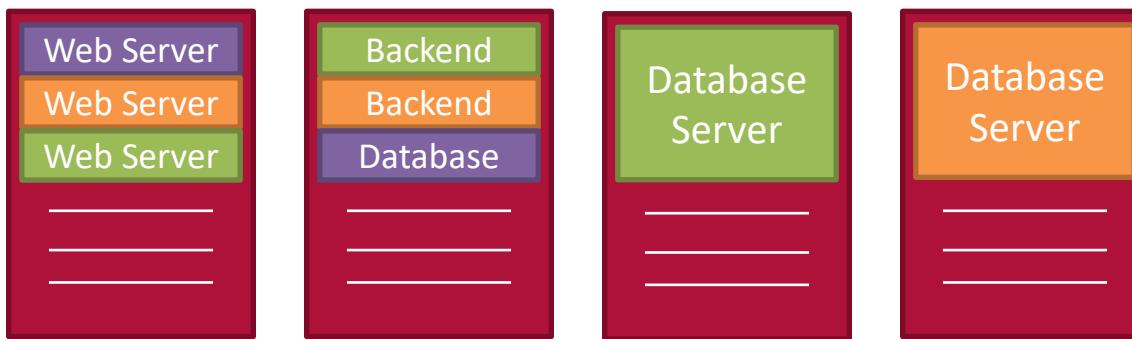
Or



- Running components of an application on one or multiple physical servers:
 - Costs and agility: Purchase, housing, maintenance
 - Fault tolerance: Servers are single points of failure
 - Resource utilization: Servers were often underutilized, but sometimes also bottlenecks (i.e. w/ dynamic load)

Bare-Metal Servers in one Organization

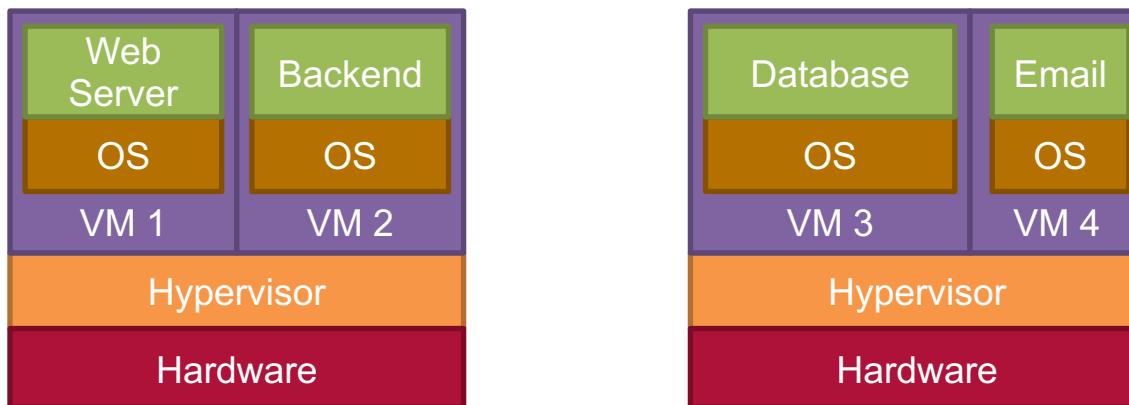
- Server consolidation: Manually assign components from different applications to set of physical nodes



- Only possible within one organization without isolation
- And static assignment, while loads may vary over time

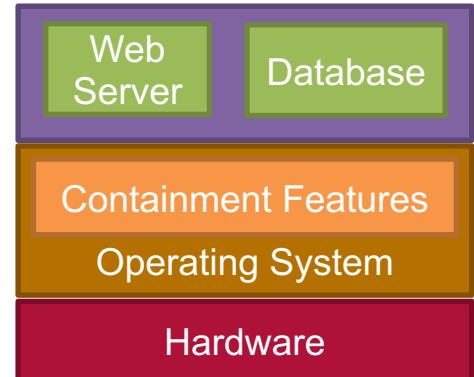
Recap: Virtual Machines

- Hardware virtualization: Same architecture, but possibly different OS and resource configuration (#cores, RAM, storage) plus isolation
- Same physical host can run multiple virtual machines
- VMs can be migrated dynamically and snapshots can provide fault tolerance



Recap: Container

- Motivation: VM images are large (contain OS) and starting VMs takes longer than processes (10x)
- Containers: Same OS kernel, but different resource configuration + isolation between contained processes
- Containment using Linux kernel features such as cgroups and namespaces

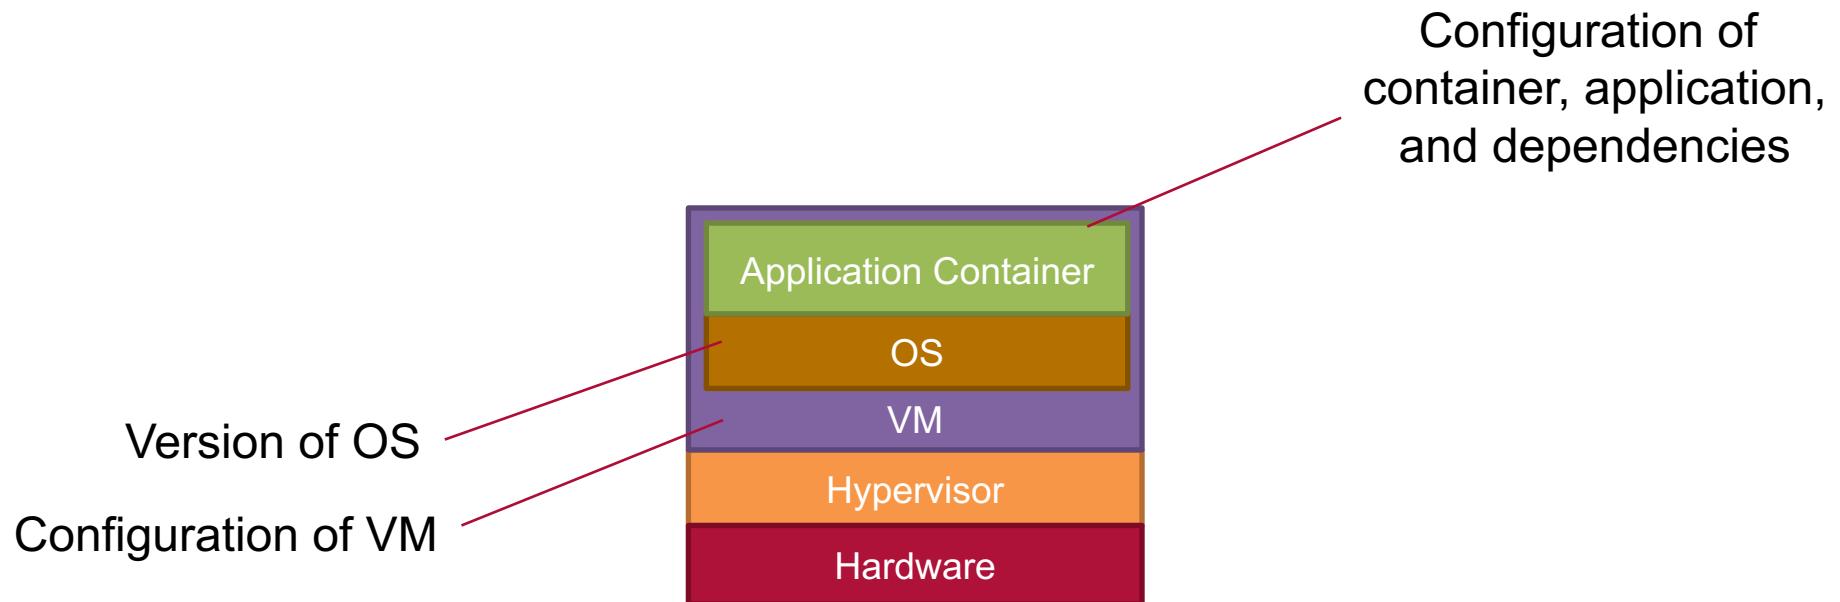


Management of Resources (1/2)

- Now we can create VMs and containers, but management of large sets of resources is still difficult:
 - Provisioning of VMs and containers
 - Configuration of systems
 - Monitoring and failure handling
 - Replication and load balancing
- Cloud Operating Systems: manage large sets of virtual resources running on large sets of physical resources

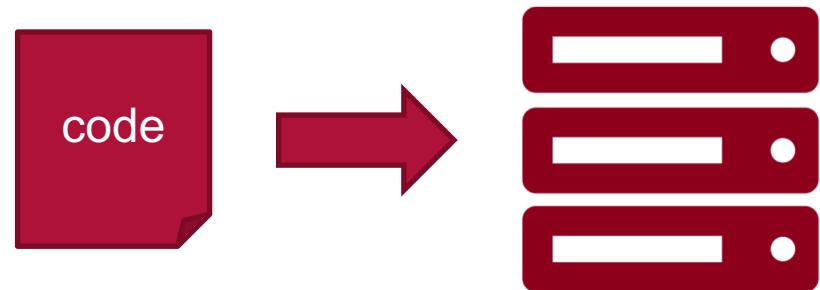
Management of Resources (2/2)

- Major remaining problems in the “Cloud Age”: server sprawl, configuration drift, snowflake servers... → Technical debt!



Infrastructure-as-Code

- Define servers, networking, and other infrastructure elements in source code:
 - Configuration of environments
 - Dependencies with specific versions
- Automation tools built around infrastructure definitions
→ easily rebuild servers (instead of updating running servers: “immutable infrastructure”)
- Versioning and sharing of infrastructure definitions



Overview

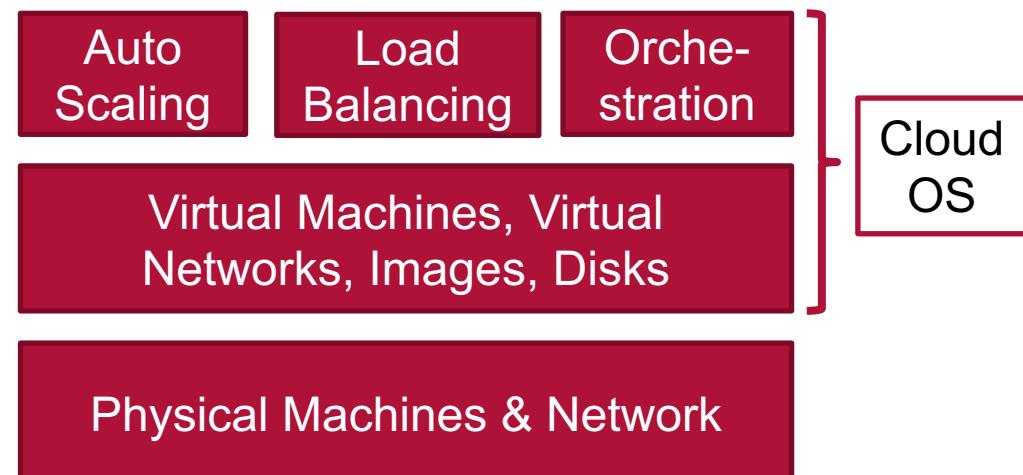
- Intro
- **Cloud Operating Systems**
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

Cloud Operating Systems (1/3)

- Virtualization is immensely useful, but managing many virtual machines manually is impractical
- Therefore: Cloud Operating systems
 - Control large pools of compute, storage, and networking resources
 - Provides dashboards and APIs for datacenter operators (administration) and users (provisioning)
- Open Source cloud systems: OpenStack, OpenNebula
- Commercial public clouds: AWS EC2, GCE, Azure, Digital Ocean

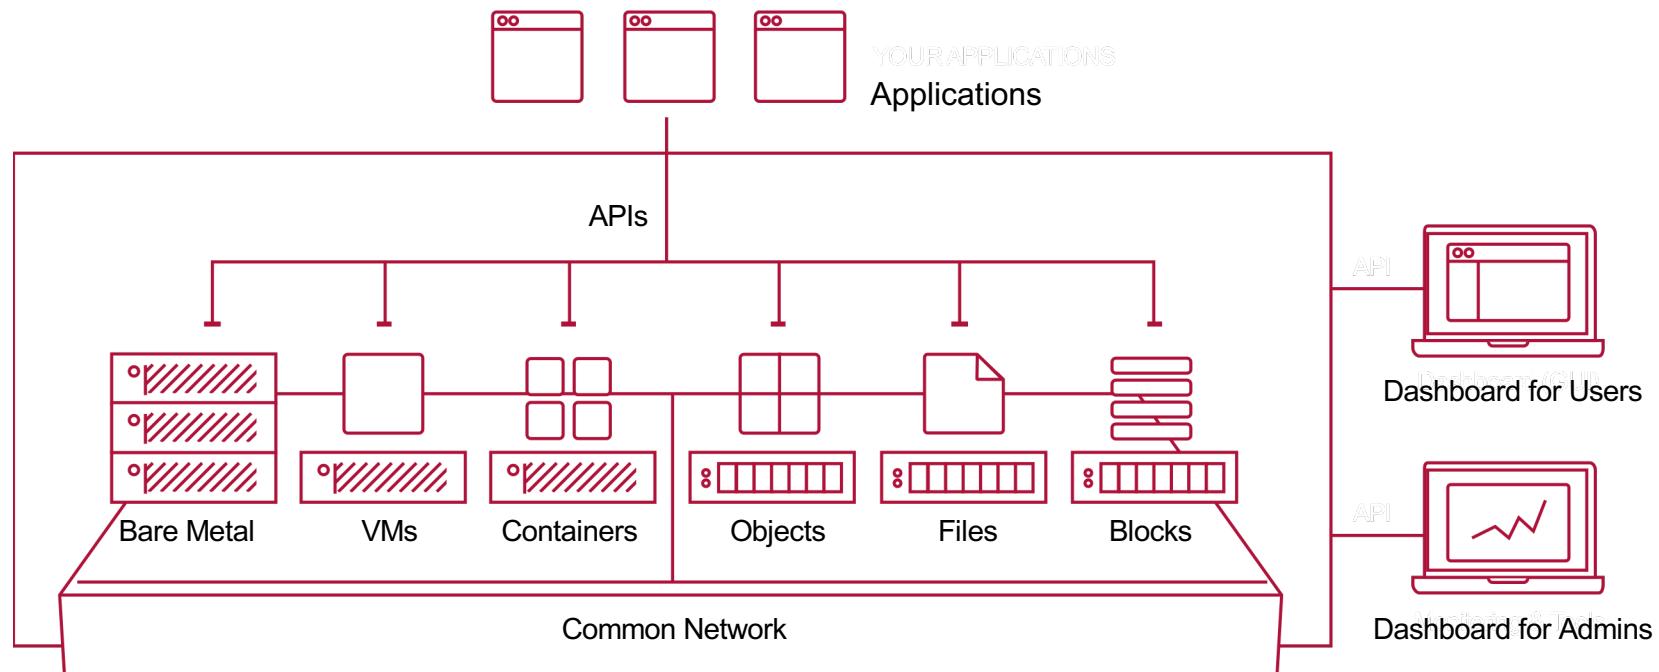
Cloud Operating Systems (2/3)

- User Interface and APIs for
 - spawn & maintain VMs,
 - virtual networking,
 - manage images & virtual disks
- Advanced features
 - load balancing,
 - auto scaling,
 - and orchestration



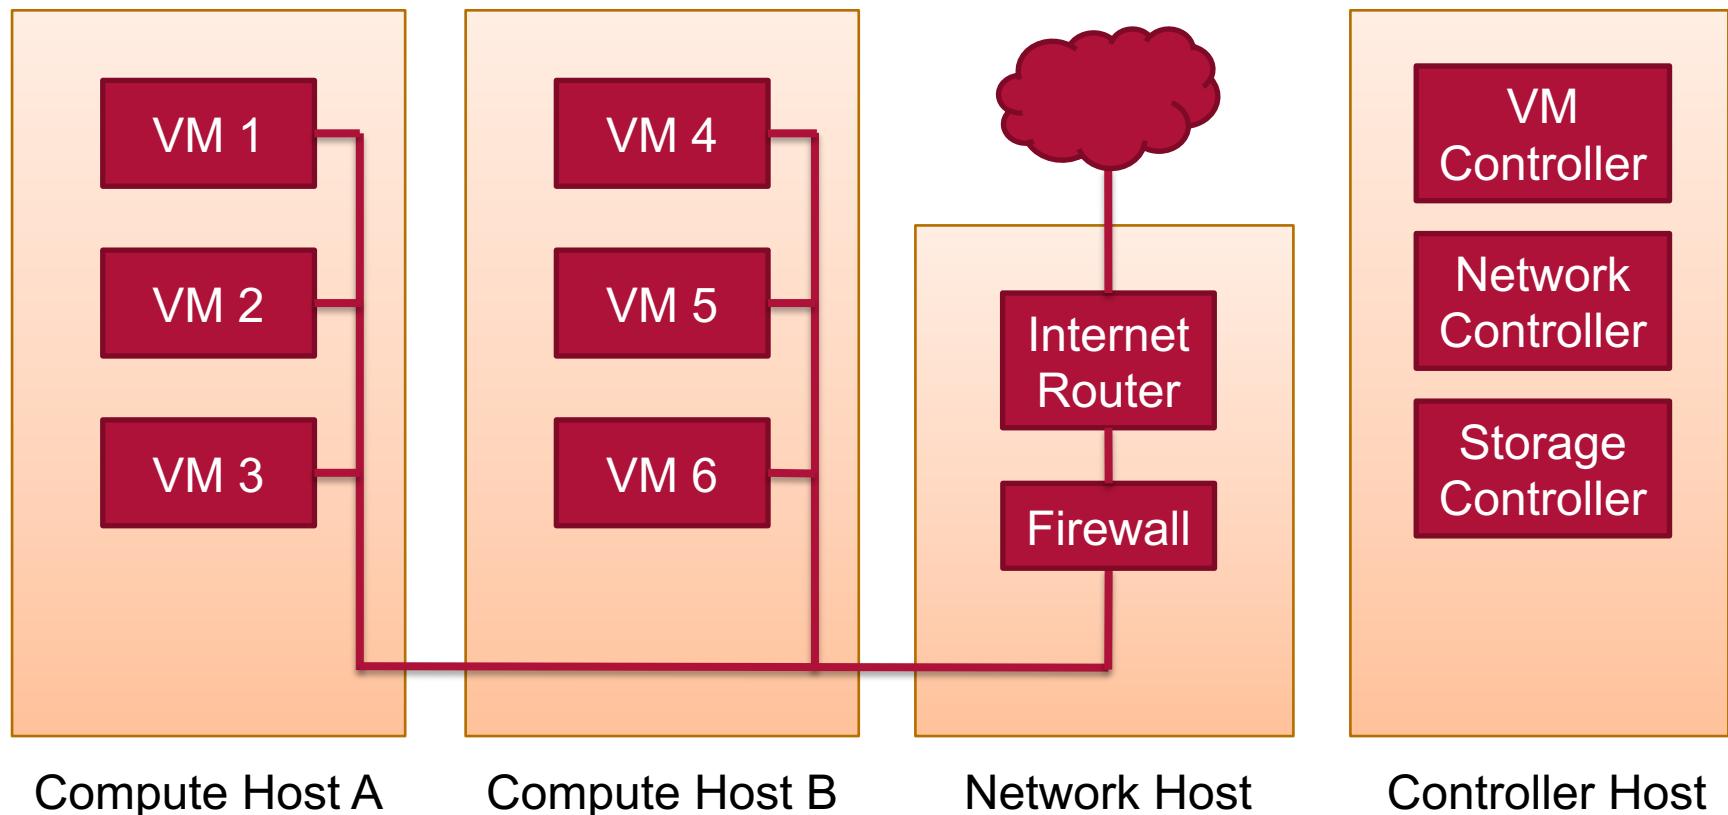
Cloud Operating Systems (3/3)

- Interfaces to manage and provision virtual resources



Managed Cloud Environment

- Different roles for physical hosts: compute hosts, storage hosts, network hosts, and controller host

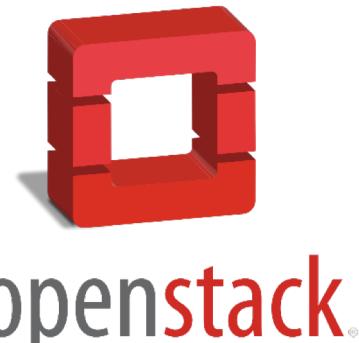


Overview

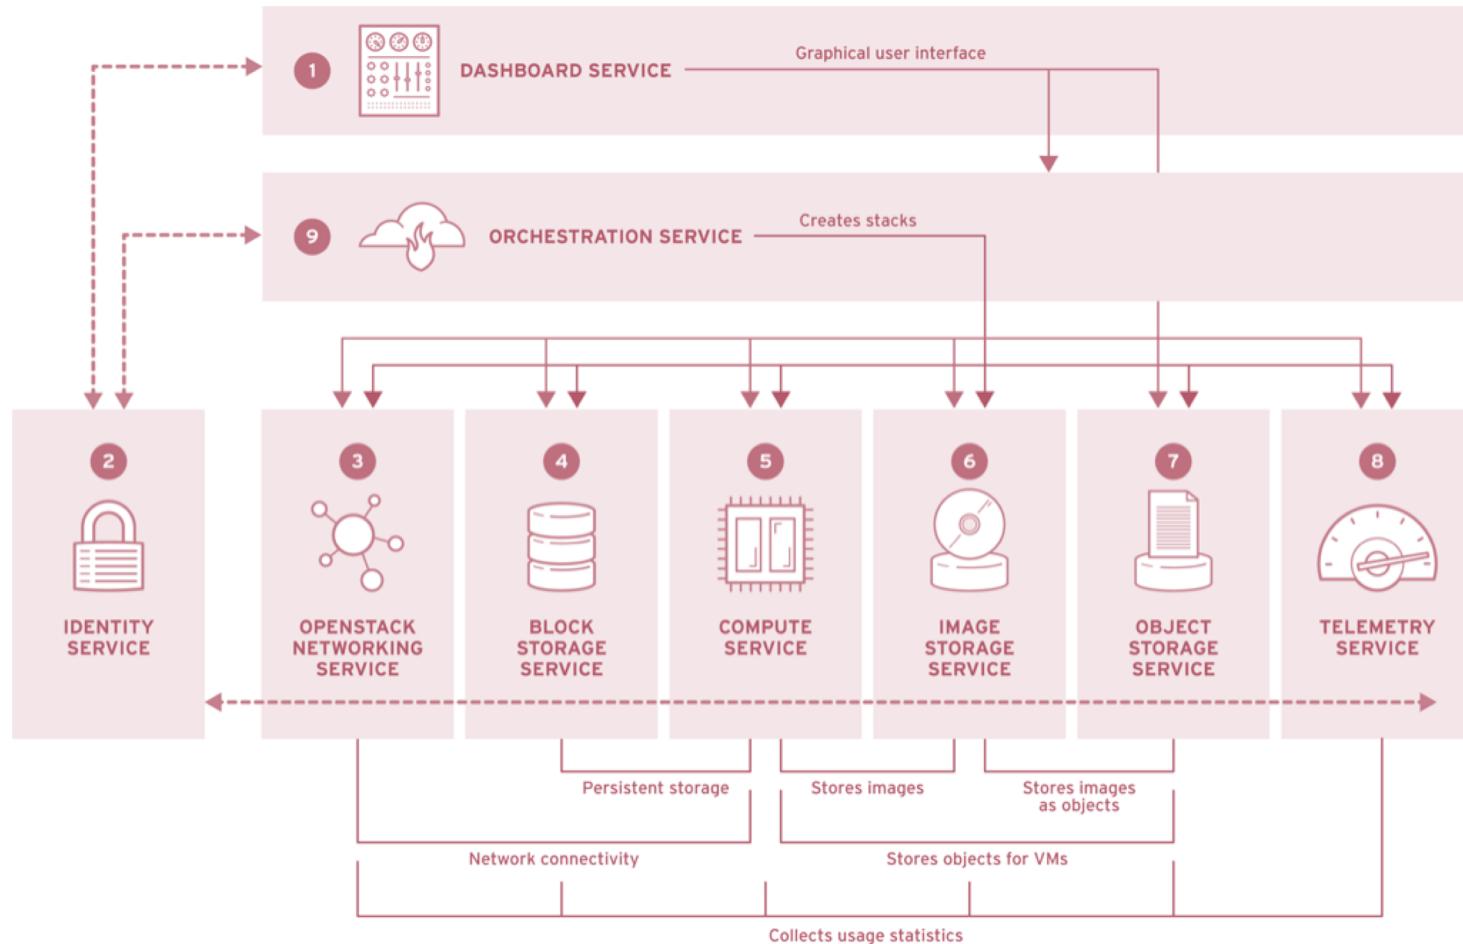
- Intro
- Cloud Operation Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

OpenStack

- Open source cloud operating system
- Mostly deployed as IaaS, providing virtual machines and other resources to users
- Started in 2010 by Rackspace Hosting and NASA
- Contributions from IBM, Red Hat, HP, Cisco, Google, Oracle, EMC, VMware
- Most popular open cloud platform in both industry and academia



Core Components of OpenStack



Compute Service: Nova

- Allows to create and manage virtual machines on demand from images
- Schedules virtual machines to run on physical nodes
- Defines drivers that interact with underlying virtualization mechanisms

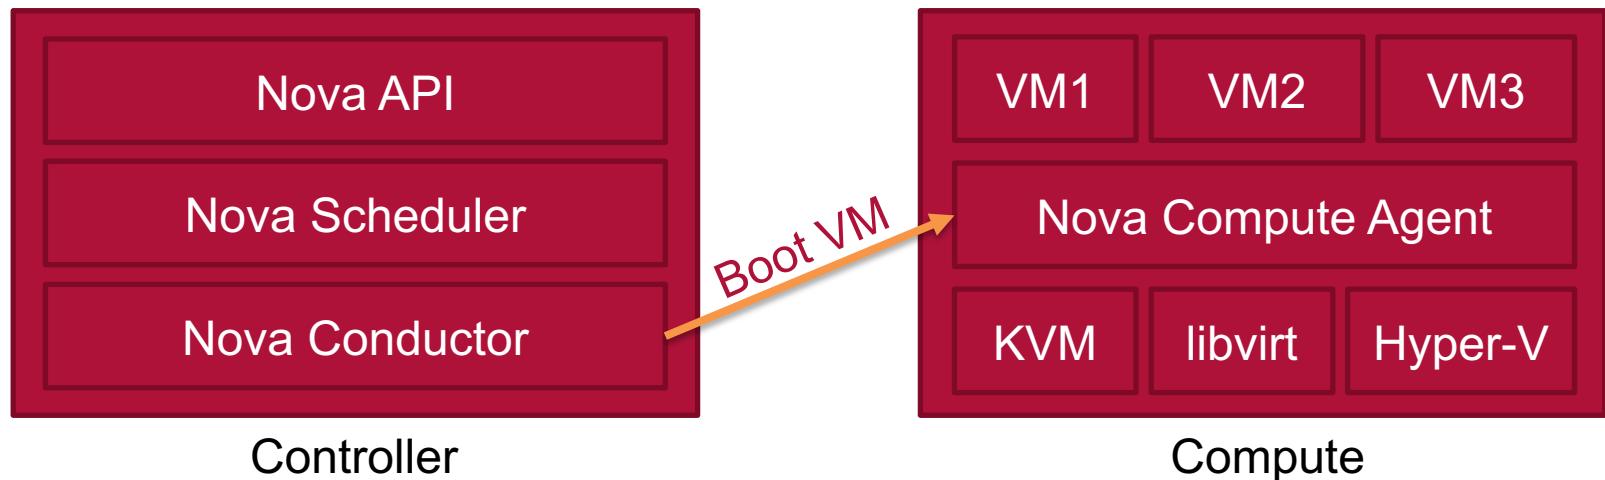
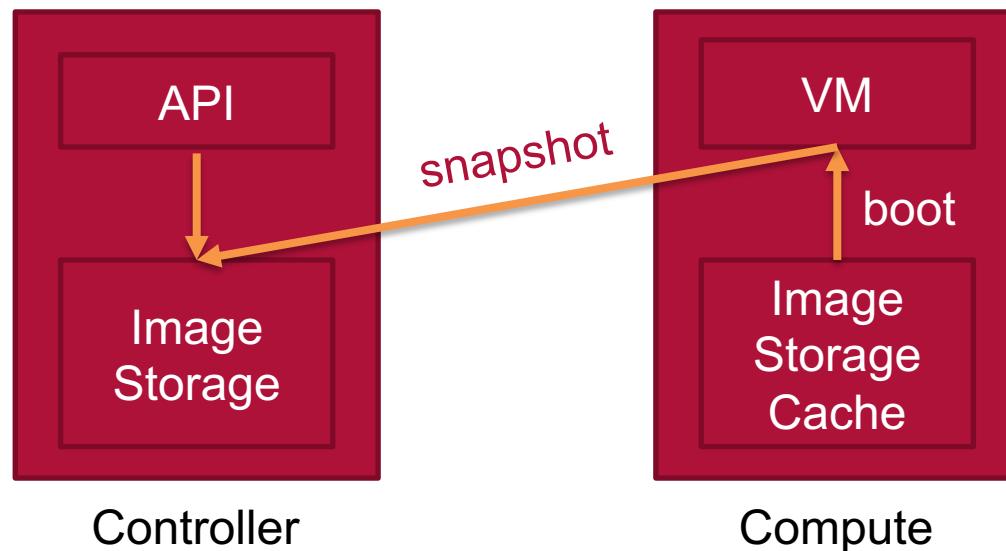


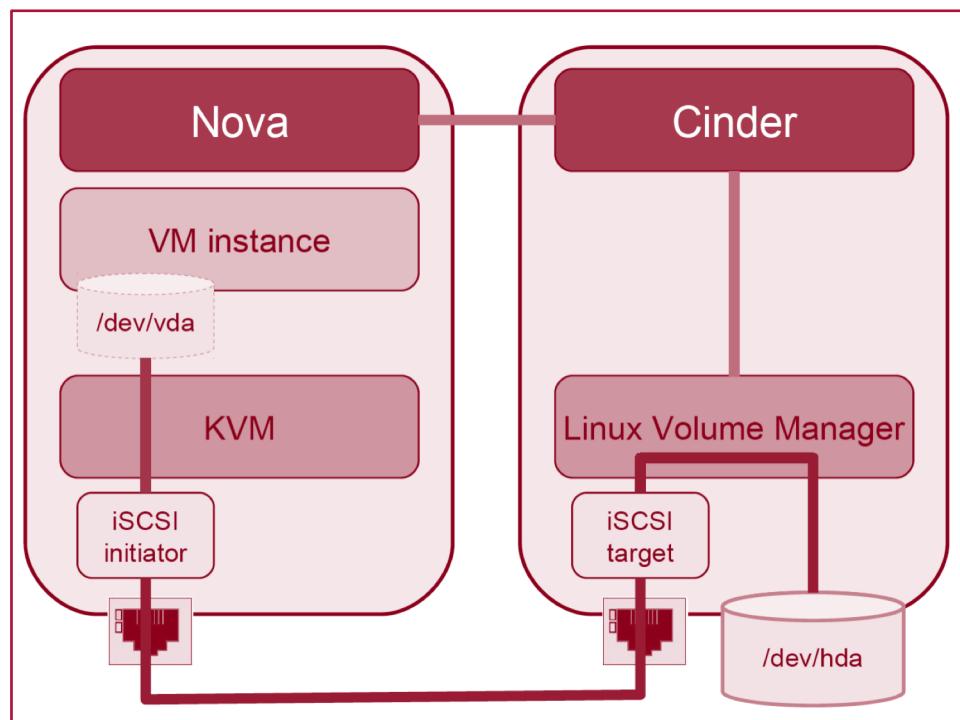
Image Service: Glance

- Registry for virtual disk images
- Users can add new images or take snapshots of existing VMs
- Snapshots can be used as templates for new servers



Block Storage Service: Cinder

- Provides network-backed virtual block storage volumes
- Enables live migration of VMs, supports replication
- Multiple storage backends



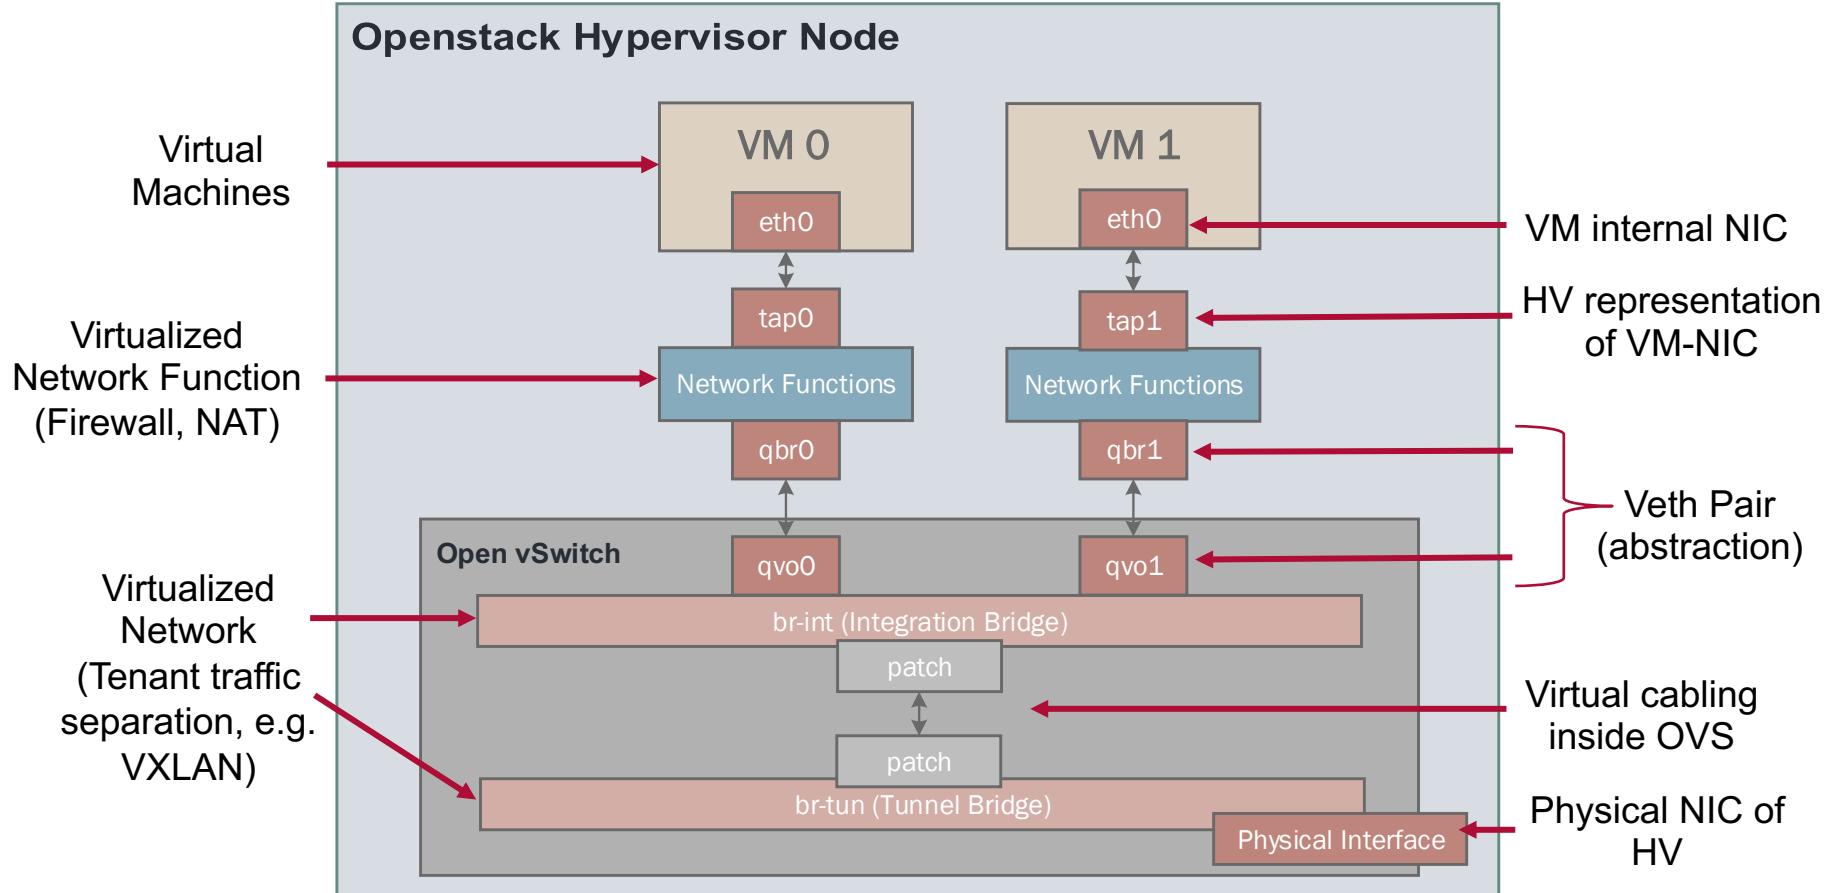
Virtual Switches & SDN

- Virtual Switches: Virtual switching stack that enables network automation through programmatic extensions and standard protocols (e.g. OpenFlow)
- Software-Defined Networking (SDN)
 - Goal: Simplify network administration
 - Separation of control plane from data plane
 - Allows on the fly configuration of switches (controller-based decision making)

Virtual Networking

- Networking between virtual machines and physical network
- Interconnects virtual machines
 - within one hypervisor
 - between several hypervisors
 - with the Internet
- Provide additional network functions
 - Like DNS, DHCP, Firewalls ...

Virtual Networking in OpenStack



Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- **Infrastructure-as-Code**
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

Motivation

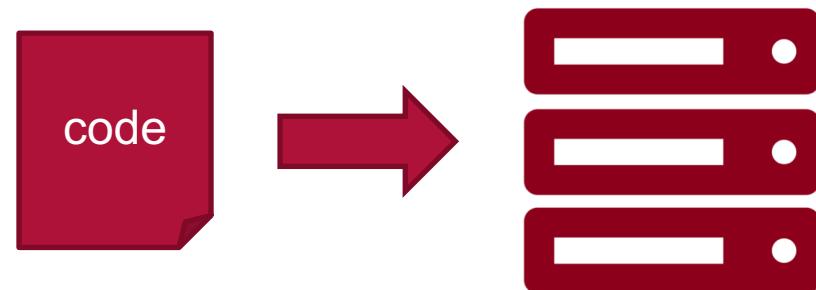
- Manual configuration of servers usually leads to
 - *Configuration Drift*: change configuration of single servers manually over time (without documentation)
 - *Snowflake Servers*:
 - ◆ Unique configuration
 - ◆ Not reproducible when hardware dies
 - ◆ Difficult to mirror for a test environment
 - ◆ Deliberate vs. default configuration
 - ◆ Host-specific vs. general configuration

Open Questions

- What if a server fails?
 - How to mirror a production environment for testing?
 - How to keep a cluster of servers consistent?
-
- Disk snapshots allow to recover and mirror servers, but
 - not straightforward to understand the state of servers
 - don't allow managing multiple servers that are mostly similar, but specific in a few configurations

Infrastructure as Code

- Goal: Make systems easily reproducible
- Solution: One single source of truth with a clear declaration of host-specific configuration → infrastructure definition files
 - That can be executed
 - That can be versioned
 - That can be shared
- Only use these definitions and automation tools, without manually updating single servers



Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

Ansible

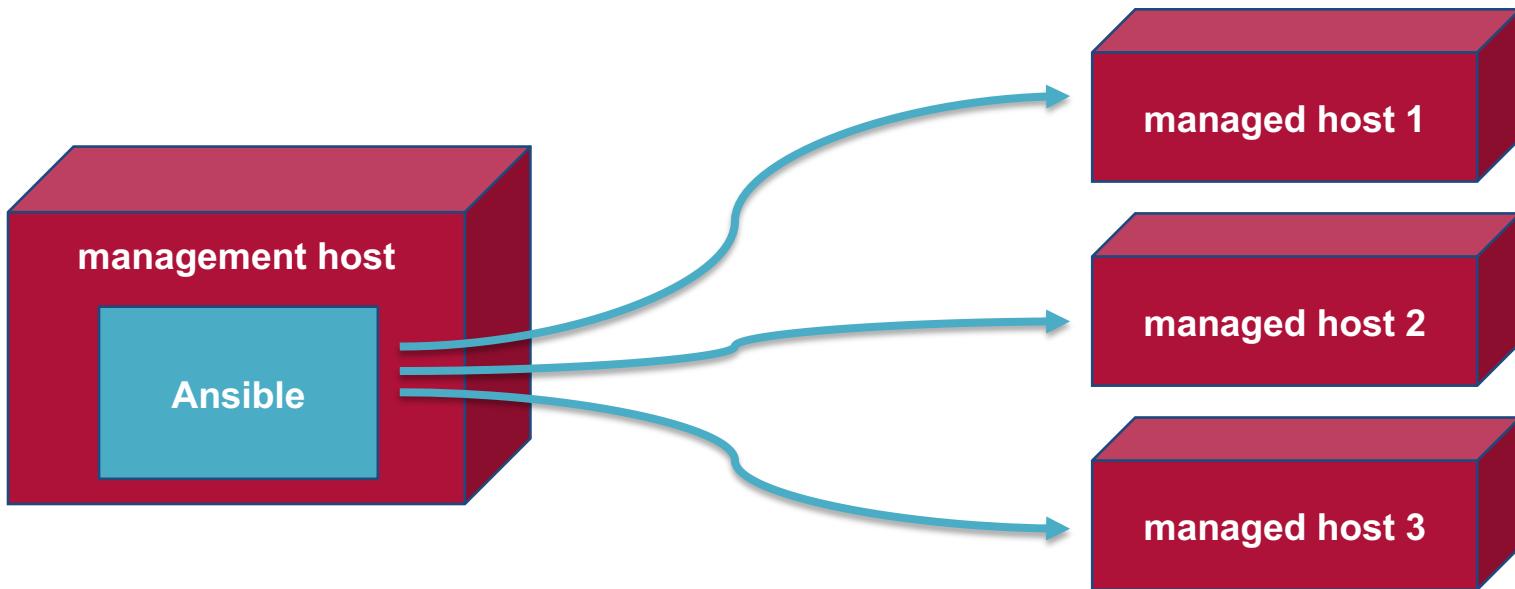
- Automation language and engine
- Configuration management
 - i.e. configure existing servers
- Orchestration
 - i.e. allocate and provision new servers
- Since 2012, acquired by Red Hat in October 2015
- Well-documented conventions and best practices, but no artificial limitations, so users can choose to ignore



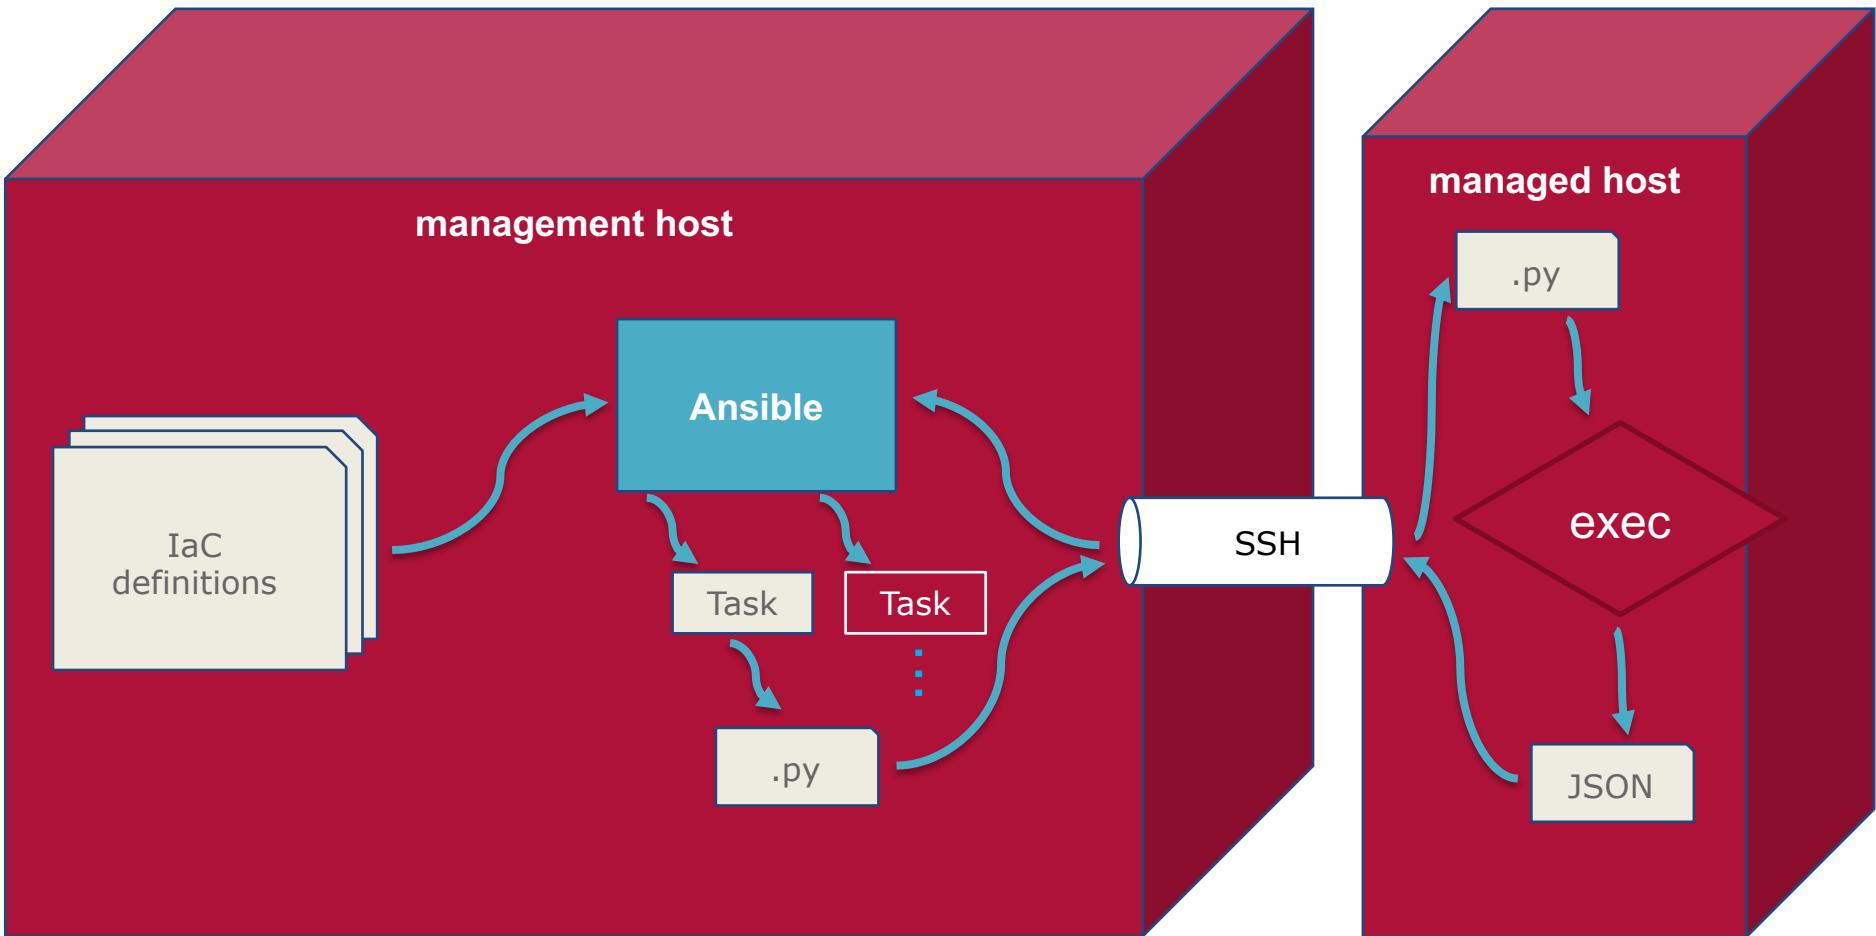
Ansible

- Multi-platform
 - Operating systems: Linux, Unix-like, MacOS
 - Orchestration-related tools: LXC, libvirt, VMware, ...
 - IaaS platforms: AWS, Azure, OpenStack, ...
- Well-established technologies as foundation
 - Python
 - SSH (PowerShell remoting, respectively)
 - YAML + Jinja templating
- Design: Declarative and agent-less

Agent-Less Architecture (1/2)



Agent-Less Architecture (2/2)



Ansible Tasks (1/2)

- Declarative
 - I.e., define the desired state
- Opposed to imperative
 - I.e., define how to achieve the desired state
- Implies idempotency
 - I.e., applying the configuration again leads to same state
 - ◆ # that's the intention; you can, of course, create non-idempotent tasks

```
# an example "task":  
- name: "enforce permissions for .ssh directory of {{ ansible_user_id }}"  
  # ↑ "name" is optional but strongly encouraged  
  file: # ← the module name, ↓ followed by module parameters  
    dest: "{{ ansible_user_dir }}/.ssh"  
    state: directory  
    mode: u+rX,go-rwx  
    recurse: yes
```

Ansible Tasks (2/2)

But to which hosts should this task be applied to?

```
# an example "task":  
- name: "enforce permissions for .ssh directory of {{ ansible_user_id }}"  
  # ↑ "name" is optional but strongly encouraged  
  file: # ← the module name, ↓ followed by module parameters  
    dest: "{{ ansible_user_dir }}/.ssh"  
    state: directory  
    mode: u+rX,go-rwx  
    recurse: yes
```

Ansible Playbooks (1/2)

- Playbooks map everything to hosts

- Tasks
- Roles
- Variables

- Further features include:

- Rolling updates
 - ◆ Configure only n hosts at a time
 - ◆ Configure only x% of host at a time
- Tolerate percentage of hosts where configuration fails
- ...

```
# file: my_playbook.yml

# a "Play":
- name: configure internal Web servers
hosts: webservers
roles:
  - webserver
tasks:
  - name: forbid external connections
    iptables:
      chain: INPUT
      source: !192.168.1.1/24
      jump: DROP
```

Ansible Playbooks (2/2)

But how does Ansible know the hosts "webservers"?

```
# file: my_playbook.yml

# a "Play":
- name: configure internal Web servers
hosts: webservers
roles:
  - webserver
tasks:
  - name: forbid external connections
iptables:
  chain: INPUT
  source: !192.168.1.1/24
  jump: DROP
```

Ansible Inventory

```
# no FQDNs required, as long as  
# ``ssh <hostname>`` works (~/.ssh/config)  
mail
```

[webservers]

```
w1.example.com  
w2.example.com
```

[dbservers]

```
d1.example.com  
d2.example.com
```

[appservers]

```
# we can override variables:  
a01.example.com www_user=web  
# we can use a pattern-like syntax:  
a[02-12].example.com
```

```
# we can define variables for groups:
```

```
[appservers:vars]  
www_user=www-data
```

```
# a group of groups
```

```
[rootservers:children]  
appservers  
dbservers
```

```
# alternatively, we can use YAML:
```

all:

hosts:

```
mail.example.com:
```

children:

webservers:

hosts:

```
w1.example.com:  
w2.example.com:
```

dbservers:

hosts:

```
d1.example.com:  
d2.example.com:
```

```
...
```

- Inventories can also be obtained from external scripts / executables (think: IaaS)

Ansible Variables (1/2)

- Variables can be defined in many many places
 - Playbooks, Tasks, Inventory, CLI args, roles, facts, ...
 - But, scattering variables all over the place is discouraged, of course
 - e.g., using role, group and host variables can suffice
- "facts"
 - A heap of auto-detected variables: e.g., Uptime, chroot?, user home, SSH host key, Python version, total and free memory, host name, FQDN, environment variables, block devices, mounts, interfaces, BIOS version, processor specs, OS name and version, ...
- Jinja templating also in variables files
 - e.g., `app_user: "{{ ansible_env.SUDO_USER|default(ansible_user_id) }}"`

Ansible Variables (2/2)

■ Host and group variables

- ◆ e.g., in project directory:

```
# e.g., set inventory to ./inventory.ini:  
.ansible.cfg  
inventory.ini  
my_playbook.yml  
group_vars/  
# applies to all hosts:  
all.yml  
webservers.yml  
host_vars/  
mail.yml
```

```
# no FQDNs required, as I  
# ``ssh <hostname>`` will  
mail
```

```
[webservers]  
w1.example.com  
w2.example.com
```

```
[dbservers]  
d1.example.com  
d2.example.com
```

```
[appservers]  
# we can override variables  
a01.example.com www_user=www-data  
# we can use a pattern-like variable  
a[02-12].example.com
```

```
# we can define variables  
[appservers:vars]  
www_user=www-data
```

```
# a group of groups  
[rootservers:children]  
appservers  
dbservers
```

Ansible Roles (1/2)

■ Roles

- ◆ e.g., in project directory:

roles/webserver

```
# default variables:  
defaults/main.yml  
# tasks to be executed:  
tasks/  
  main.yml  
  # ↑ might include:  
  certbot.yml
```

...

Ansible Roles (2/2)

■ Roles

- ◆ e.g., in project directory:

roles/webserver

...

```
# e.g., restart service on
# configuration changes:
handlers/main.yml
# e.g., declare dependencies to
# other roles:
meta/main.yml
```

Summary of IaC in Ansible

- Inventory
 - List of managed hosts
 - Groups and groups of groups
- Roles
 - Concept to reuse tasks and variables for similar hosts
 - Usually contain (default) variables
- Host and group variables
 - This is where variables should go
- Playbooks (Ansible's execution unit)
 - "map" roles, tasks, variables to hosts from inventory

Run a Playbook

```
$ cat my_playbook.yml
```

```
- name: test SSH connection
hosts: pim:static
tasks:
  - name: test if logging into host and running a module remotely works
    ping: # does not require any arguments
```

```
$ ansible-playbook my_playbook.yml
```

```
PLAY [test SSH connection] ****
TASK [Gathering Facts] ****
ok: [static]
ok: [pim]
```

```
TASK [test if logging in and running a module remotely works] ****
ok: [pim]
ok: [static]
```

```
PLAY RECAP ****
pim                  : ok=2    changed=0   unreachable=0   failed=0
static               : ok=2    changed=0   unreachable=0   failed=0
```

Run adhoc commands

```
$ ansible vs2_and_containers -m file -a "path=/run/log/mylogs state=directory"
vs2 | FAILED! => {
    "changed": false,
    "msg": "There was an issue creating ...: [Errno 13] Permission denied: '/run/log/mylogs'",
    "path": "/run/log/mylogs",
    "state": "absent"
}
static | CHANGED => {
    "changed": true,
    "gid": 0,
    "group": "root",
    "mode": "0700",
    "owner": "root",
    ...
}
rproxy | SUCCESS => {
    "changed": false,
    ...
}
...
lists | UNREACHABLE! => {
    "changed": false,
    "msg": "SSH Error: data could not be sent .... Make sure this host can be reached over ssh",
    "unreachable": true
}
```

"But my task is dead-simple ...

... using Ansible would be overkill."

- How does your shell script handle ...
 - ... all the return codes?
 - ... whitespaces in variables and undefined variables?
 - ◆ e.g., \$ rm -rf /\$x
 - ... being re-run – do we still get the desired result?
 - ... changes in e.g. CLI APIs?
 - ... host-specific configurations?
 - ... shell interoperability?
 - ... different platforms?

"But my task is dead-simple ...

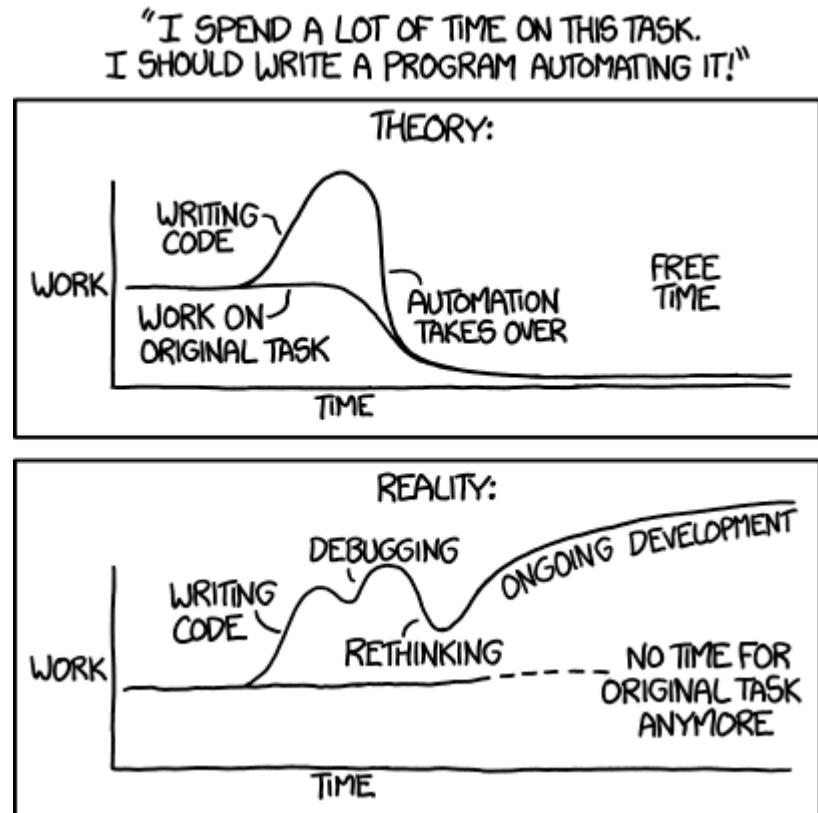
- How does your shell script handle ...
 - ... if you need to run it only partially?
 - ... if it needs to be understood by humans?
 - ... if it needs to be run sequentially on multiple hosts?
 - ... if it needs to be run on multiple hosts in parallel?
 - ...
- → simple tasks usually turn out harder than expected

we're talking about Ansible here, but other configuration management/IaC tools would get the job done as well – a matter of requirements and taste

Open Questions

not specific to Ansible

- When to automate?
 - When expected to do task more than two times?
 - Depends on complexity
- Programming vs. markup
 - When to write a, e.g., Python program instead of using, e.g., Ansible?



<https://xkcd.com/1319/>

Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

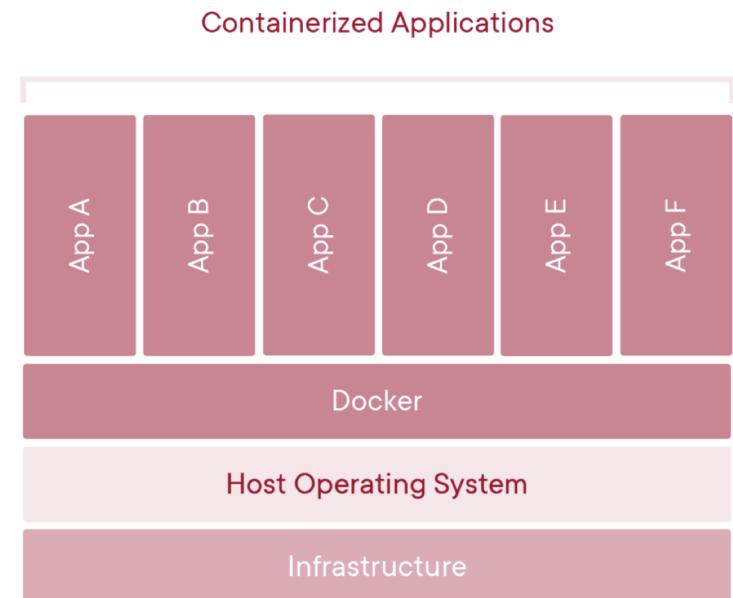
Container Analogy

- Containers
 - Standardized
 - Easy to move
 - Isolated
 - Many containers fit hosts
- Dependency management: libs and config bundled with the application, to run it everywhere → build your app in a container, test the container locally or your staging environment, then ship the container



Recap: Docker

- Container instance: running isolated process, started from a Docker image
- Docker image: snapshot of a container, packaging the application code and all dependencies
- Docker file: instructions for creating images
- Multiple containers share single OS kernel, isolated by kernel containment features (e.g. cgroups)



Dockerfiles

- Creating docker images from configuration files
- Example: Dockerfile for an Express web server image using NodeJS

```
$ docker build --tag hello .
$ docker run --rm hello
```

```
# Create image based on the official Node 6 image from dockerhub
FROM node:9

# Create a dir where our app will be placed
RUN mkdir -p /usr/src/app

# Change dir so that our commands run inside this new dir
WORKDIR /usr/src/app

# Copy dependency definitions
COPY package.json /usr/src/app

# Install dependecies
RUN npm install

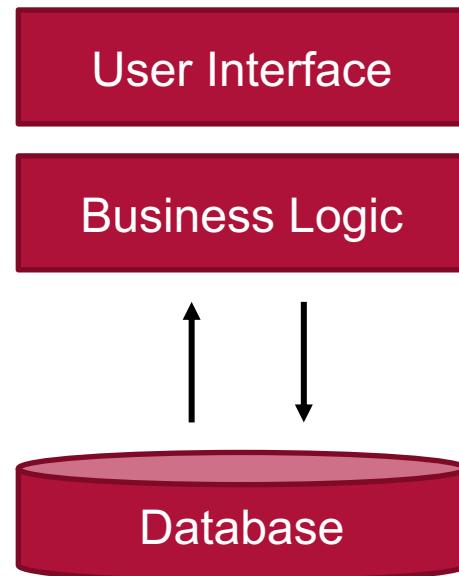
# Get all the code needed to run the app
COPY . /usr/src/app

# Expose the port the app runs in
EXPOSE 3000

# Serve the app
CMD ["npm", "start"]
```

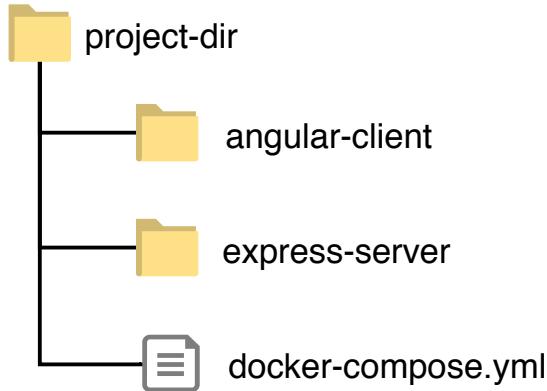
Container Composition

- Applications typically consist of multiple connected components intended to run as a single service
- e.g. Web Application
 - Front-end framework
 - Back-end framework
 - Persistent data store



Docker Compose

- Tool for defining and running multi-container Docker applications



```
$ docker-compose up
```

```
version: '2' # specify docker-compose version

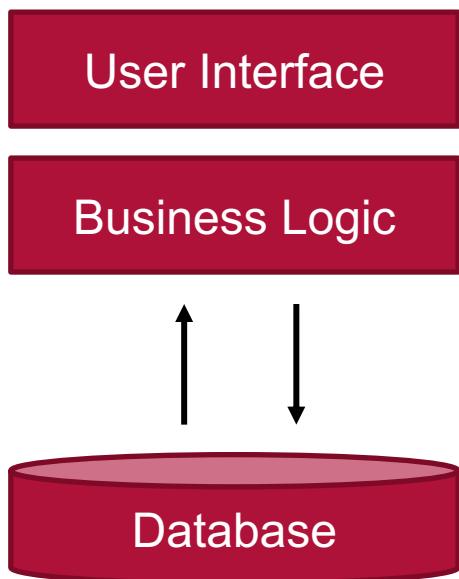
# Define the services/containers to be run
services:
    angular:
        build: angular-client # Dockerfile directory
        ports:
            - "4200:4200" # port forwarding

    express:
        build: express-server # Dockerfile directory
        ports:
            - "3000:3000" # ports forwarding

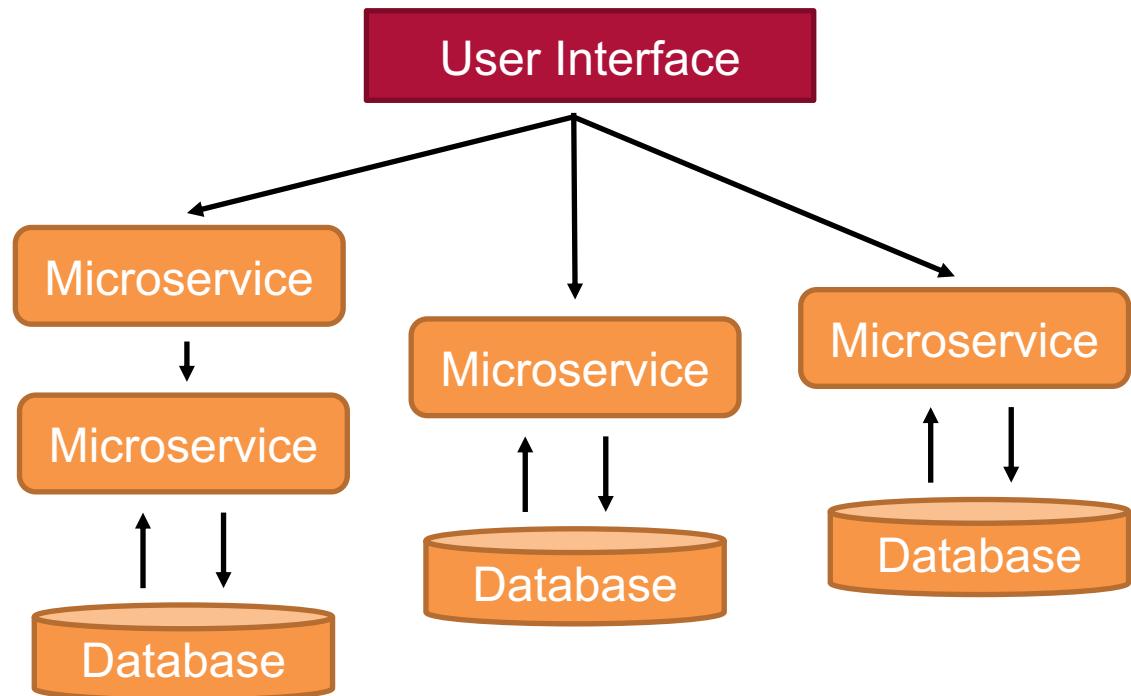
    database:
        image: mongo # image to build container from
        ports:
            - "27017:27017" # port forwarding
```

Microservices (1/2)

Monolith

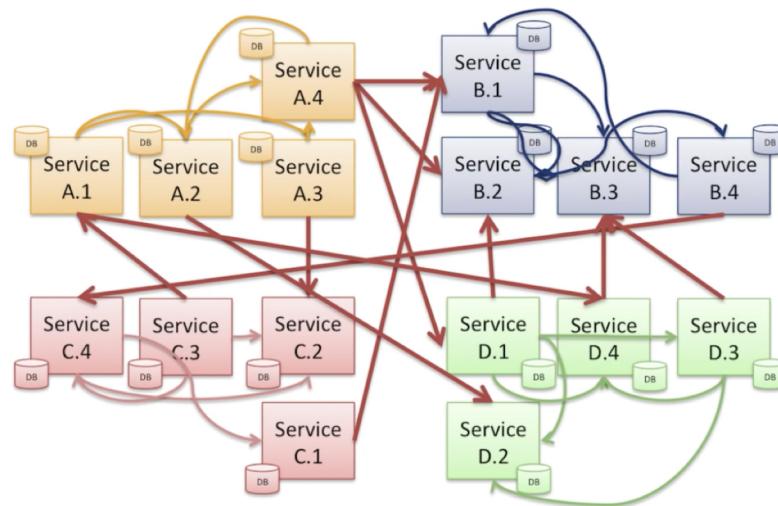


Microservices



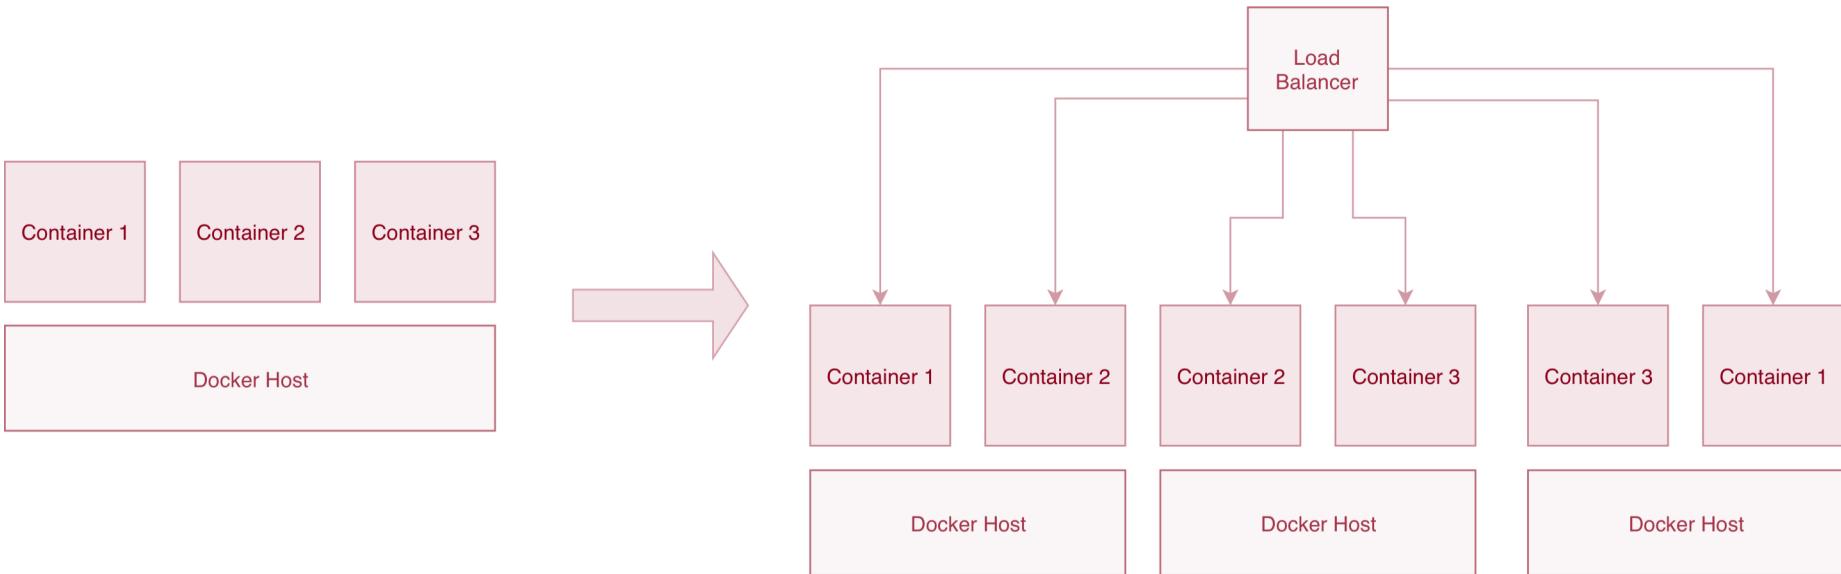
Microservices (2/2)

- Advantages
 - Independent development
 - Small teams
 - Fault isolation
 - Scalable!
- Disadvantages
 - Overhead (duplicated tech)
 - Management of services and networking



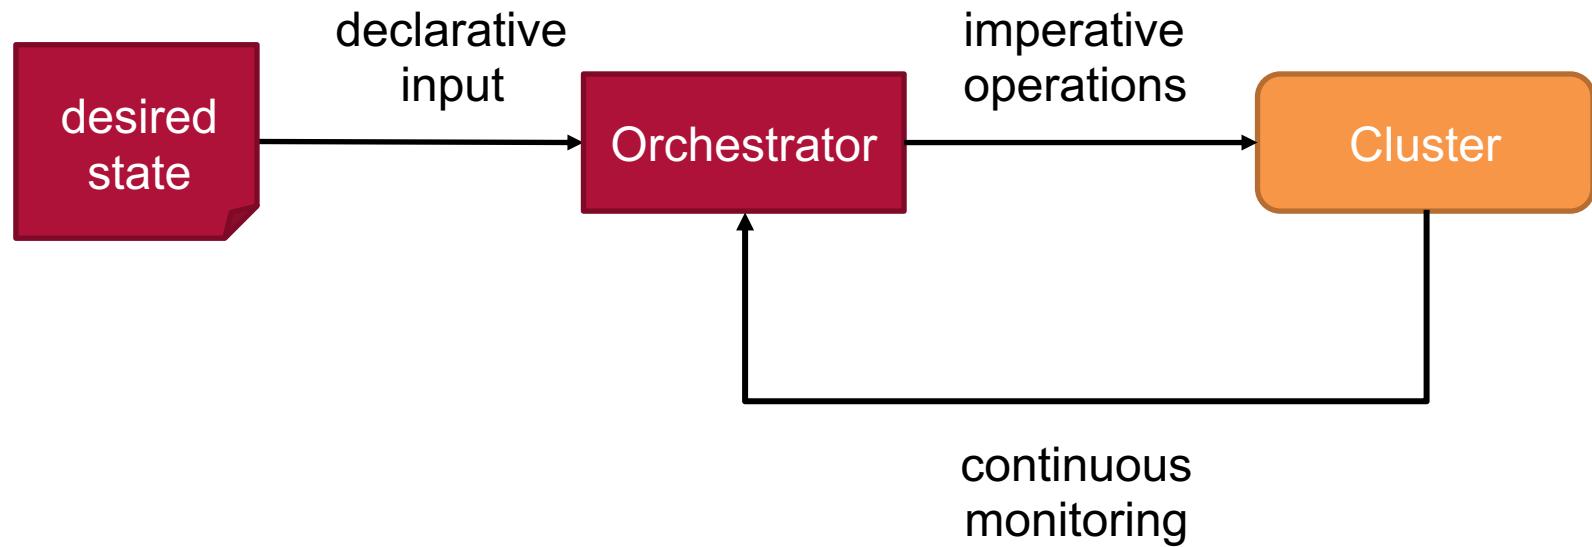
Docker Cluster

- Running an application on a cluster of docker hosts for fault tolerance and scalability



Orchestration

- Orchestration tools: Control systems for clusters



Main Orchestration Tasks

- Scheduling and placement
- Service configuration
- Networking and storage management
- Monitoring and logging
- Replication and scaling
- Re-scheduling and load balancing
- Rolling updates

Container Orchestration

Container Orchestration

- Distributed container management

Container Runtime

- Local container management

Infrastructure

- Container-agnostic infrastructure

- Provisioning and deployment of containers
- Replication and availability of containers
- Configuration and networking of containers
- Load balancing across replicated containers
- Monitoring of replicated containers

Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - **Kubernetes**
- DevOps, Continuous Integration, and Continuous Delivery

Kubernetes

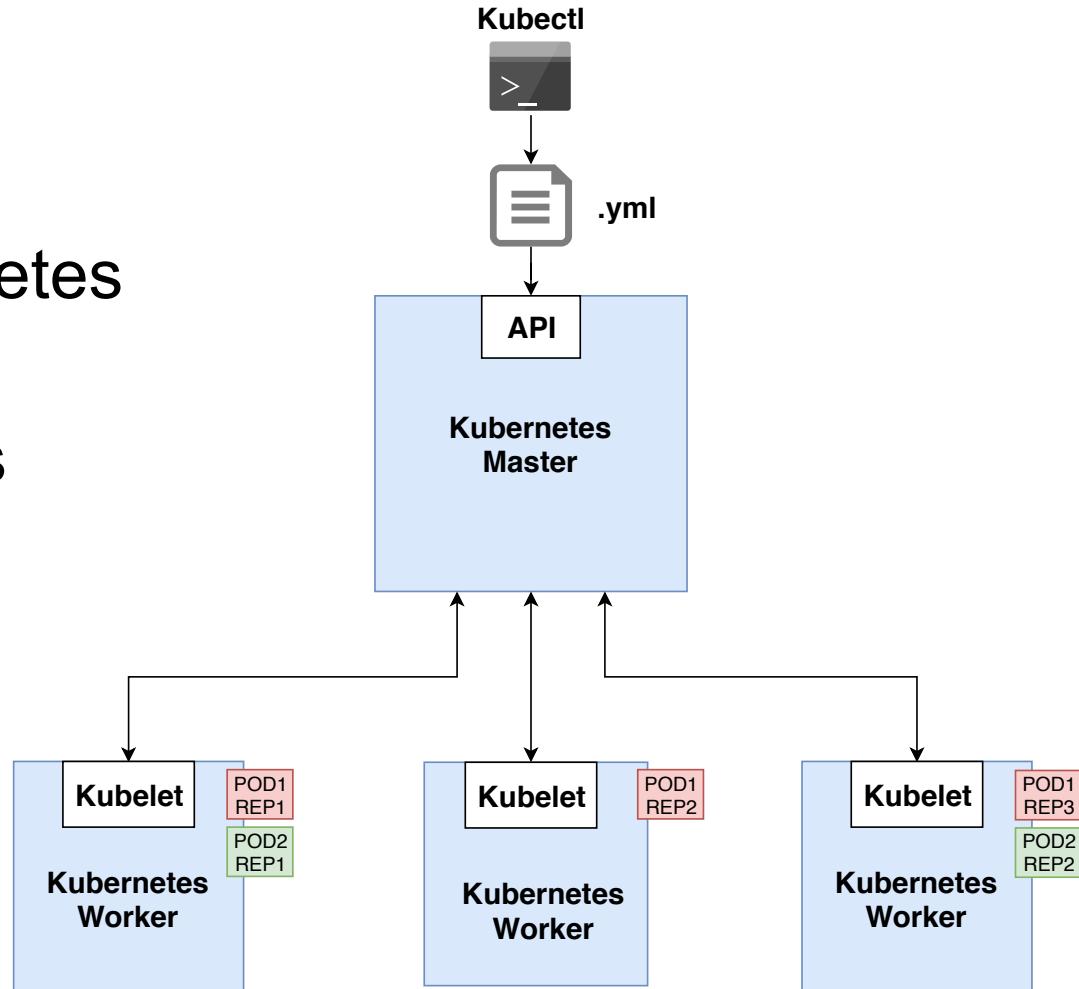
- Greek for “helmsman” or “pilot”, abbreviated “k8s”
- Open-sourced by Google in 2014
- Version 1.0 released in July 2015
- Google and the Linux Foundation created the Cloud Native Computing Foundation (CNCF) to offer k8s as an open standard
- Google Kubernetes Engine (GKE), Azure Container Service (AKS), AWS Elastic Container Service (EKS)

What Does k8s Do for You?

- Automation engine for cluster management, from a declarative description of the desired cluster state
- Deploying, monitoring, and scaling
 - Instantiating sets of containers
 - Linking containers together through agreed interfaces
 - Exposing services to machines outside the cluster
 - Monitoring, logging, and rescheduling of failed containers
 - Scaling out or down the cluster by adding or removing containers

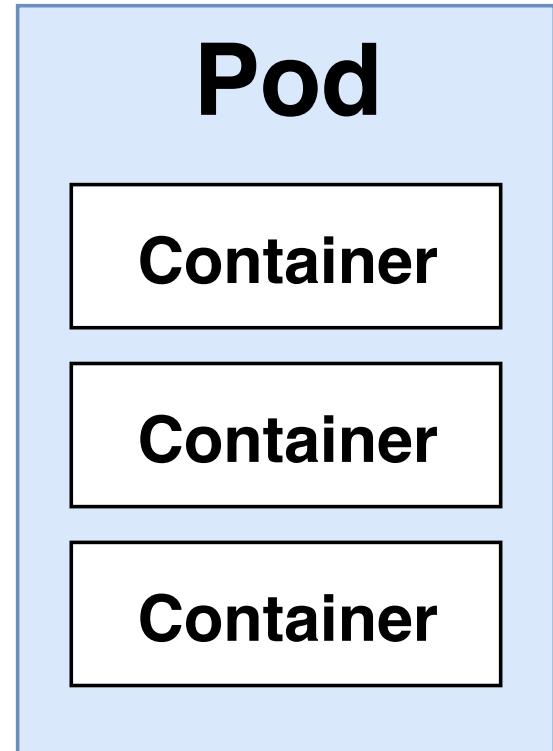
Kubernetes Cluster

- Collection of nodes (either bare-metal or virtual machines), managed by Kubernetes
- Runs groups of replicated containers (called *Pods*)



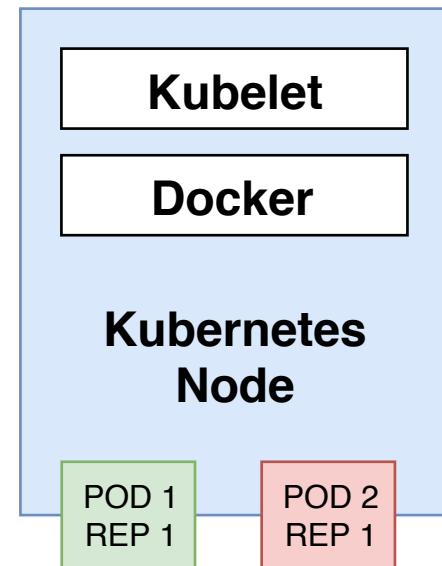
Kubernetes Pods

- Pods consist of
 - Group of containers
 - Container configurations
 - Shared storage
- Containers in a pod
 - are scheduled together
 - are guaranteed to be on the same node
- Replicas of pods



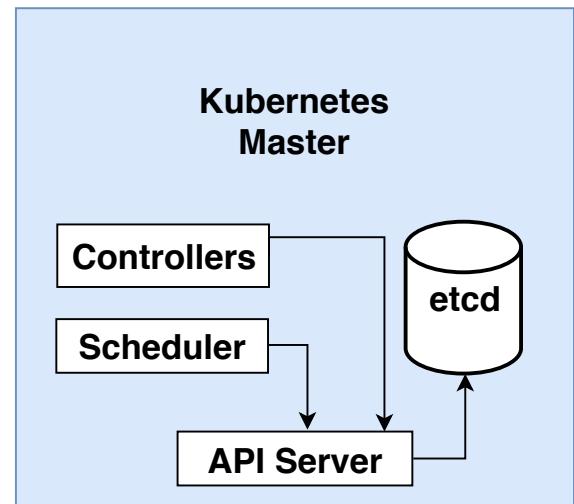
Kubernetes Nodes

- Each worker node in the cluster runs two processes:
 - kubelet: communicates with the Kubernetes master
 - kube-proxy: reflects Kubernetes networking on each node



Kubernetes Master

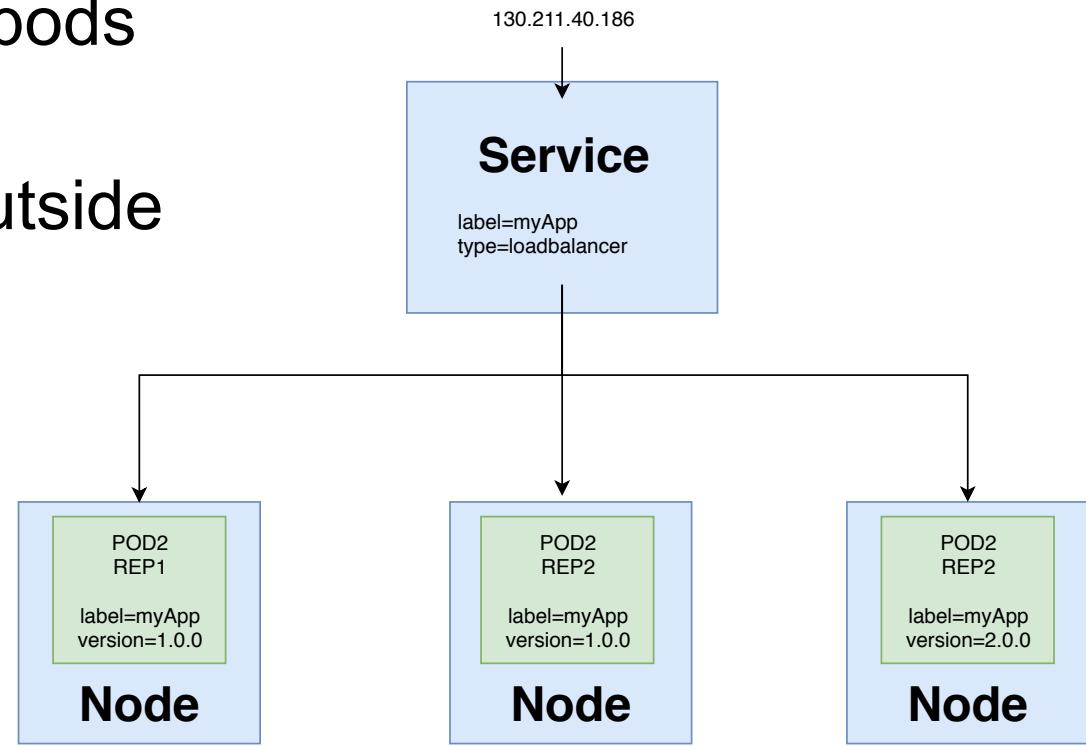
- Collection of processes managing the server state on a single node of the cluster
- Controllers, e.g. replication controller
- Scheduler: places pods based on resource requirements, hardware and software constraints, data locality, deadlines...
- etcd: key-value store for cluster state



Kubernetes Services

- Services handle and load-balance requests to pods
- Services expose functionality to the outside

```
apiVersion: v1
kind: Service
metadata:
  name: myApp-service
spec:
  type: LoadBalancer
  ports:
  - protocol: TCP
    # accessible externally
    nodePort: 80
    # port in Pod
    targetPort: 8080
  selector:
    app: myApp
```



Kubernetes Deployment

- Defines desired state, k8s then establishes and re-establishes it

```
$ kubectl create -f nginx-deployment.yaml
```

```
$ kubectl get deployments
```

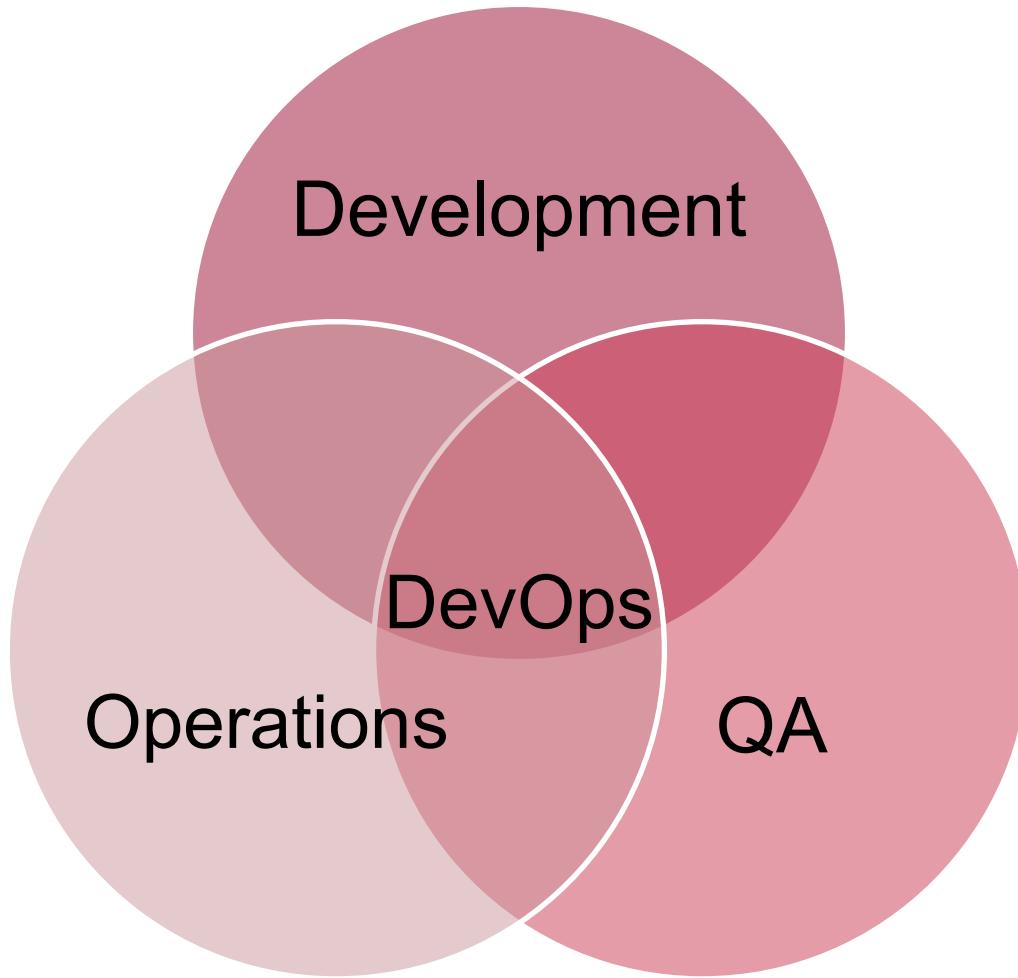
| NAME | DESIRED | CURRENT | UP-TO-DATE | AVAILABLE | AGE |
|------------------|---------|---------|------------|-----------|-----|
| nginx-deployment | 3 | 3 | 3 | 3 | 18s |

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.4
          ports:
            - containerPort: 80
```

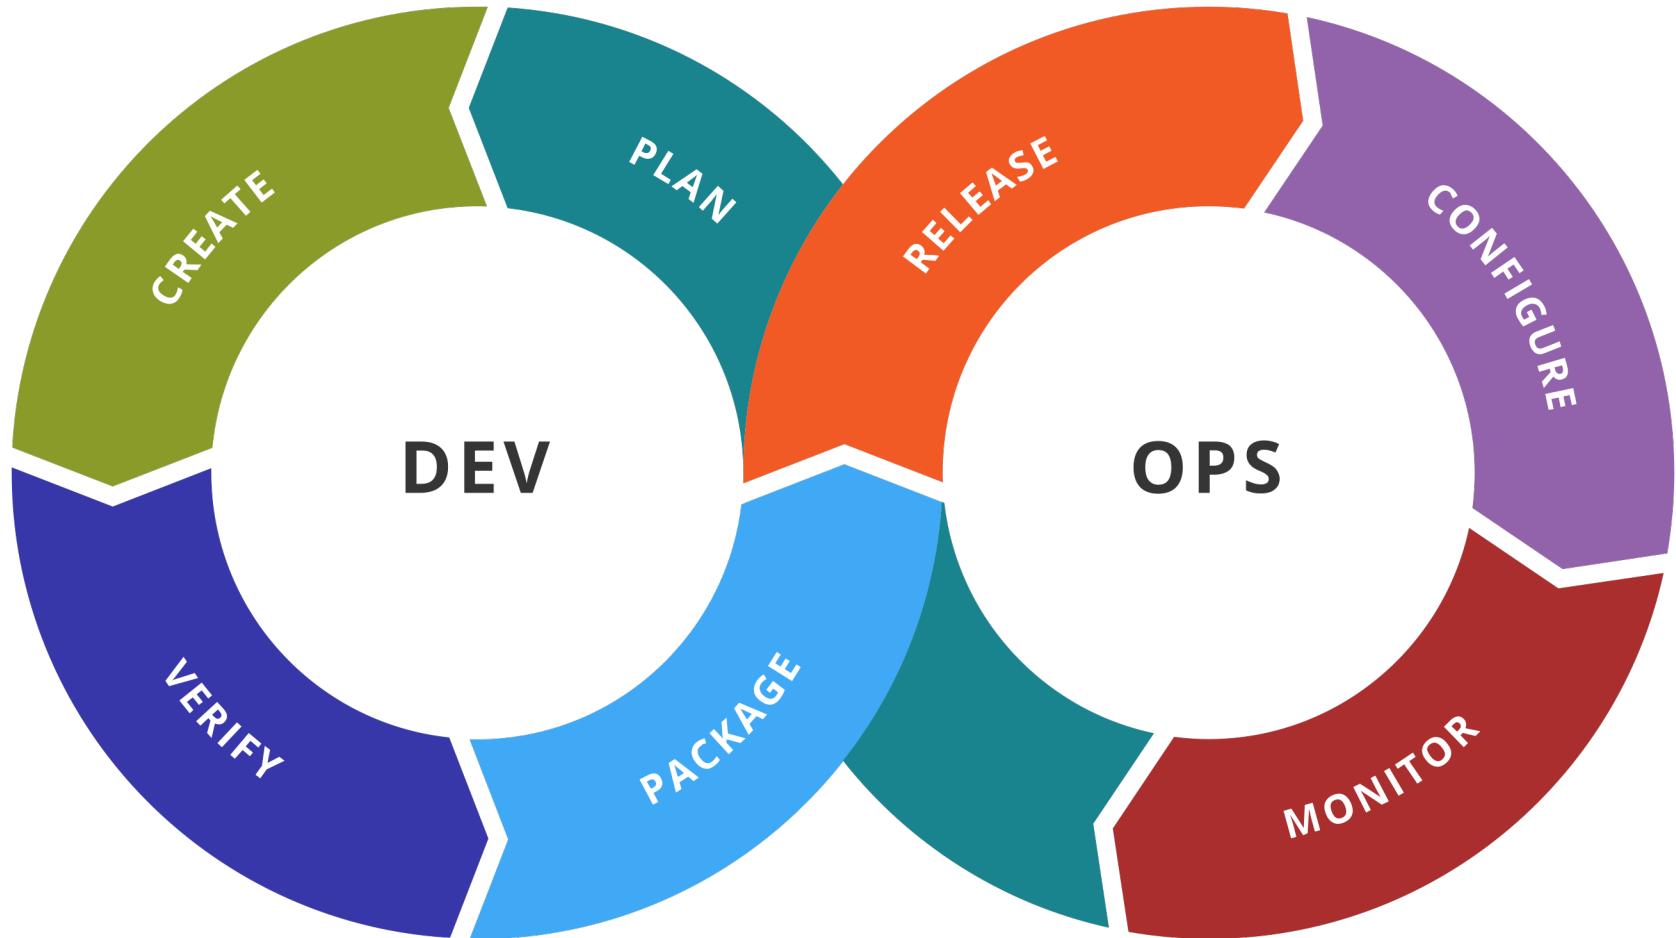
Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- **DevOps, Continuous Integration, and Continuous Delivery**

DevOps



DevOps



Continuous Integration

- Automation of integration *and testing* of software
- First proposed by Booch, 1991
- Adopted and popularized as part of XP
- Motivation:
 - Evade “integration hell”: Find bugs early, while still easy to fix
 - Immediate feedback on system-wide impact of local changes
 - Team has a common view of the state of a software project

Continuous Integration: Practices

- One common (versioned) code repository
- Build automation (and short builds)
- Self-testing builds
- Regular commits
- Building every commit (usually on a CI server)
- Test & production environments as similar as possible
- Test results visible for everyone



Continuous Delivery

- Fast and reproducible software releases
- From the first principle of the agile manifesto:

“Our highest priority is to satisfy the customer through early and continuous delivery of valuable software.”
- CD Metric: How long does it take to release a single-line-code-change to production? (*cycle time*)
- Business case: Software development as investment
→ optimizing ROI through small cycle time

Continuous Delivery: Problems of Delivering Software

- Antipatterns
 - Deploying software manually
 - Deploying to a production-like environment only after development is complete
 - Manual configuration management of production environments

Continuous Delivery vs. Continuous Deployment

Automatic

Manual

Build

Test

Deploy to
staging

Acceptance
tests

Deploy to
production

Smoke tests

Continuous Delivery

Continuous Integration

Build

Test

Deploy to
staging

Acceptance
tests

Deploy to
production

Smoke tests

Continuous Deployment

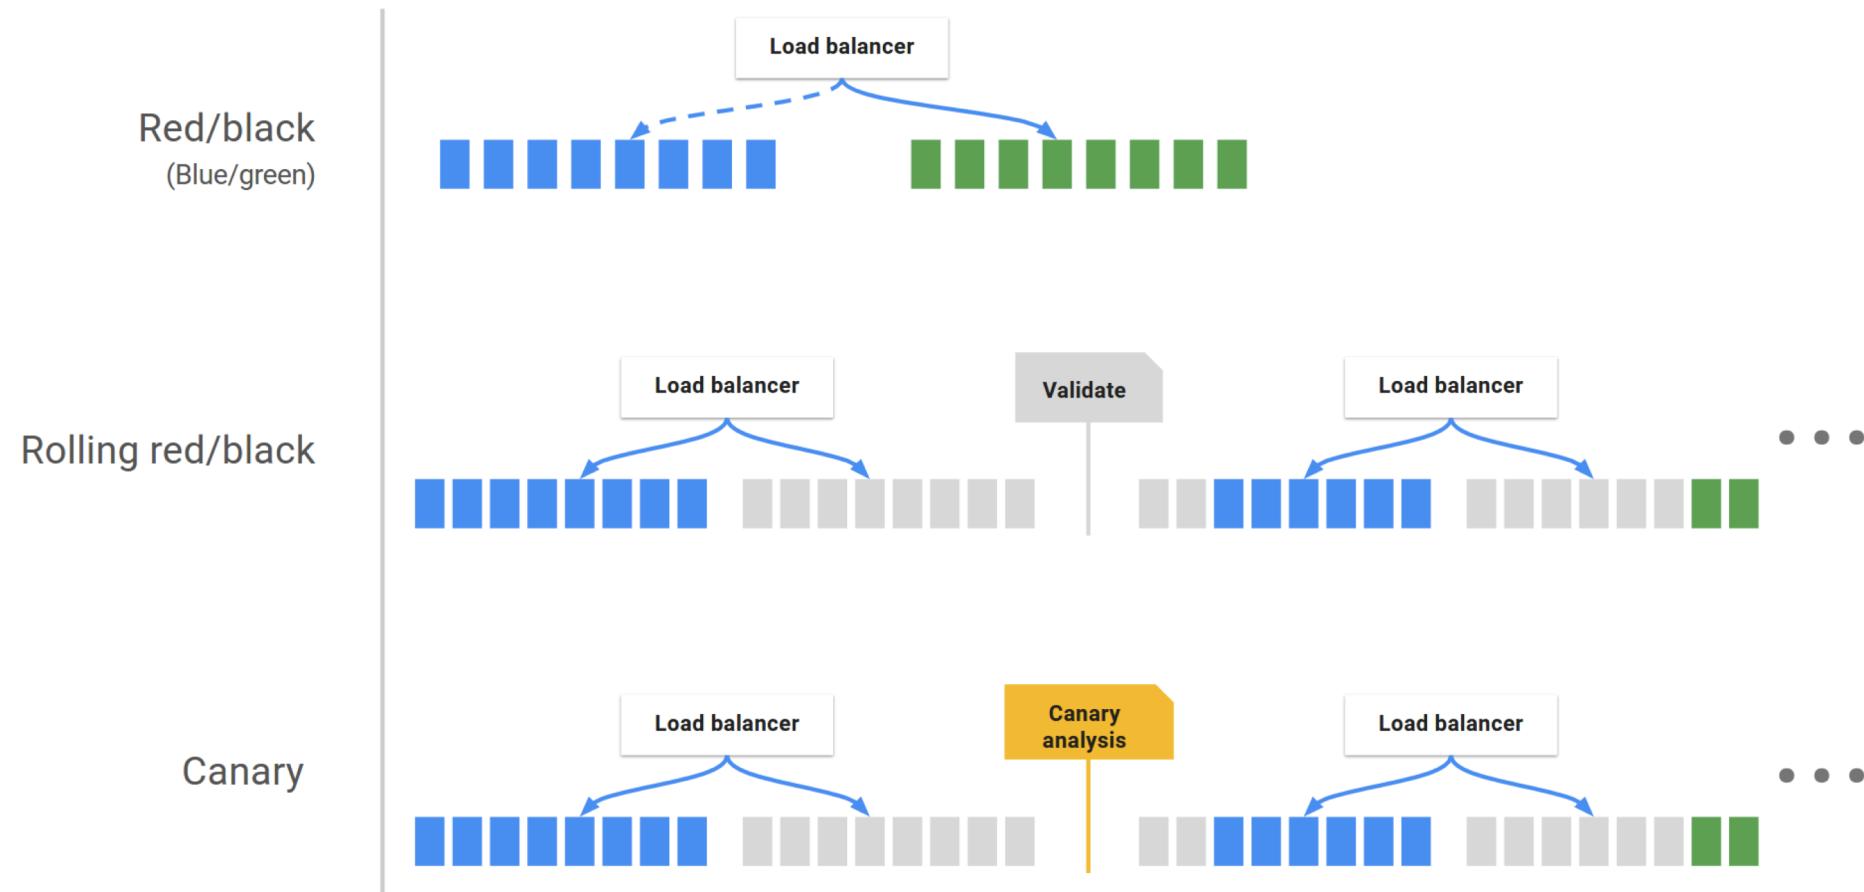
Continuous Delivery: Principles

- Parts of a working software application:
 - Executable code
 - Configuration
 - Host environment
 - Data
- Feedback on changes to any part
- Feedback as fast as possible: unit tests, acceptance tests, capacity tests

Deployment Strategies: Blue/Green

- Strategy to deploy an upgrade
- Two production environments: Blue and green
- Only one environments is active at a time
 - 1. Deploy to the non-active environment
 - 2. Switch the traffic coming through the load balancer over to the new environment
 - 3. Monitor new version, maybe switch back
- One environment always runs the last version, the other the newest version: easy roll-backs

Deployment Strategies

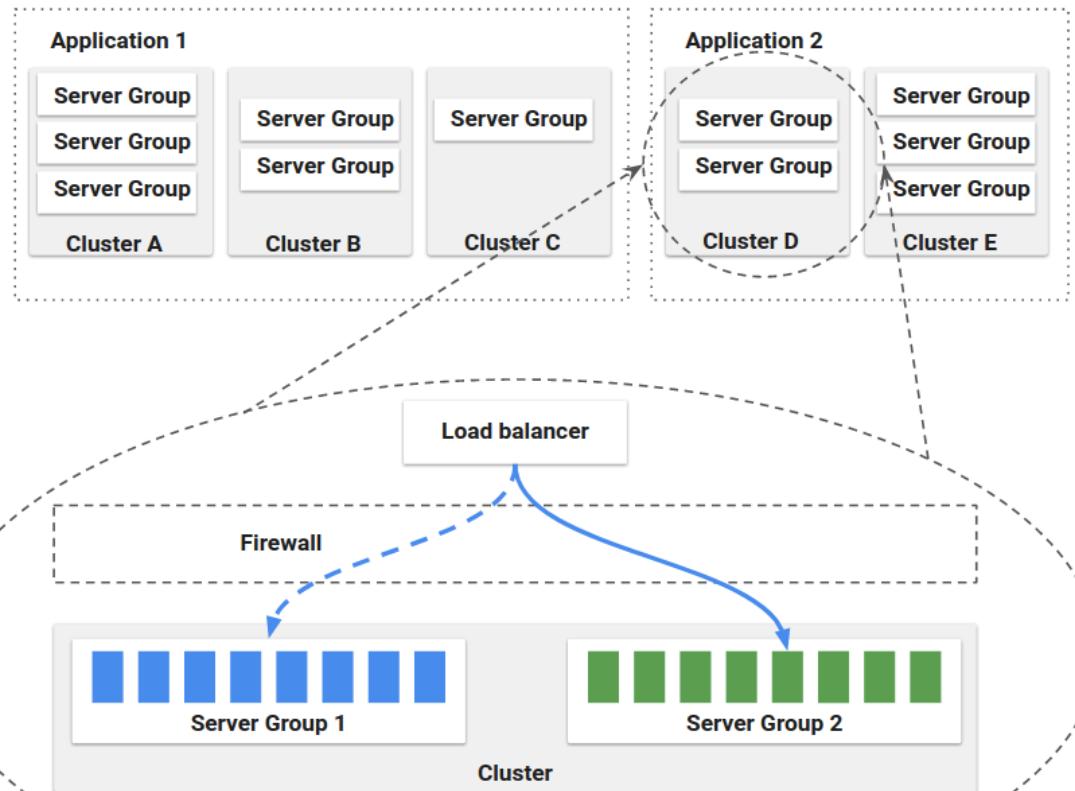


Spinnaker

- Continuous delivery platform
- Started by Netflix, collaboration with Google
 - Successor to Asgard, a web interface to manage AWS
 - From baking to cloud deployment
- Support and common abstractions for EC2, GCE, Kubernetes, Azure, ...

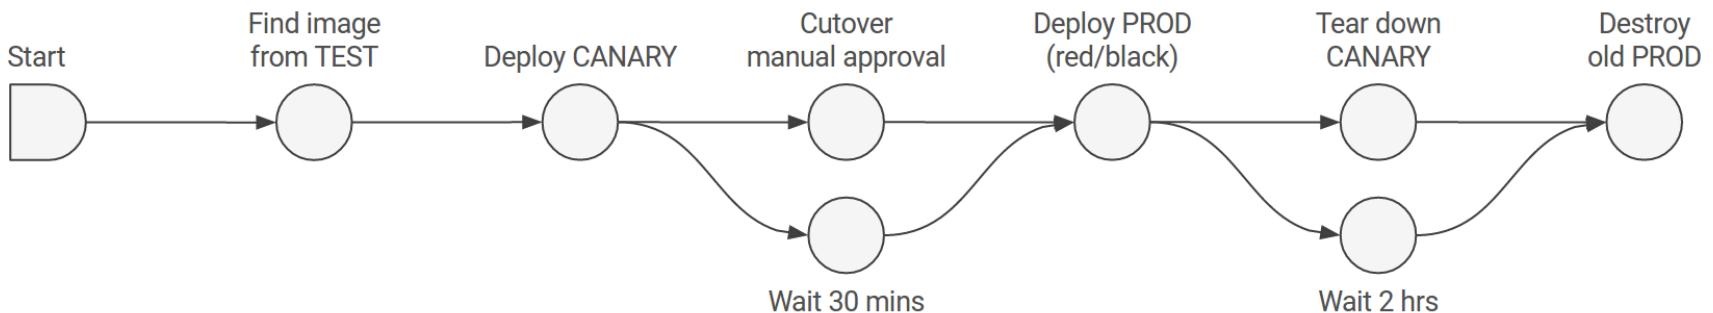
Spinnaker

- Deploy applications / (micro)services
- Organize servers into
 - Clusters
 - Server groups

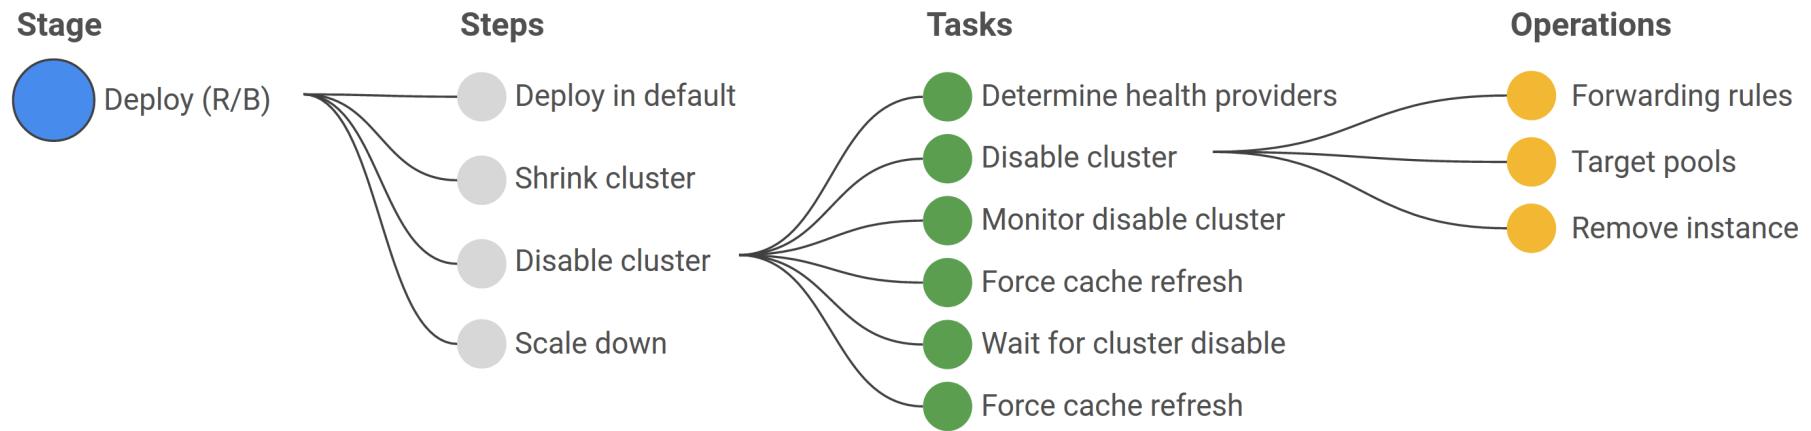


Spinnaker Pipeline

- Deployment happens via pipelines
- Pipelines are built from stages
- Predefined stages exist for many things
 - Bake
 - Manual Judgement
 - Clone server group
 - Deploy



Spinnaker Stage Example



GitLab

- Includes a CI/CD system
- CI Jobs are defined in .gitlab-ci.yml

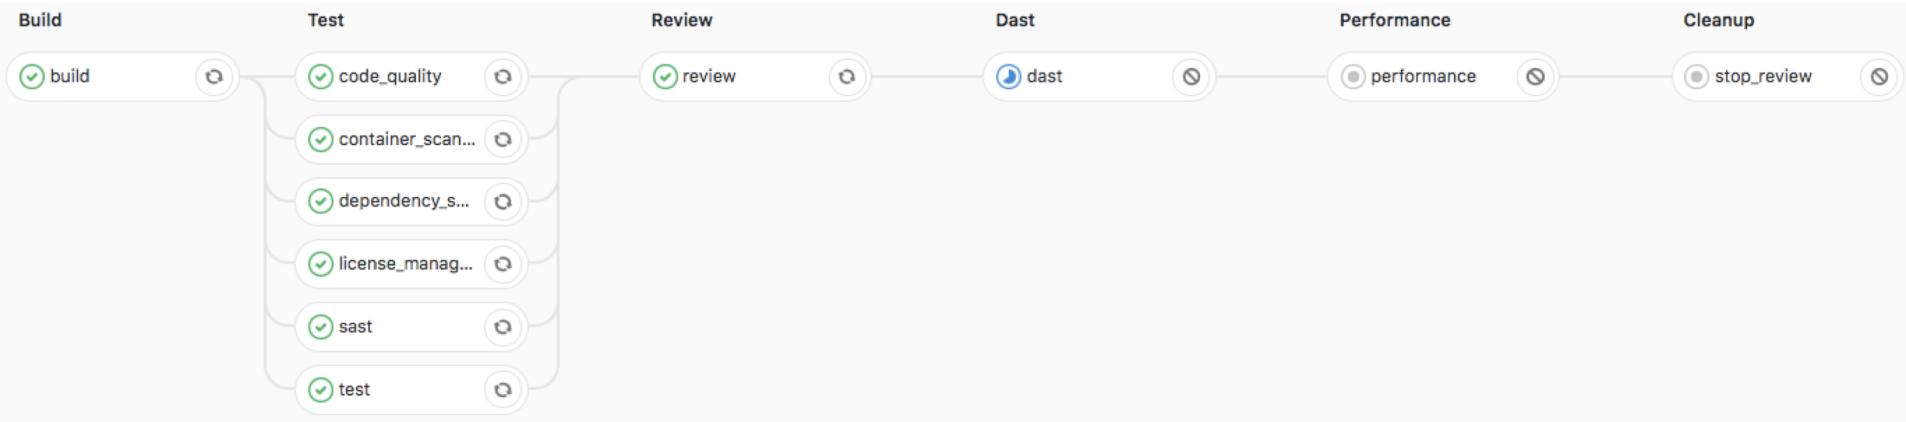
```
#gitlab-ci.yml

image: blang/latex
Job Name build:
    script: Commands for this job; required property
            - latexmk -pdf
only: Only execute on certain branches
      - master
artifacts:
    paths: Save these artifacts back to server
          - "*.*pdf"
```

GitLab: Auto DevOps

- Based on Docker, Kubernetes and Gliderlabs Herokuish
- Defined as one long .gitlab-ci.yml
- Builds based on
 - Dockerfile or
 - Herokuish buildpacks
- Tests
 - Code Quality
 - Dependencies
 - Static and dynamic application security testing
 - Test via buildpacks
- Deploy to production or new review app

GitLab: Auto DevOps, Example



Pipeline for Branches

GitLab: Auto DevOps

- Integrated with the merge request flow

The screenshot shows a GitLab merge request interface. At the top, it displays a 'Request to merge' from the branch 'beilharz-master-pat...' into the 'master' branch. Below this, a green checkmark indicates that 'Pipeline #37212462 passed for 7a62a8a4 on beilharz-master-pat...'. A message below the pipeline status states 'Deployed to review/beilharz-master-... 6 hours ago' and 'Memory usage decreased from 105.35MB to 104.44MB'. To the right of this message is a 'View app' button, which is highlighted with an orange rectangle. The next section shows a warning icon with the text 'Requires 1 more approval'. Following this are several green checkmarks indicating 'No changes to code quality' and 'No changes to performance metrics', both of which are also highlighted with orange rectangles. Below these, a warning icon indicates 'Security scanning detected 2 new vulnerabilities', with a 'View full report' button highlighted with an orange rectangle. The final section shows a green checkmark for 'License management detected no new licenses' with a 'Manage licenses' button and a 'View full report' button. At the bottom, a warning icon prevents the merge with the message 'You can only merge once the items above are resolved'.

Overview

- Intro
- Cloud Operating Systems
 - OpenStack
- Infrastructure-as-Code
 - Ansible
- Container Orchestration
 - Kubernetes
- DevOps, Continuous Integration, and Continuous Delivery

Summary

- Virtual machines and containers allow to quickly provision new servers, i.e. by using the interface of a Cloud OS
- Yet, server management and configuration remain very challenging (server sprawl, configuration drift, snowflake servers)
- Central Ideas
 - Infrastructure-as-Code: executable, understandable, and versioned descriptions of desired states as input to automation and orchestration tools
 - Cultural change to DevOps, for continuously and automatically deploying new versions to production

References

- Humble, Jez, and David Farley. *Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation*. Addison Wesley, 2010.
- <https://martinfowler.com/bliki/SnowflakeServer.html>
- <https://martinfowler.com/bliki/BlueGreenDeployment.html>
- <https://www.spinnaker.io/concepts/>
- <https://www.spinnaker.io/reference/architecture/>
- <https://docs.gitlab.com/ee/ci/>
- <https://docs.gitlab.com/ee/topics/autodevops/index.html>