

Methods of Cloud Computing

Programming Cloud Resources 1: Scalable and Fault-Tolerant Applications



Complex and Distributed Systems
Faculty IV
Technische Universität Berlin



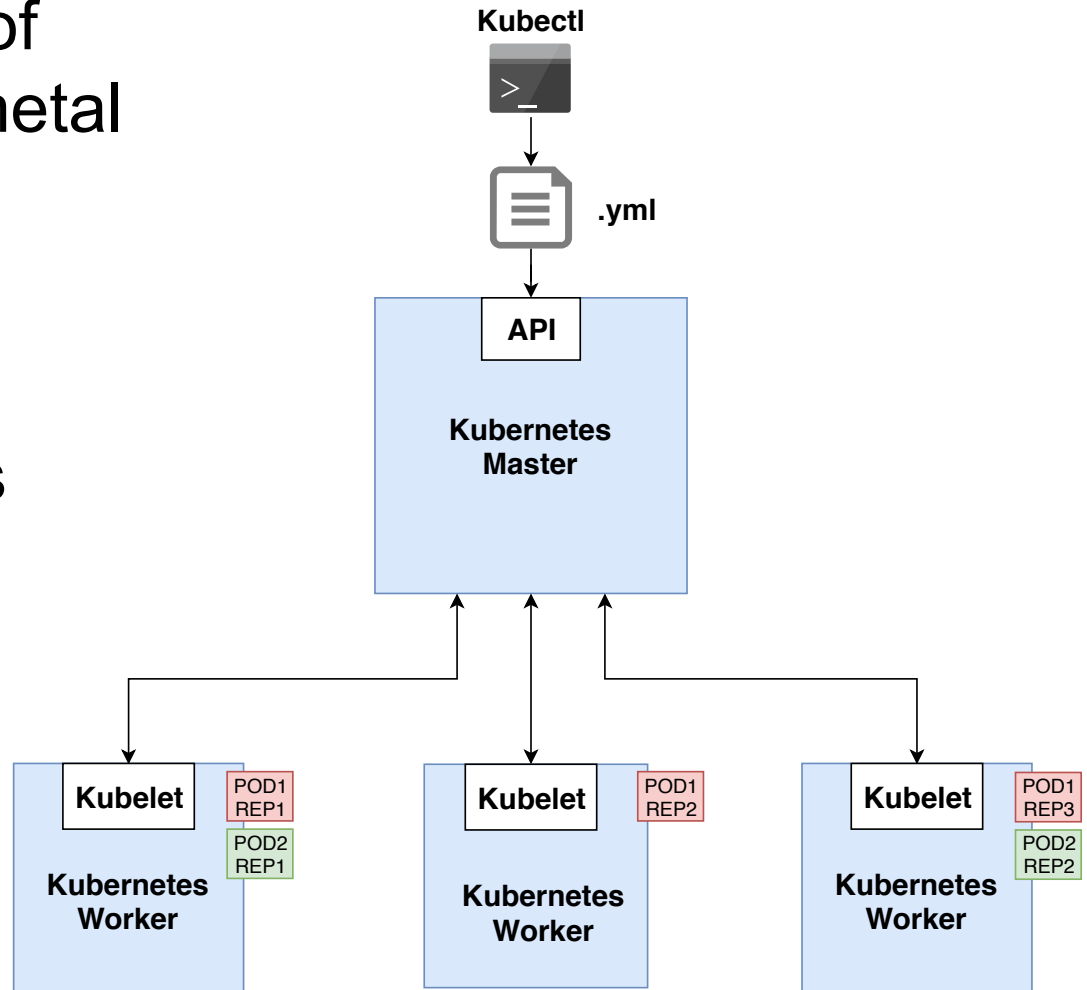
Operating Systems and Middleware
Hasso-Plattner-Institut
Universität Potsdam

Overview

- Intro
- Partitioning
- Replication and Consistency
- CAP Theorem
- Case Studies
 - **Kubernetes Auto-Scaling**
 - Amazon DynamoDB
 - Blockchain

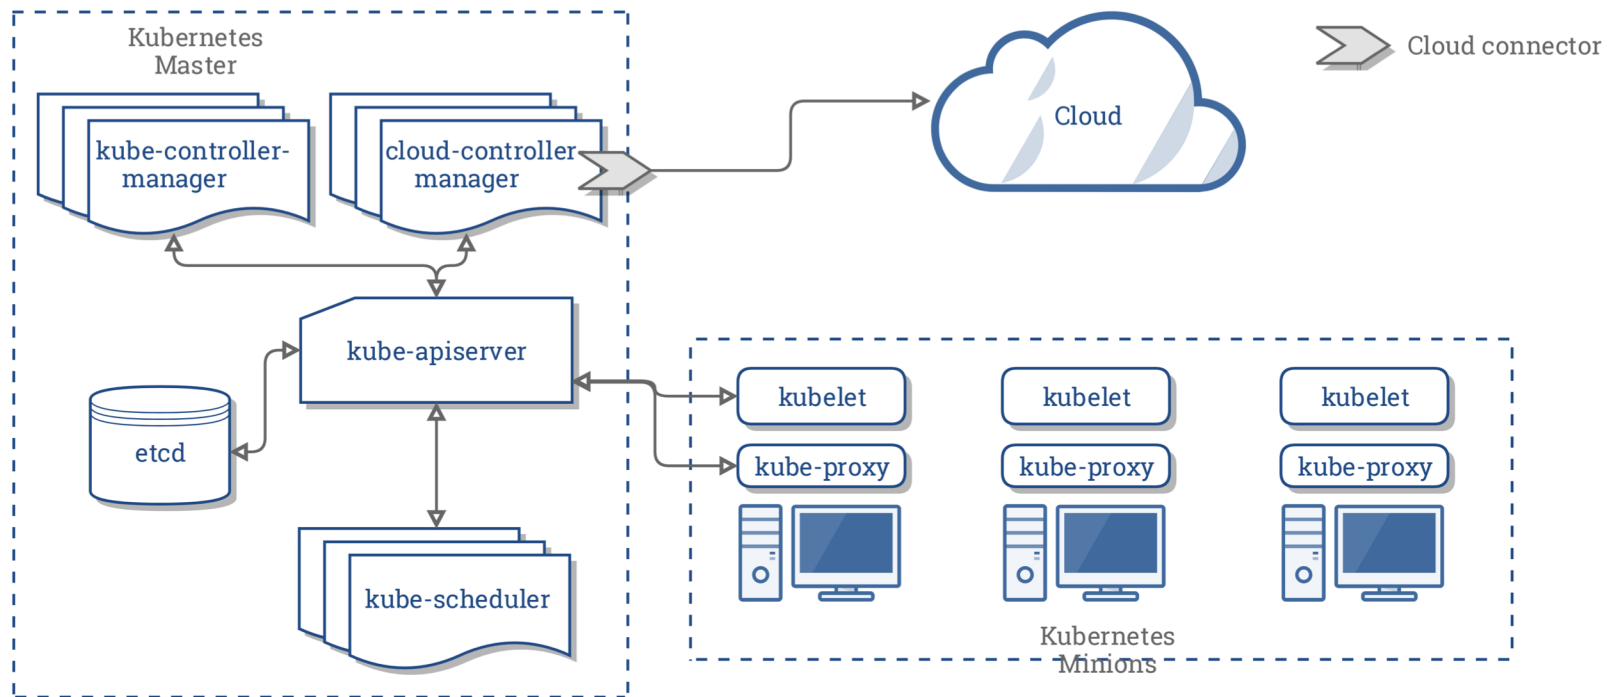
Recap: Kubernetes Cluster

- Manages collection of nodes (either bare-metal or virtual machines)
- Runs groups of replicated containers (called *Pods*)



Kubernetes

Horizontal Pod Autoscaler




Kubernetes

Horizontal Pod Autoscaler

- Automatically scales number of pods in a replication controller
- Based on
 - CPU utilization or custom metrics
 - Target value for the metric
- Autoscaler checks metrics every 30s
- Sets the number of replications to optimize the metric towards the target value
 - Creating more pods
 - Or reducing the number of pods

Kubernetes

Horizontal Pod Autoscaler



Kubernetes

Horizontal Pod Autoscaler

- Number of replicas is scaled by the quota of average current metric value and target value

```
desiredReplicas =  
    [ currentReplicas * (currentMetricValue/desiredMetricValue) ]
```

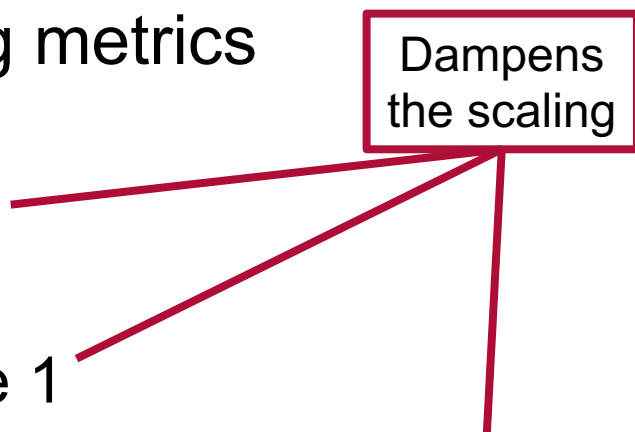
- No scaling if this quota is within **0.1 tolerance**

Jitter avoidance using
tolerance

➔ This assumes linear scaling!

- Autoscaler scales to the highest number of desired replicas in a 5 minute sliding window
 - Quick responses to more load, but reduces *thrashing*

Sometimes metrics are not available

- Discard pods that are being shut down
 - Normally running pods with missing metrics
 - If result without these would be to scale up: assume metric to be 0
 - If result without these would be to scale down: assume metric to be 1
 - Assume metric to be 0 for pods that are not yet ready
- 
- Dampens the scaling

Overview

- Intro
- Partitioning
- Replication and Consistency
- CAP Theorem
- Case Studies
 - Kubernetes Auto-Scaling
 - **Amazon DynamoDB**
 - Blockchain

Amazon DynamoDB^[4]

- Highly-available key-value store, provided as a managed service by Amazon
- Primary design goals
 - High scalability
 - ◆ E.g. 1000s of servers
 - High availability
 - ◆ Particularly, support for “always write”
 - High performance
 - ◆ Particularly, small latency (single digit milliseconds) and number of requests (20 million requests per second)
- Sacrifices consistency to achieve these goals



Amazon DynamoDB: Design Principles


- Dynamo follows **peer-to-peer** approach
 - No server is more important than any other
 - No single point of failure
- Nodes can be added/removed incrementally at runtime
 - In terms of the CAP theorem, DynamoDB is **AP**
 - Weak consistency guarantee: **eventual consistency**
- No hostile environment, all servers obey rules
 - No security issues must be considered

How to Partition Data Among the Servers?

- Dynamo designed to be a key-value store
 - No support for range queries needed
 - No need for range partitioning
- Hash partitioning has good load balancing properties
 - But how to avoid data reorganization when servers are added or removed?
Consistent hashing: we assign virtual servers in the ring and only had to move data for the server lost instead of all the servers
- Solution: Hash function no longer maps to partitions
 - Instead function maps to a fixed number of slots
 - Variation of classic hashing called *consistent hashing*


Consistent Hashing_[6]

- Idea: Additional mapping between slots and partitions
 - Hash function maps keys to large but fixed number of slots (for example 2^{128} slots)
 - A partition now covers a consecutive number of slots
 - ◆ A server can be in charge of one or more partitions
 - ◆ Size of a partition is variable




Distributed Hash Tables

- Consistent Hashing is basis for distributed hash tables (DHTs)
 - Each server takes at least one partition
 - Therefore each server is responsible for a continuous range of slots
 - Upon arrival/departure of server only $O(\#Keys/\#Servers)$ data items must be reorganized




Amazon DynamoDB's Partitioning Algorithm (1/2)

- Slots form a circular ID space
 - All servers hash their ID with an MD5 function
 - Hashing result determines server's position on the ring




Amazon DynamoDB's Partitioning Algorithm (2/2)

- To distribute data, key of data also hashed with MD5
 - Result of hashing is a position in the circular ID space
 - Rule: Server is responsible for all preceding IDs (i.e. slots) up to and including its own ID



Routing in Amazon DynamoDB

- Each server maintains full routing table (ID to IP)
 - Each server can determine which server is responsible for a data item based on routing table and mapping rule!
 - One hop routing keeps latencies small




Example routing table in DynamoDB

ID	IP Address
5	192.168.1.2
14	192.168.1.18
17	192.168.1.115
27	192.168.1.98

Virtual Servers for Load Balancing

- Despite uniform distribution over ID space, the servers may receive skewed number of requests
- Idea: Each server appears on multiple positions on the ring (known as virtual servers)



Replication in Amazon DynamoDB

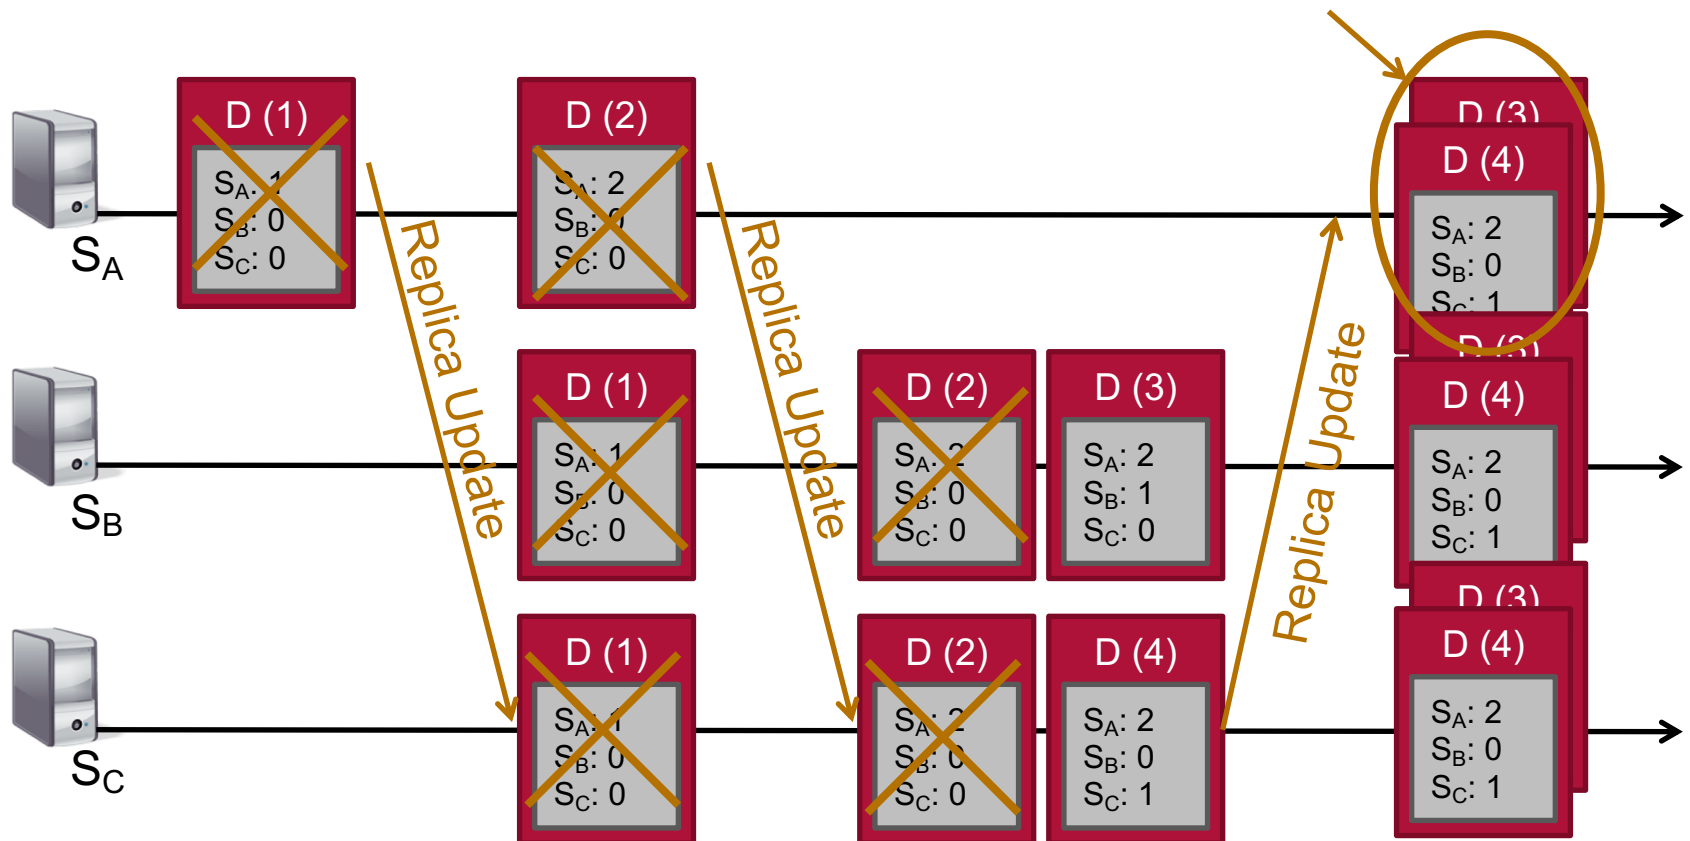
- Servers replicate data to their N successors on the ring
- Replication is adjusted whenever a server is added or removed from the ring
- Servers use heart beat protocol to determine availability
 - Periodic messages exchanged between ring neighbors
 - When server does not answer heart beat request in a given time span, it is considered gone
- Dynamo allows reads and writes on every replica!

Data Versioning in Amazon DynamoDB

- Dynamo allows „always write“ paradigm
 - Write operation allowed on every replica
 - put-() operation returns after one replica has been written
 - ◆ Lazy update of replicas in the background
- Result: Different version of a data item may exist
 - Dynamo treats each version of the data item as an immutable object
 - Vector clocks are used to reconcile different versions
 - When system cannot reconcile different versions, the versions are presented to client for reconciliation

Example of Version Reconciliation

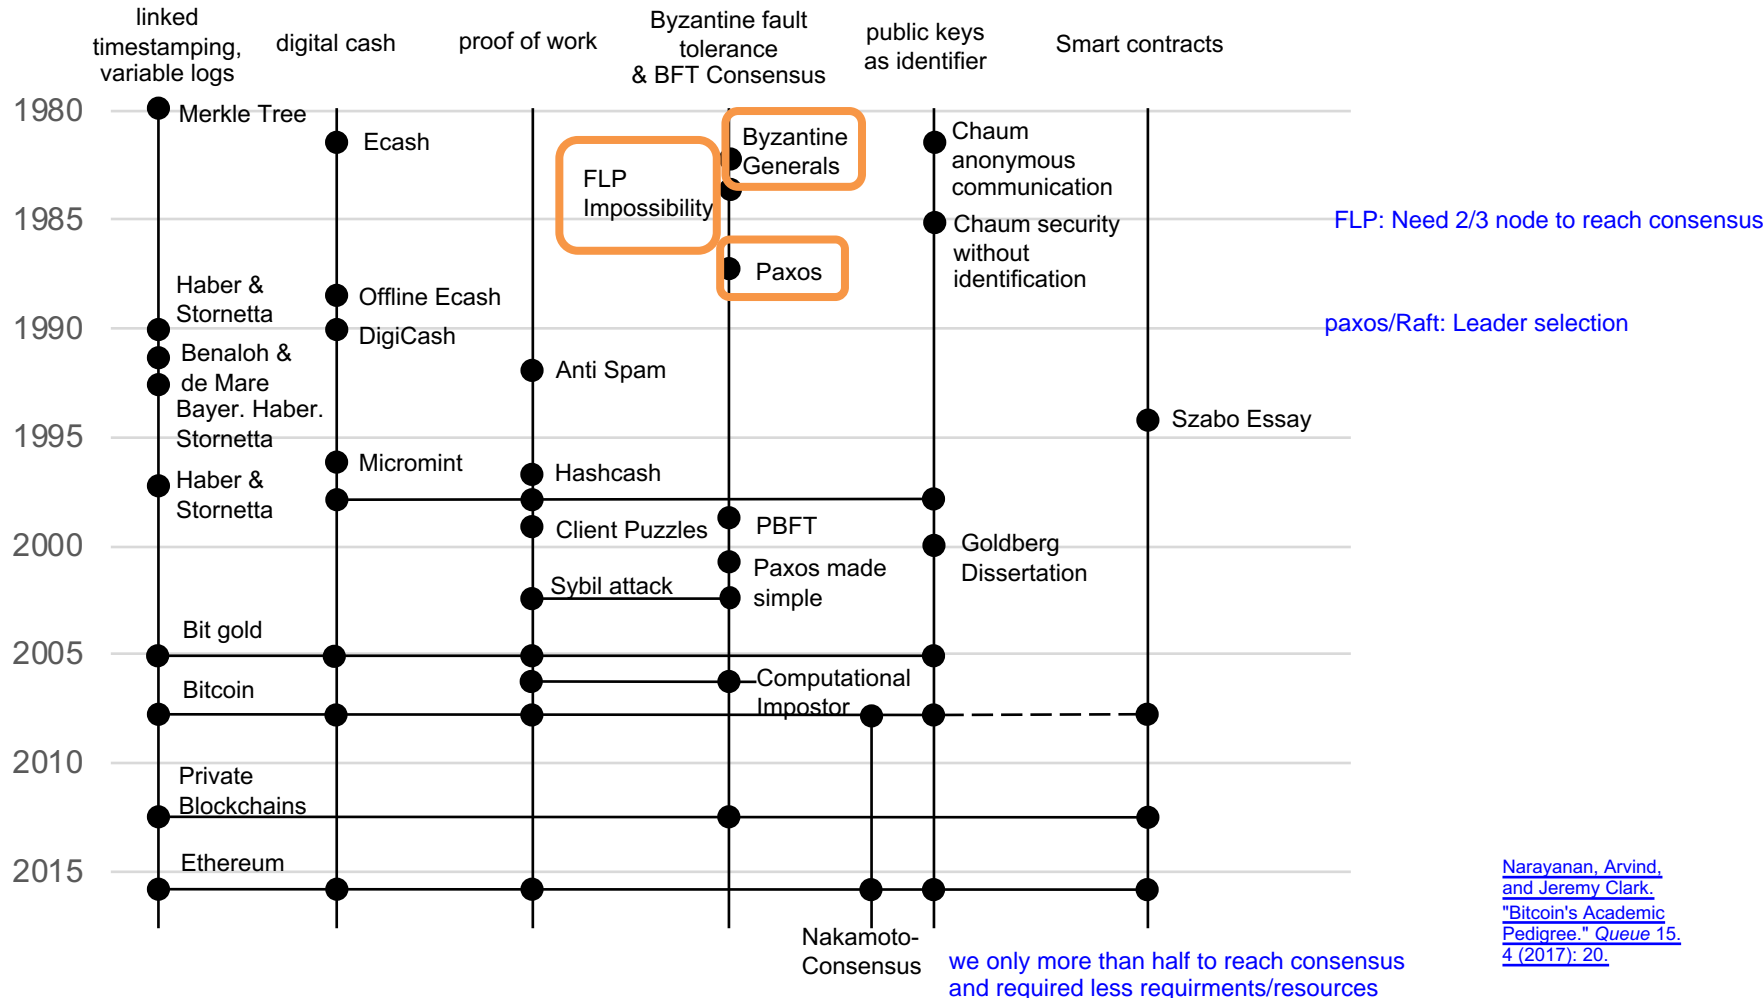
- Let's assume data item D is replicated among three servers: S_A , S_B , and S_C



Overview

- Intro
- Partitioning
- Replication and Consistency
- CAP Theorem
- Case Studies
 - Kubernetes Auto-Scaling
 - Amazon Dynamo Load balancing
 - **Blockchain** consistency

Blockchain: Background




A Ledger

Amount	Sender	Receiver
...
2 BTC	2bf12	4c2dd
2 BTC	4c2dd	1156f
1 BTC	4c2dd	2bf12
...

1. Transactions are signed by the sender
2. Nobody is allowed to send more money than he has
 - we are done?

Blockchain




Nakamoto-Consensus

Nakamoto-Consensus based on Proof-of-Work and fixed consensus rules

- Proof-of-Work:


- For a given target:
- Find a *nonce*, so that



$$\text{hash}(\text{hash}(\text{block}_{n-1}) \oplus \text{transactions} \oplus \text{nonce}) < \text{target}$$

➔ whoever is the first to find an appropriate nonce gets to propose the next block

Double Spends



once a week happens for block chain

stores want to find chain with 6 blocks since it's very highly unlikely to happen with 6 blocks

As as long 50% server are behaving we are good comparing to byzantain where we need 2/3

Nakamoto-Consensus

Voting is not explicit but implicit by signing blocks

→ voting weight proportional to computer performance

Overview

- Intro
- Partitioning
- Replication and Consistency
- CAP Theorem
- Case Studies
 - Kubernetes Auto-Scaling
 - Amazon Dynamo
 - Blockchain

Summary

- Scaling-out to more virtual resources: scalability and fault tolerance
 - Load balancing for replicated stateless components
 - Load balancing *and* data consistency models for replicated stateful components
- The higher the consistency level, the less scalable replicated services are
- System components fail eventually, decide on either availability or consistency

Literature and References

- Literature
 - A.S. Tannenbaum, M. Van Steen: “Distributed Systems: Principles and Paradigms”, Prentice Hall, 2016, Chapter 7
 - M. Kleppmann, “Designing Data-Intensive Applications”, 2017, Chapter 5 and 6
- References
 - [2] Eric. A. Brewer: “Towards Robust Distributed Systems”, PODC Keynote 2004, <http://www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf>
 - [3] S. Gilbert, N. Lynch: “Brewer’s Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services”, ACM SIGACT News, 33 (2), 2002
 - [4] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, W. Vogels: “Dynamo: Amazon’s Highly Available Key-Value Store”, in Proc. of the 21st ACM SIGOPS Symposium on Operating Systems Principles, 2007
 - [6] D. Karger, E. Lehman, T. Leighton, R. Panigrahy, M. Levine, D. Lewin: “Consistent Hashing and Random Trees: Distributed Caching Protocols for Relieving Hot Spots on the World Wide Web”, in Proc. of the 29th ACM Symposium on Theory of Computing, 1997