

Methods of Cloud Computing

Beyond Cloud



Complex and Distributed Systems
Faculty IV
Technische Universität Berlin



Operating Systems and Middleware
Hasso-Plattner-Institut
Universität Potsdam

Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Overview

- **Intro**
- **Federated, Hybrid and Multi-Clouds**
 - Concepts
 - Policy enforcement
- **Fog, Edge and IoT**
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Background

Cluster

- Many computers in one room

Grid

- Loosely coupled computers or clusters all over the world

Cloud

- IT resources as a utility

Edge

- Extending the Cloud to the edge of the network

Cluster

- Mostly homogeneous compute resources and software stacks
- Interconnected by a low-latency and high-bandwidth network
- Goal: improving availability and price/performance
- Examples
 - Analytics cluster at Facebook

Grid

- Heterogeneous compute resources
 - Connected via a slow network (i.e. the internet)
 - Heterogeneous software stacks unified by a middleware
-
- Examples
 - PanDA (CERN)
 - BOINC

Background

Cluster

- Many computers in one room

Grid

- Loosely coupled computers or clusters all over the world

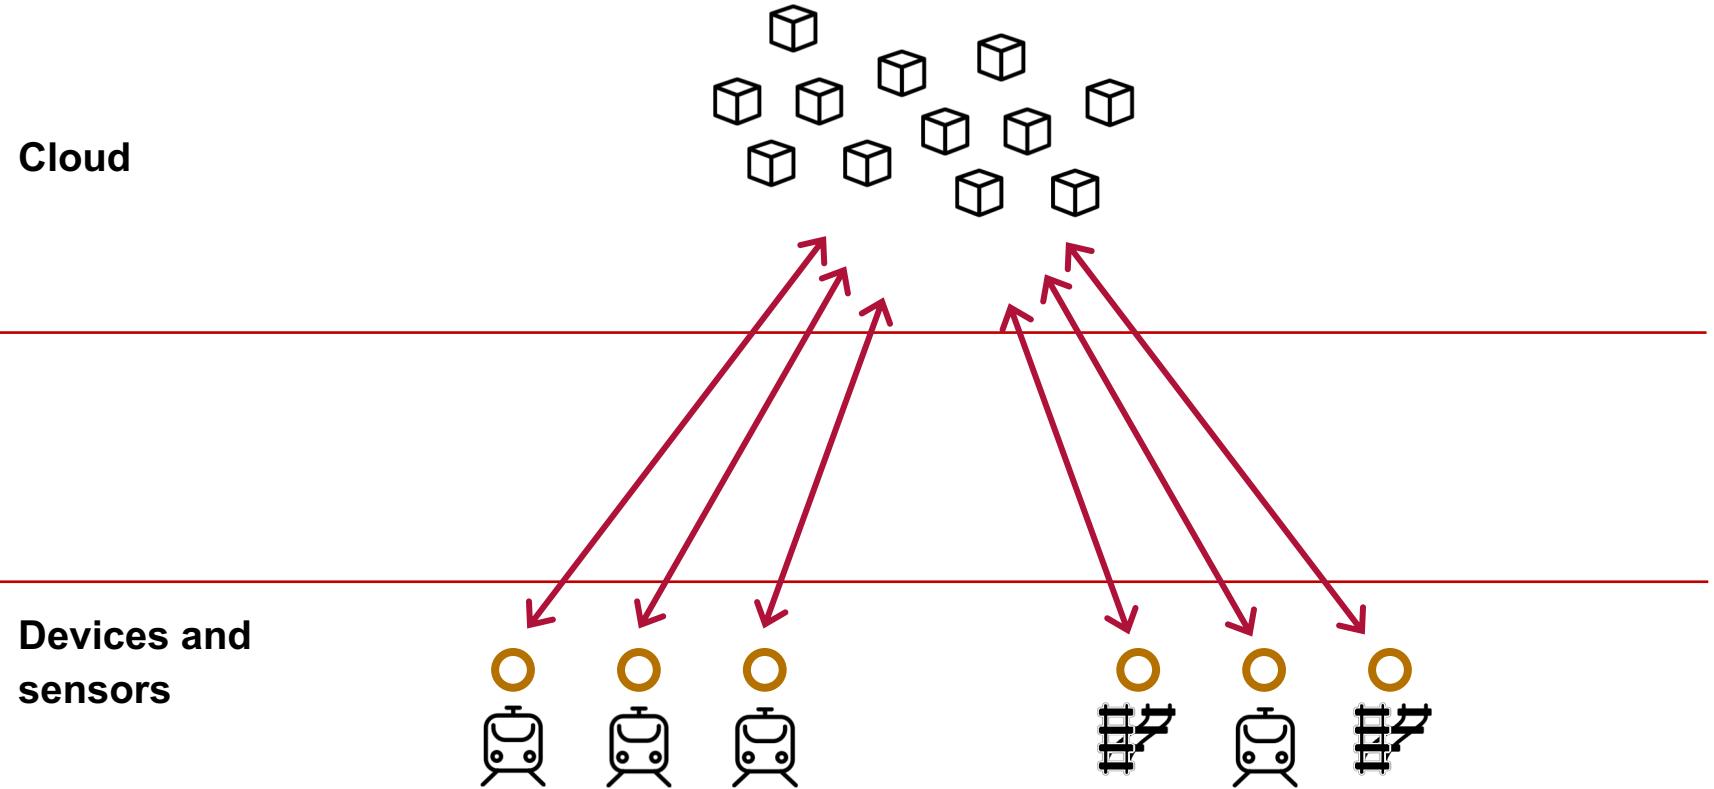
Cloud

- IT resources as a utility

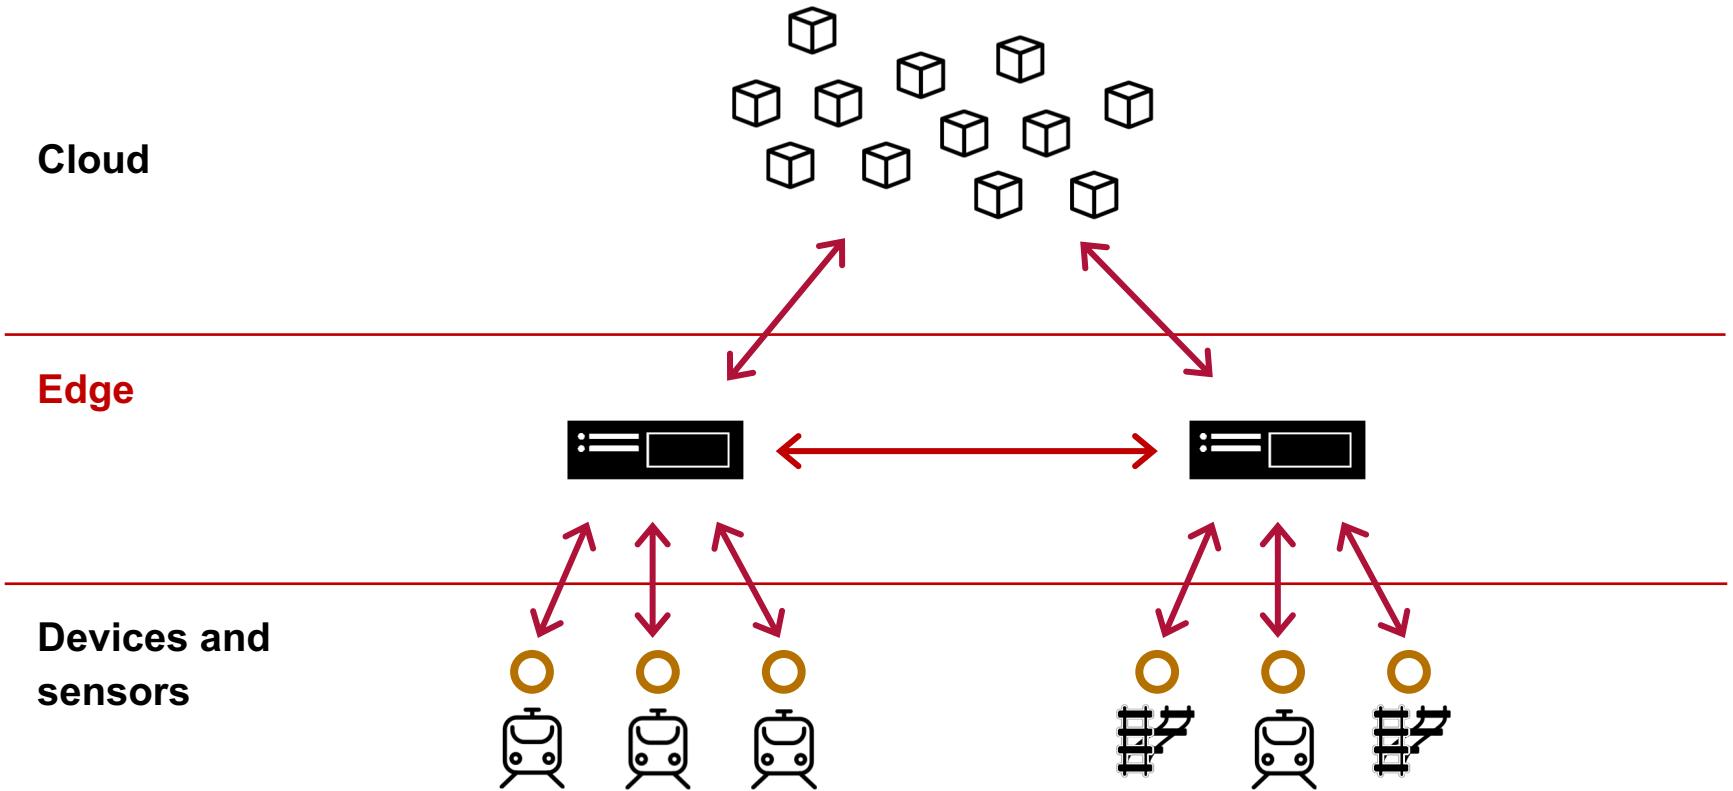
Edge

- Extending the Cloud to the edge of the network

Edge Computing



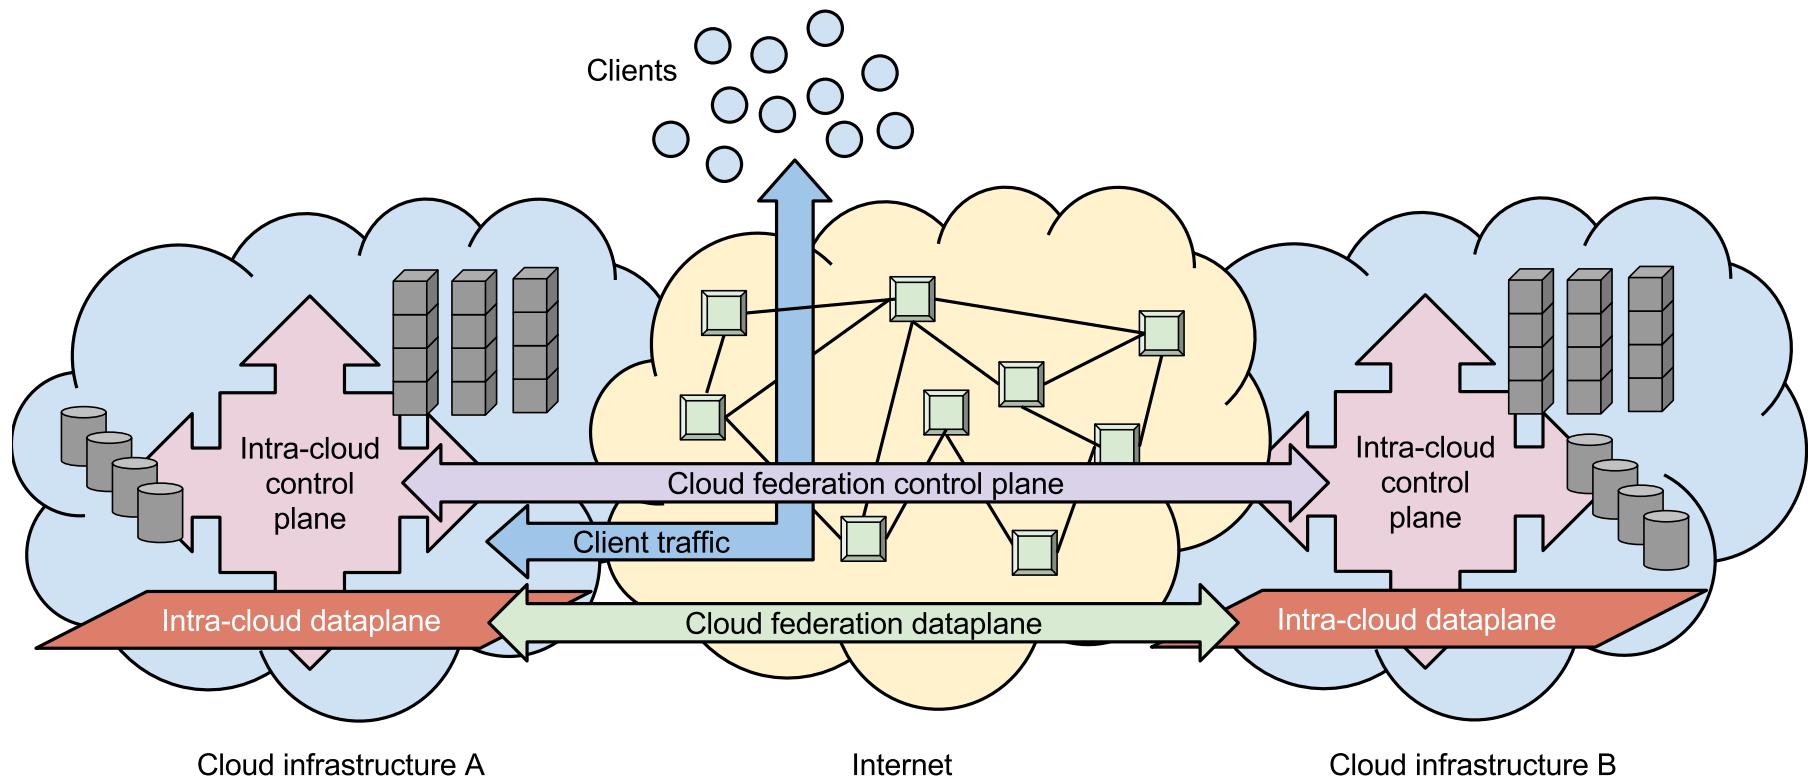
Edge Computing



Overview

- Intro
- **Federated, Hybrid and Multi-Clouds**
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Cloud Federation



Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

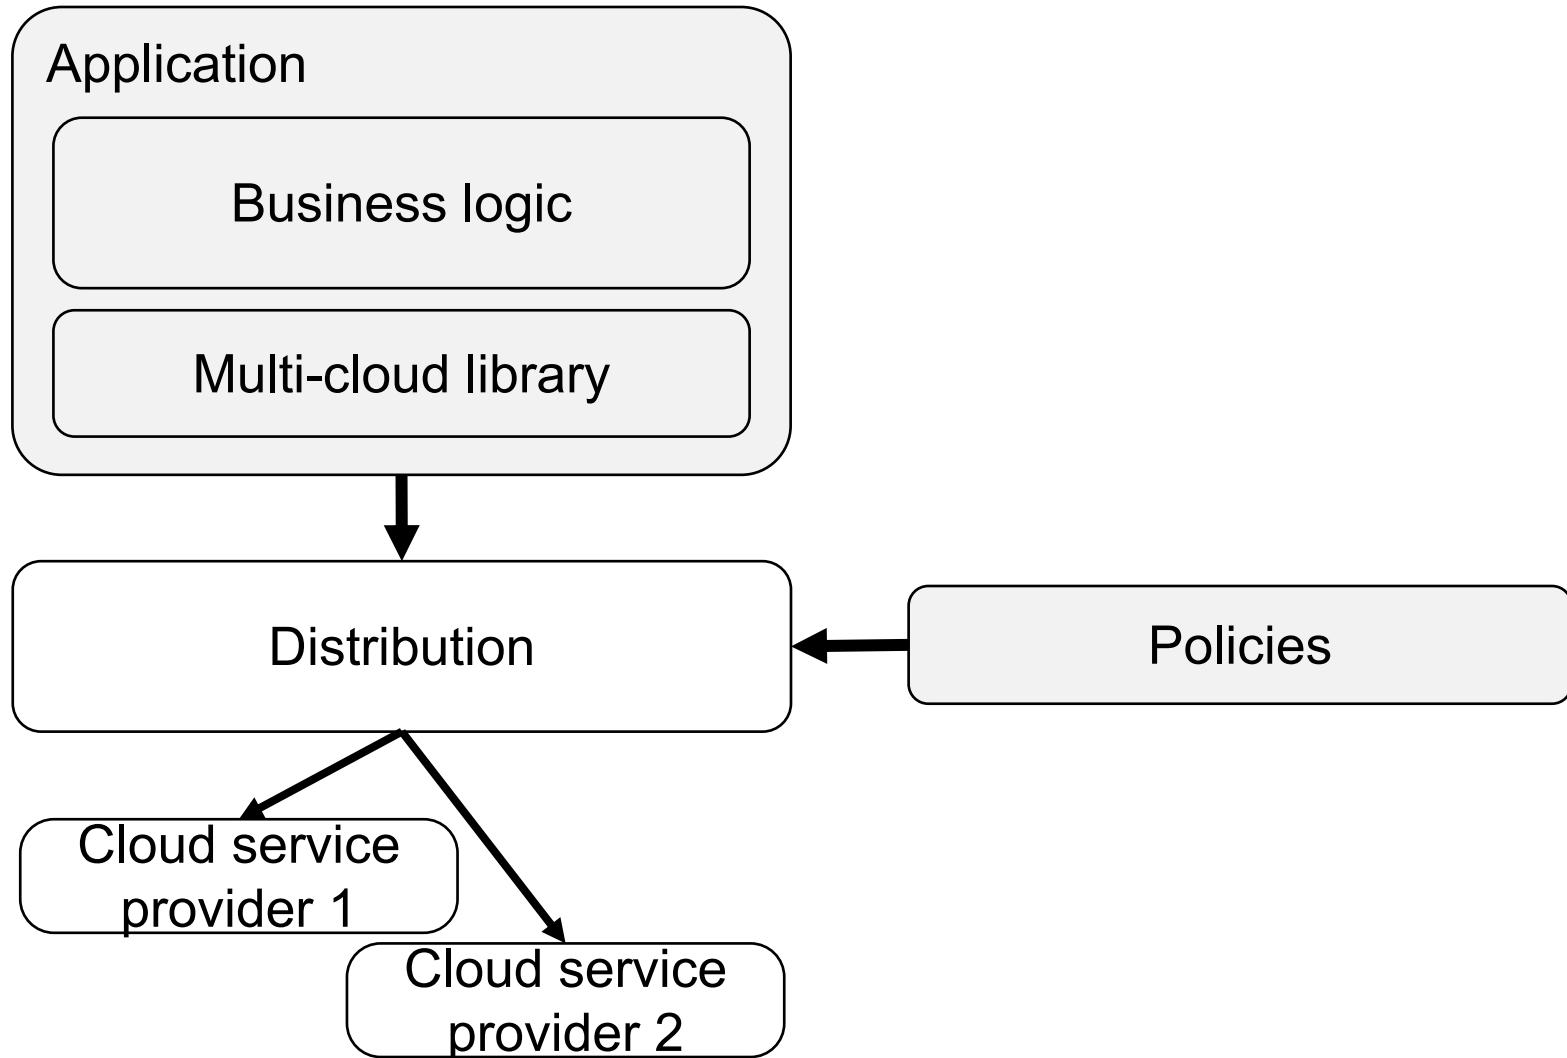
Concepts

- Combining providers and models
- Multi-cloud:
 - Using multiple cloud providers in one application
- Federated Cloud:
 - Integrating multiple cloud providers to provide a service
- Hybrid Cloud:
 - Mixing public and private cloud models

Motivation for Cloud Federation

- Resilience against vendor outages
- Fighting vendor lock-in
- Harness the diversity of cloud services by different providers
- Harness different locations of providers
- Compliance with policies
 - i.e. data needs to be stored in EU
→ Needs a way to express and enforce the policies

Cloud Federation



Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - **Policy enforcement**
- Fog, Edge and IoT
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Customer Requirements

CPPL Policy Language



- Location
- Time of deletion
- Notification on access
- Protection
- Replication
- ...

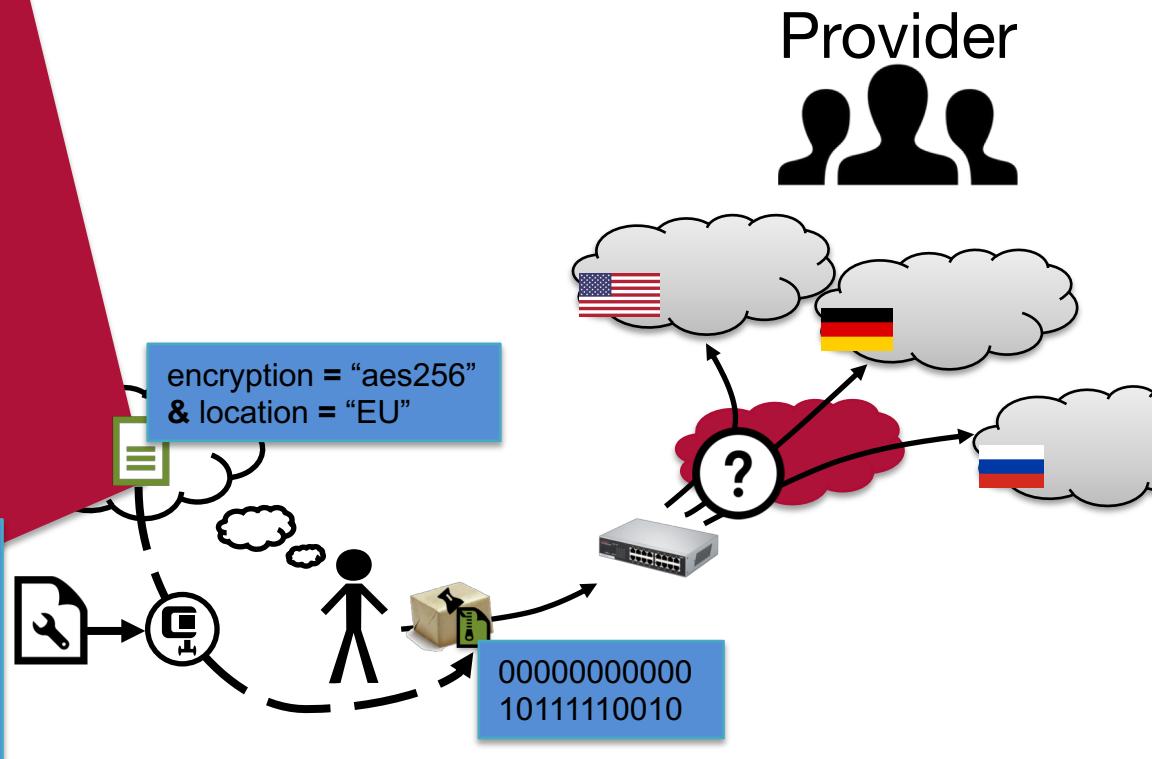


- Location
- Type
- ...



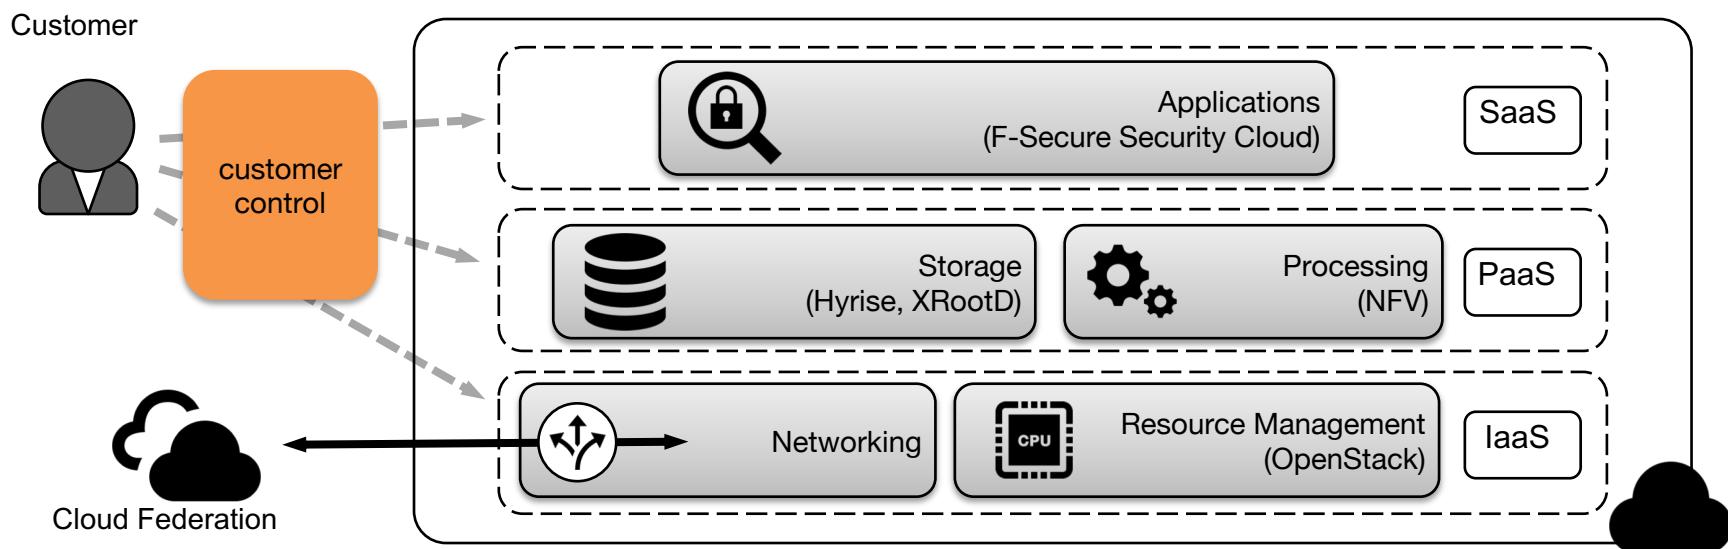
- Restricted paths
- Transport protection
- ...

```
{  
  "name": "encryption", "type": "string",  
  "values": ["aes192", "aes256"]  
},  
{  
  "name": "location", "type": "string",  
  "values": ["US", "DE", "RU", "EU"]  
}
```



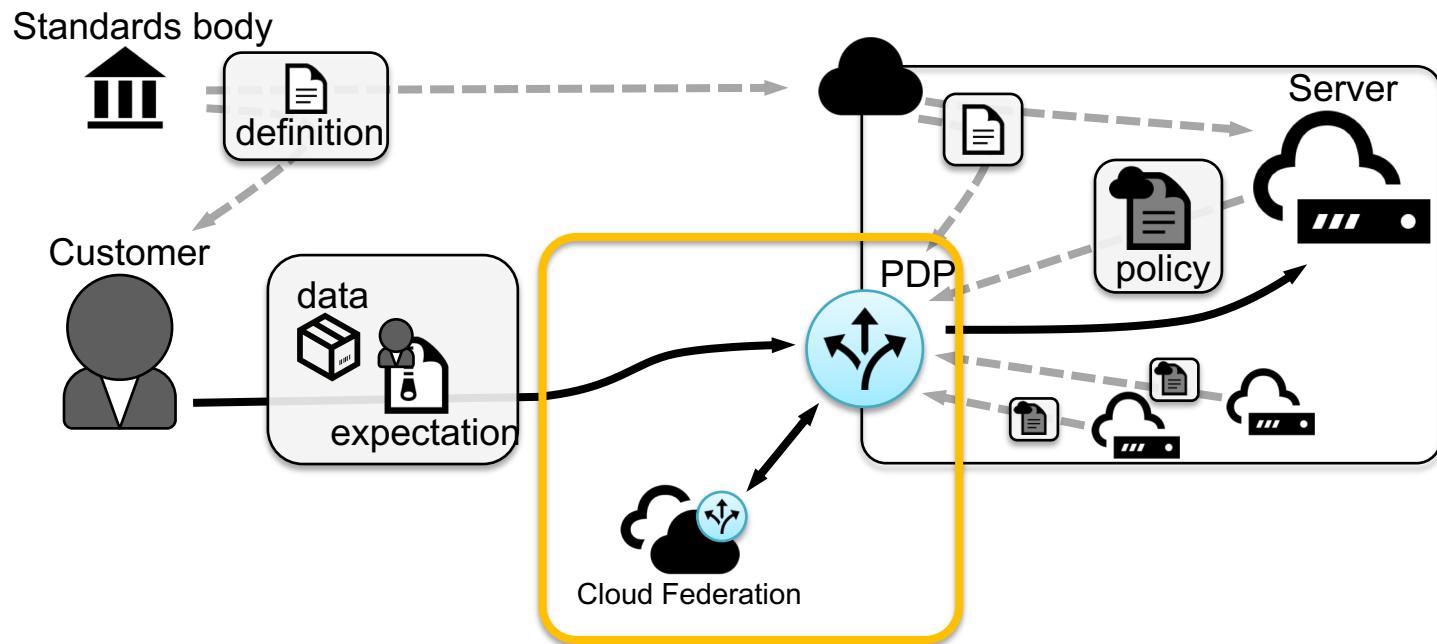
Policy-Aware Cloud Architecture

- Policy-awareness must be addressed at all cloud layers



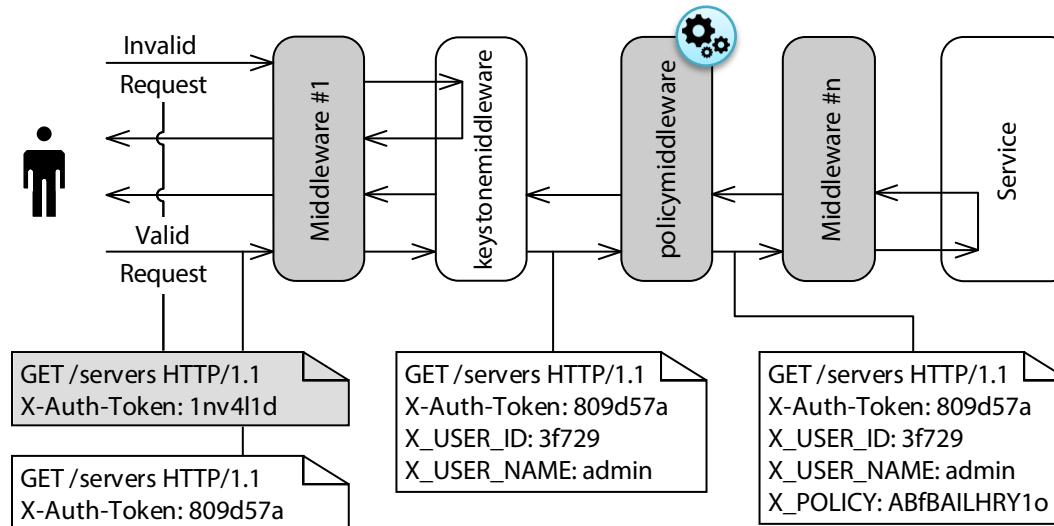
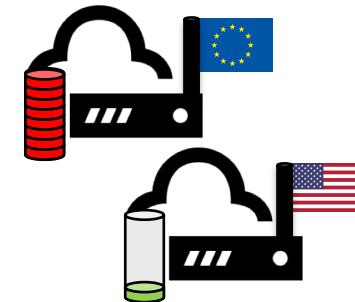
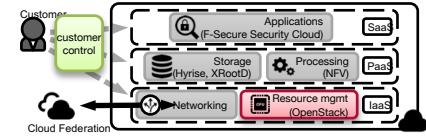
Decisions at the Edge

- Policies often restrict set of eligible servers
- Policy decision points:
 - Check policies and assign data to complying systems
 - Realize policy compliance across clouds of a federation



OpenStack – Policy-Aware Resource Management

- Elasticity in a policy-aware cloud
 - Often requested but seldom provided attributes → High load
 - Traditional elasticity is **not** policy-aware
- Policy-aware resource management for **OpenStack**



OpenStack – Policy-Aware Resource Management

The screenshot shows the OpenStack dashboard interface. The top navigation bar includes the 'openstack' logo, a 'demo' dropdown, and a user dropdown for 'demo'. The left sidebar has a 'Project' dropdown, an 'Identity' dropdown, and a 'Settings' section with 'User Settings' and 'Change Password' options. A red vertical bar highlights the 'Edit Policy' link under the 'Settings' section. The main content area is titled 'Edit Policy' and contains a 'Policy' section with the following JSON code:

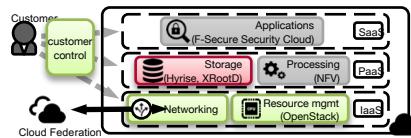
```
{  
    "availability_zones": [  
        "az2"  
    ],  
    "storage": {  
        "encryption": true  
    }  
}
```

Below the policy is a 'Description:' field with the placeholder 'Edit your policy.' and a 'Save' button at the bottom right.

On the far right, a vertical menu is open, showing 'Settings' (selected), 'Help', 'Themes' (with 'Default' checked), 'Material', and 'Sign Out'.

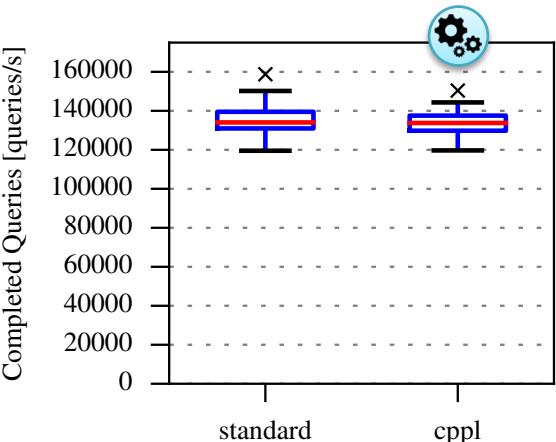
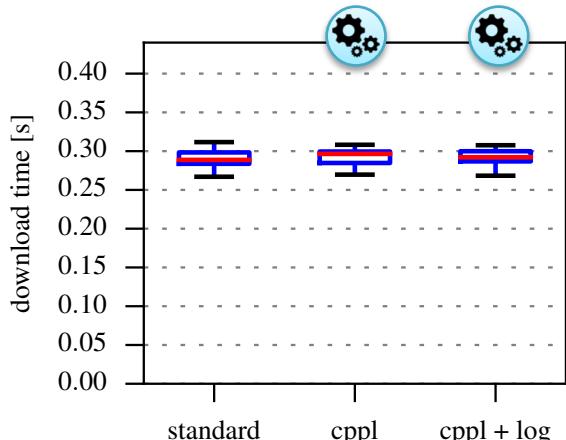
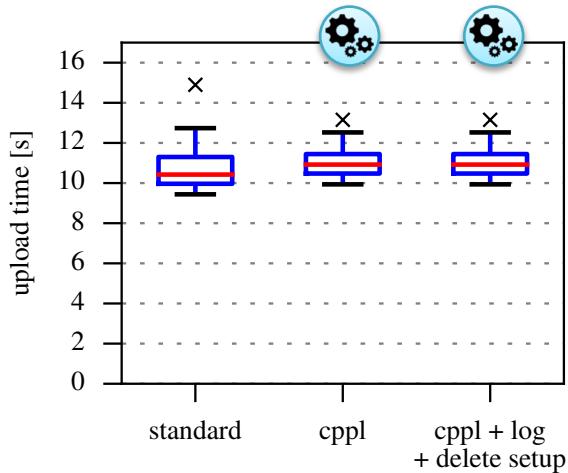
OpenStack – Policy-Aware Resource Management

The screenshot shows the OpenStack Compute (Nova) dashboard. The left sidebar is collapsed, showing 'Project' and 'Compute' dropdowns, and three main menu items: 'Overview', 'Instances', and 'Volumes'. The 'Volumes' item is highlighted with a red vertical bar. The main content area is titled 'Volumes' and contains three tabs: 'Volumes' (selected), 'Volume Snapshots', and 'Volume Consistency Groups'. Below the tabs is a search/filter bar with a 'Filter' input field, a magnifying glass icon, and two buttons: '+ Create Volume' and 'Accept Transfer'. A table header follows, listing columns: Name, Description, Size, Status, Type, Attached To, Availability Zone, Bootable, Encrypted, and Actions. A message box at the top right displays an error: 'Error: Your policy requires using an encrypted volume type.' with a close button.



Policy-Aware Storage

- Realized in **XRootD** (CERN) and **Hyrise**
 - Big data
 - High transactional workloads

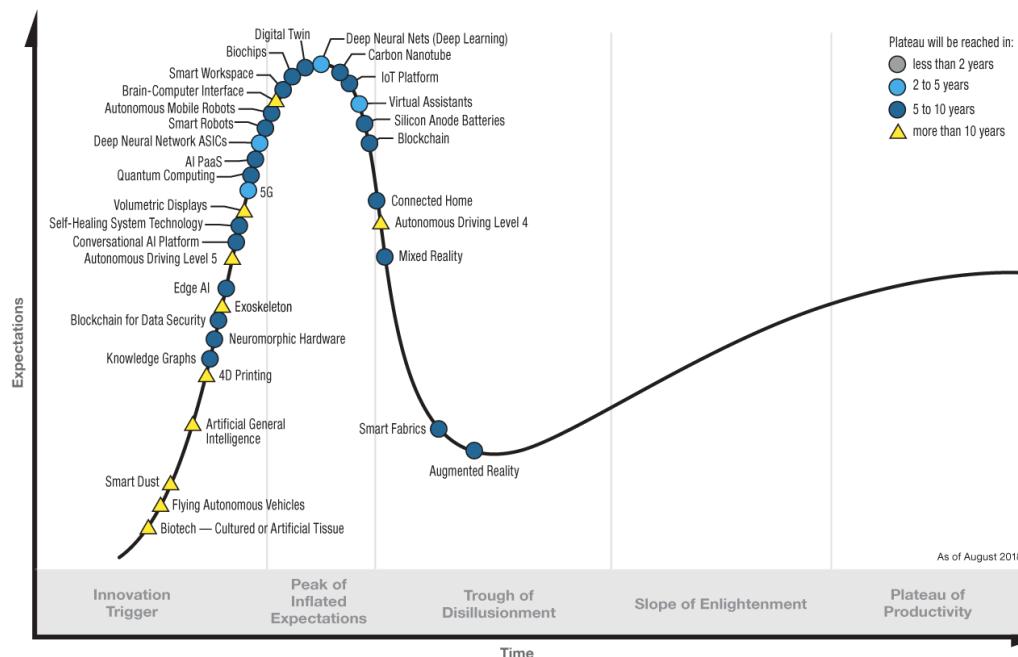


Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- **Fog, Edge and IoT**
 - Concepts
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Fog, Edge and IoT

Hype Cycle for Emerging Technologies, 2018



gartner.com/SmarterWithGartner

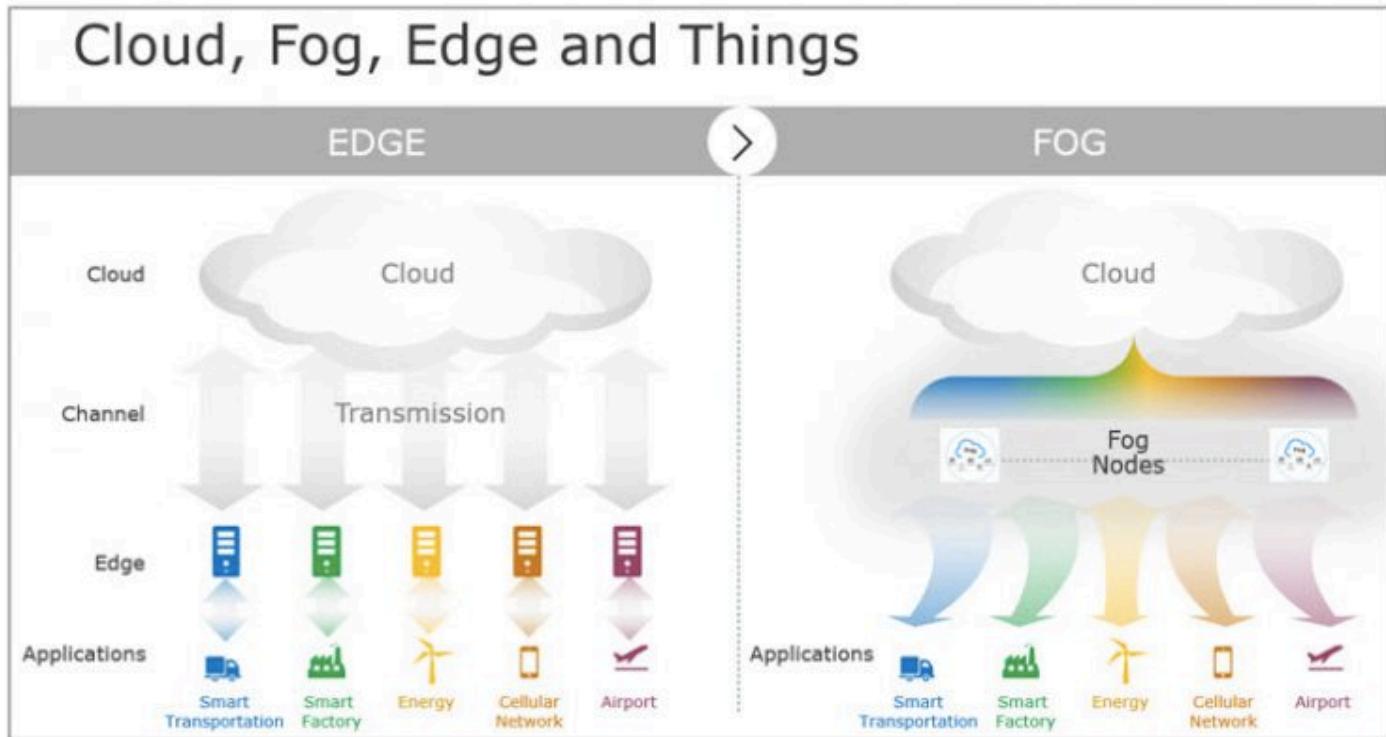
Source: Gartner (August 2018)
© 2018 Gartner, Inc. and/or its affiliates. All rights reserved.

Gartner®

Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - **Concepts**
 - Placement in fog architectures
 - Development practices: Testbed for fog computing

Edge and Fog



<https://www.openfogconsortium.org/>

Cloudlets

- Paper from 2009

The Case for VM-based Cloudlets in Mobile Computing

Mahadev Satyanarayanan[†], Paramvir Bahl[‡], Ramon Caceres[•], Nigel Davies[◦]

[†]Carnegie Mellon University, [‡]Microsoft Research, [•]AT&T Research, [◦]Lancaster University

- Idea: deploy micro data centers (called cloudlets) everywhere for usage in mobile applications
- Geo-distributed like wifi routers
- Term now also describes nodes in a fog architecture (i.e. [1])

Cloudlets

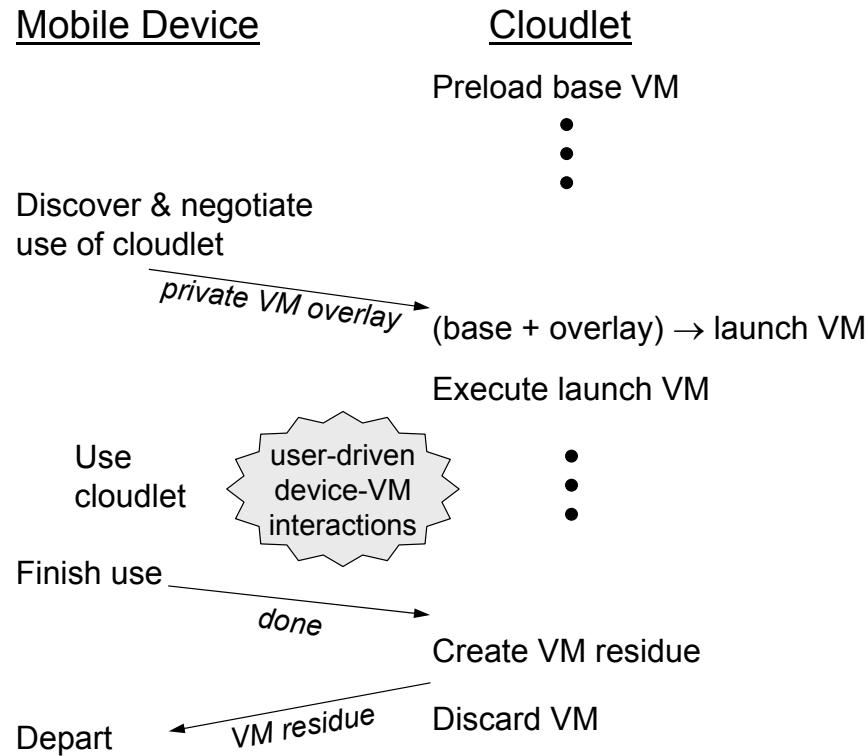
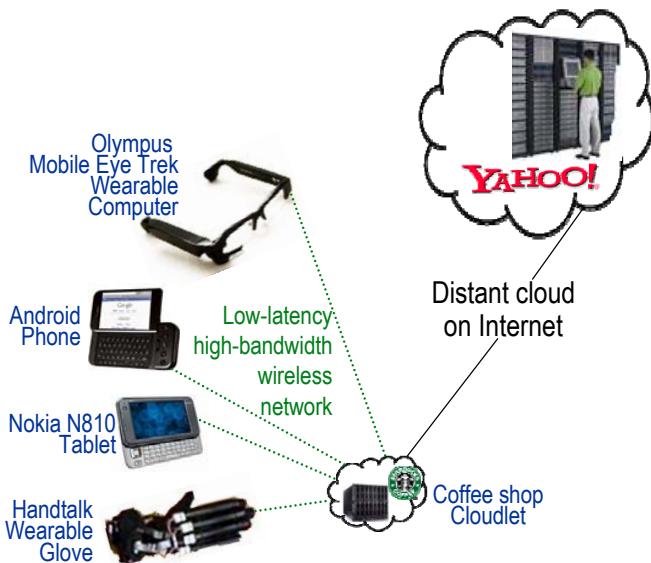


Figure 5: Dynamic VM Synthesis Timeline

Cloudlets and Clouds



| | Cloudlet | Cloud |
|-------------------------|---|--|
| <i>State Management</i> | Only soft state | Hard and soft state |
| <i>Environment</i> | Self-managed; little to no professional attention “Datacenter in a box” at business premises | Professionally administered, 24x7 operator Machine room with power conditioning and cooling |
| <i>Ownership</i> | Decentralized ownership by local business | Centralized ownership by Amazon, Yahoo!, etc. |
| <i>Network Sharing</i> | LAN latency/bandwidth Few users at a time | Internet latency/bandwidth 100s-1000s of users at a time |

(b) Key Differences: Cloudlet vs. Cloud

Cloudlets: Proof of Concept Architecture

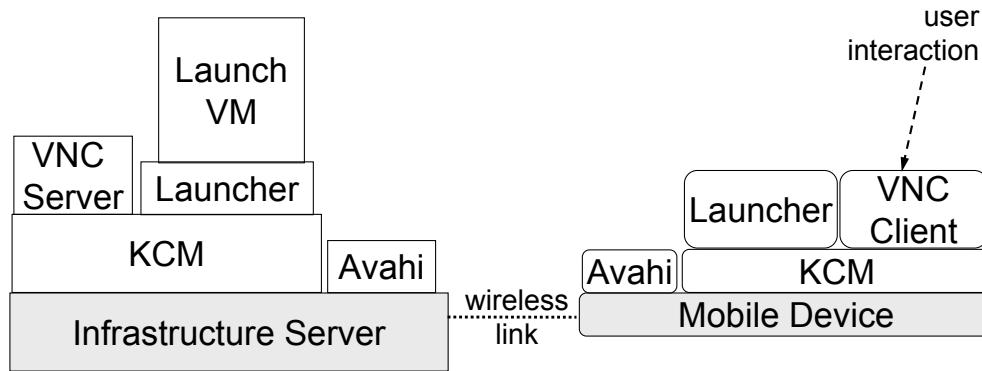


Figure 6: Runtime Binding in Kimberley

Overview

- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - Concepts
 - **Placement in fog architectures**
 - Development practices: Testbed for fog computing

Scheduling in Fog Architectures

- Workshop paper from 2018

Scheduling Stream Processing Tasks on Geo-Distributed Heterogeneous Resources

Gerrit Janßen, Ilya Verbitskiy, Thomas Renner, and Lauritz Thamsen

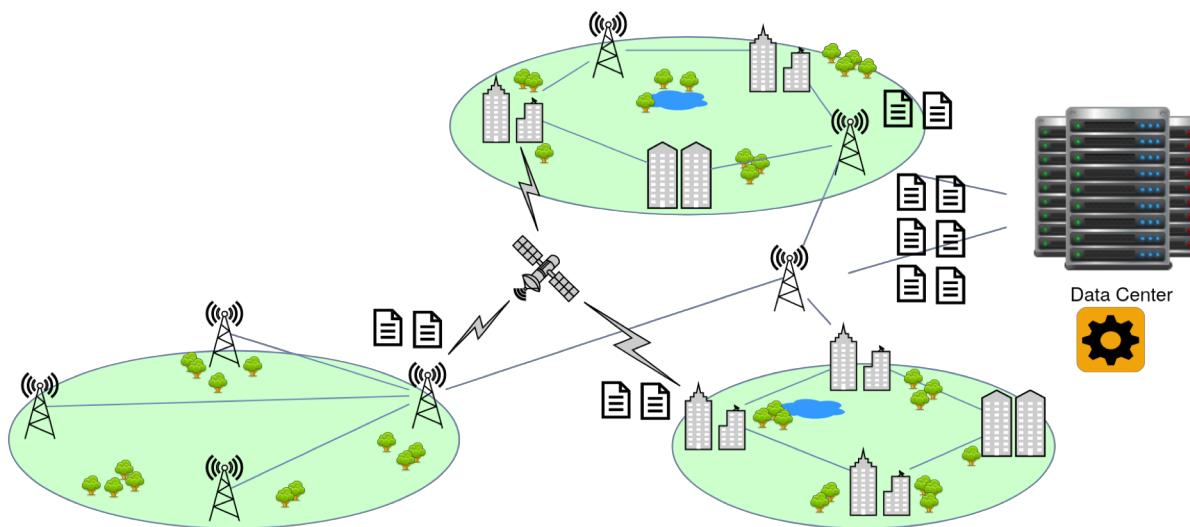
Technische Universität Berlin, Germany

{firstname.lastname}@tu-berlin.de

- Work done within the Berlin Big Data Center (BBDC), in Phase 2 on Distr. Processing of IoT Data Streams
- Empirical evaluation of different methods for scheduling Flink tasks across heterogeneous distributed nodes

Data Streams of the IoT

- Various sensors deployed across factories and urban infrastructures, continuously recording data



Quality of Service Requirements

- Distributed sensor systems increasingly deployed even in critical infrastructures → critical performance requirements, e.g.
 - end-to-end latencies
 - throughput
 - resource usage
 - reliability
 - availability



ENERGY



HEALTH



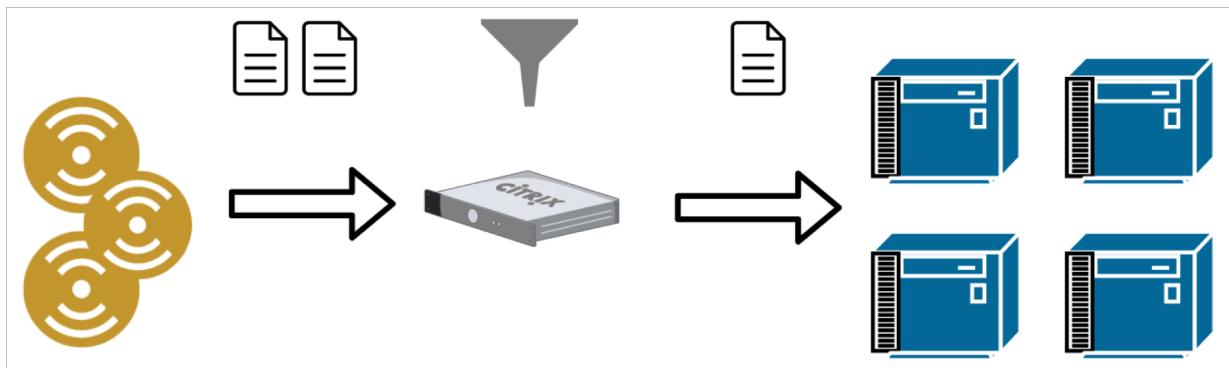
TRANSPORT



WATER

Edge and Fog Computing

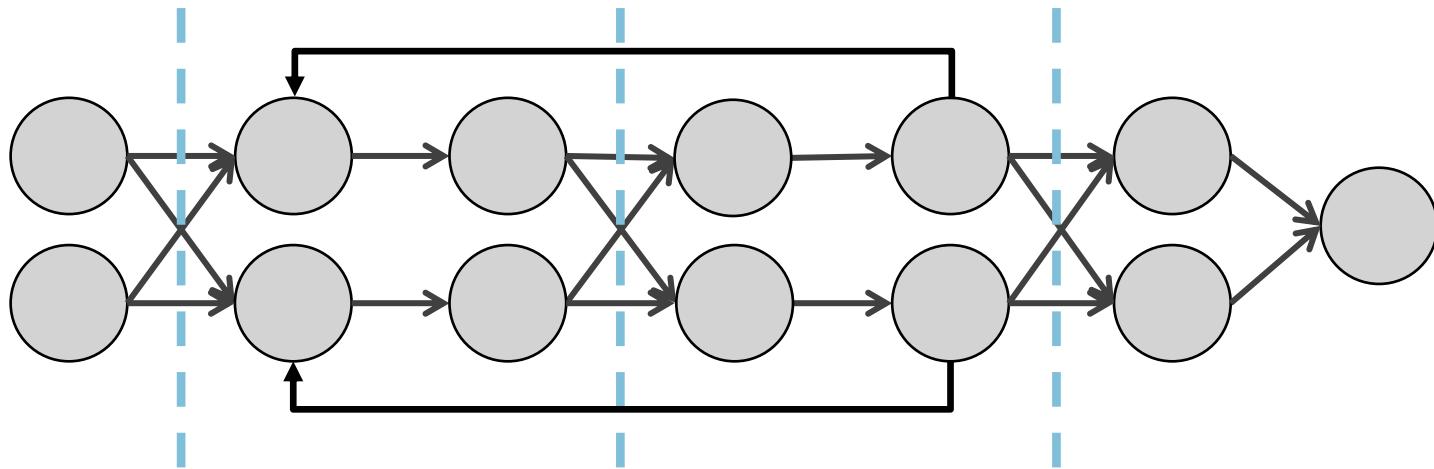
- From Edge to Fog to Cloud → execution of tasks on resources in-between sensors and Clouds



- Processing in increasingly heterogeneous, widely distributed computing environments

Distributed Dataflow Systems

- Popular tools for general-purpose scalable data processing, e.g. MapReduce, Spark, Flink, Beam
- Distributed execution of data-parallel operator graphs on shared-nothing commodity clusters



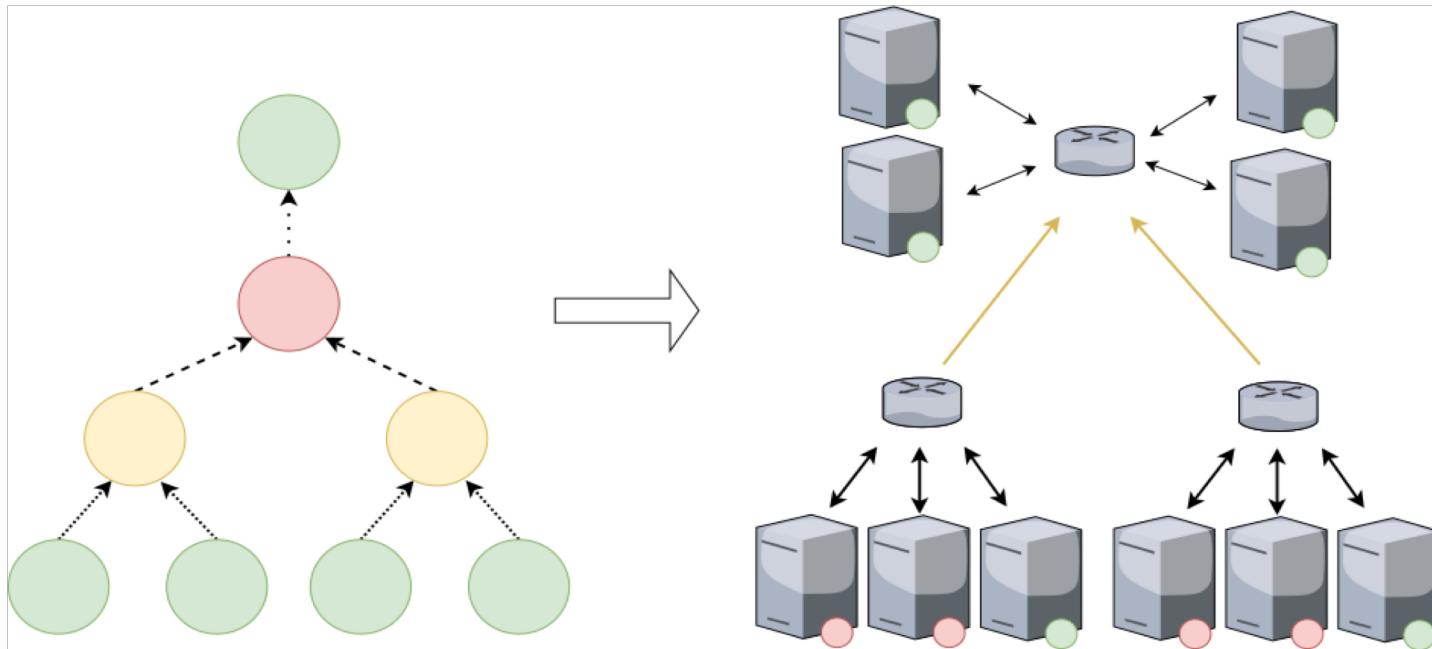
Processing of Continuous Streams

- Dedicated streaming systems (e.g. Storm), Microbatching (e.g. Spark Streaming), and unified dataflow engines (e.g. Flink and Beam) with
 - Low-latency processing of incoming data
 - Flexible windows for aggregation
 - Event time and out-of-order processing
 - Stream processing guarantees

Dataflow Task Scheduling

- At the moment, systems like Spark and Flink mostly expect to run on homogeneous clusters:
 - Same resource configuration per worker
 - Task scheduling without regard to topologies
 - Aggressive task chaining
- Yet, systems need to take the heterogeneity and geo-distribution of resources into account for IoT

Goal: Effective Task Placement



Taskgraphs:

- Resource requirements
- Task output

Infrastructure topologies:

- Latency
- Bandwidth
- Compute power

Comparison of Heuristic Methods

- Metric-based search methods:
 - Steepest descent
 - Tabu search
- Topology-based constructive methods:
 - Highest-level first (Edge first)
 - Top down (Cloud first)

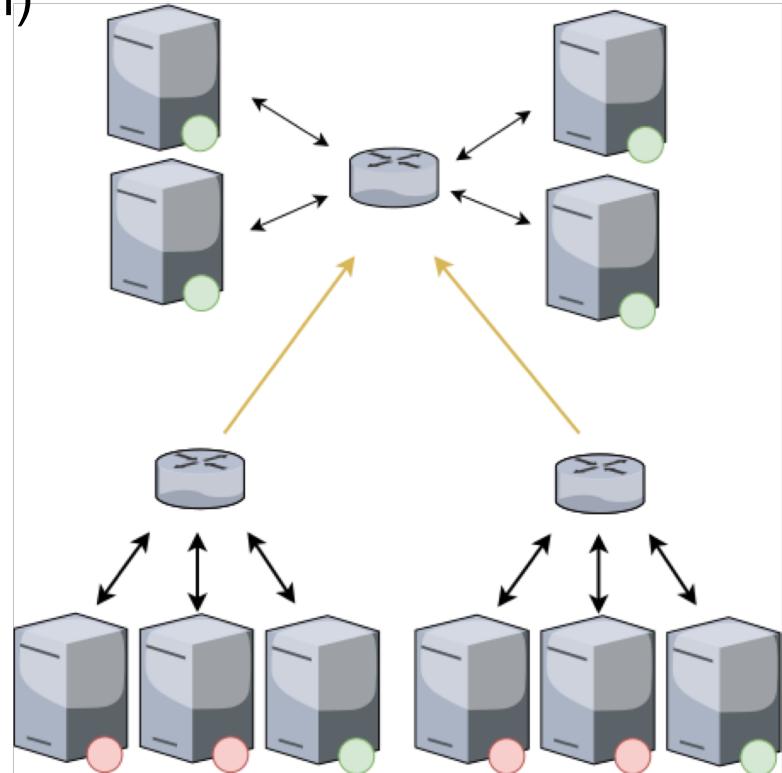
Quality of Service Metrics

- Response Time (RT): time to transmit and process data items from source to sink tasks
- Resource Fitting (RF): assessment of the task placement by a distance function
- Bandwidth Congestion (BW): accumulated underprovisioning of the network
- Overall optimization function for placement P:

$$F(P) = \frac{RT(P)-RT_{\min}}{RT_{\max}-RT_{\min}} + \frac{RF(P)-RF_{\min}}{RF_{\max}-RF_{\min}} + \frac{BW(P)-BW_{\min}}{BW_{\max}-BW_{\min}}$$

Evaluation: Testbed

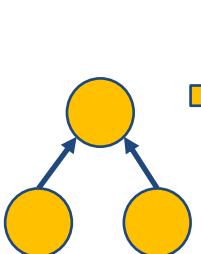
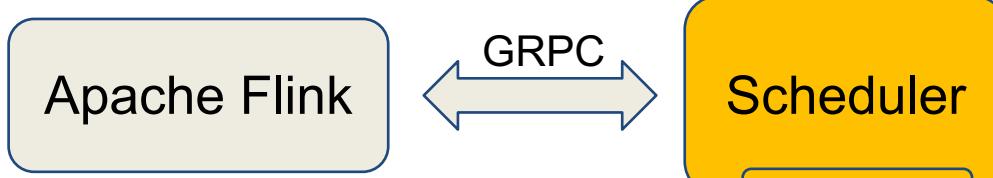
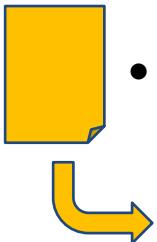
- Network emulation with GNS3 deployed on 9 machines (Intel Xeon E3-1230 V2 3.30GHz, 16 GB RAM, 1 Gbit/s Ethernet)
- Network and bandwidths (Open vSwitch)
 - Yellow (100 MB/s)
 - Black (no limitation)
- Latencies (NetEm)
 - Yellow (80ms)
 - Black (no limitation)
- CPU limitations (Docker)
 - RED - 1 core
 - Green - 4 cores



Evaluation: Prototype

Job metrics

- Resource requirements
(low, medium, high)
- Operator output
rates/size



Topology metrics

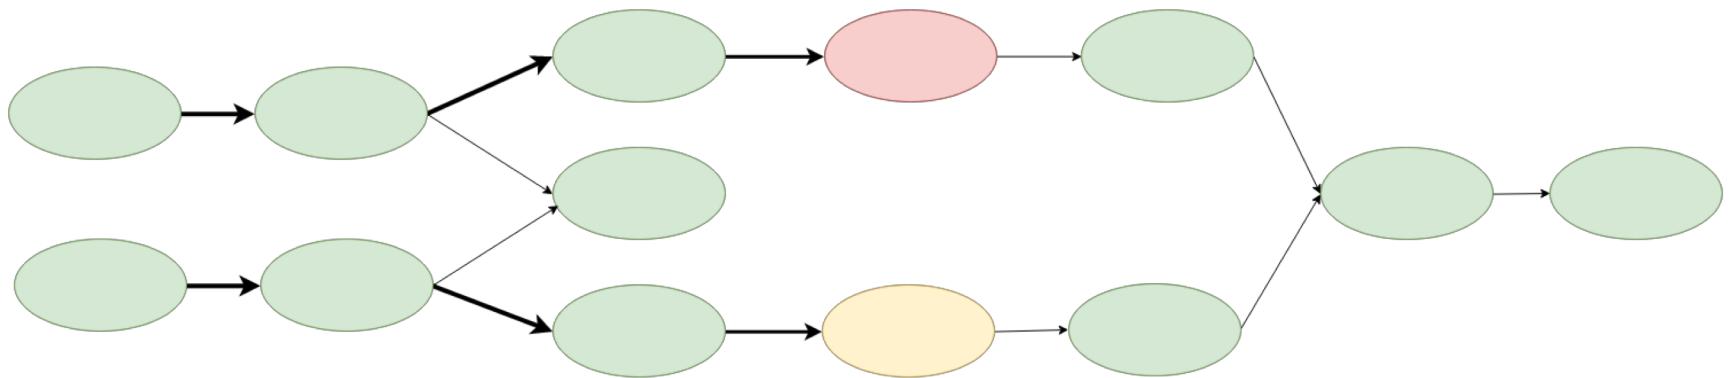
- Neighbor informations
 - Bandwidth
 - Latency
- Resource classification
(low, medium, high)

Scheduler

- converts job/topology informations to graphs
- algorithms for topological- and metrics-based scheduling

Evaluation: Test Job

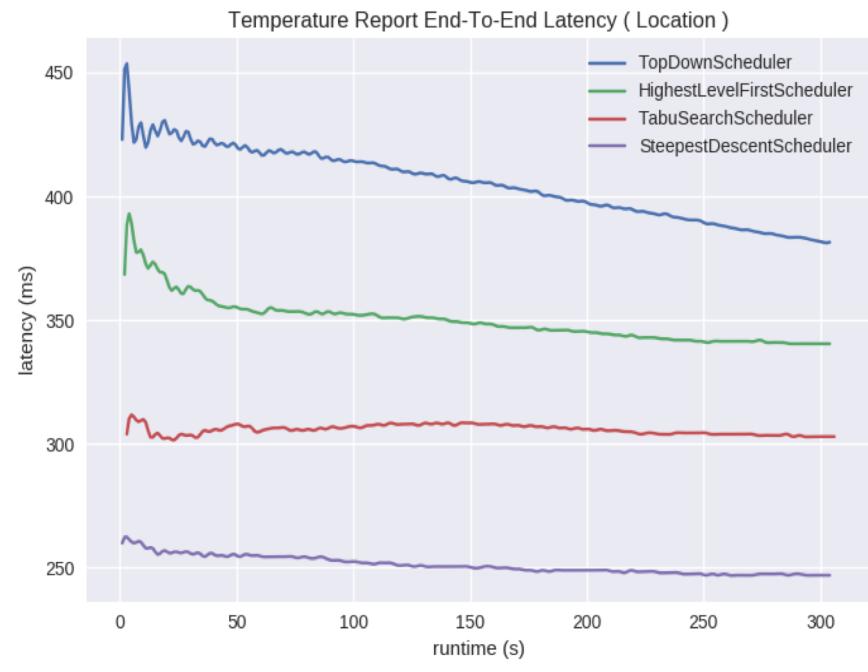
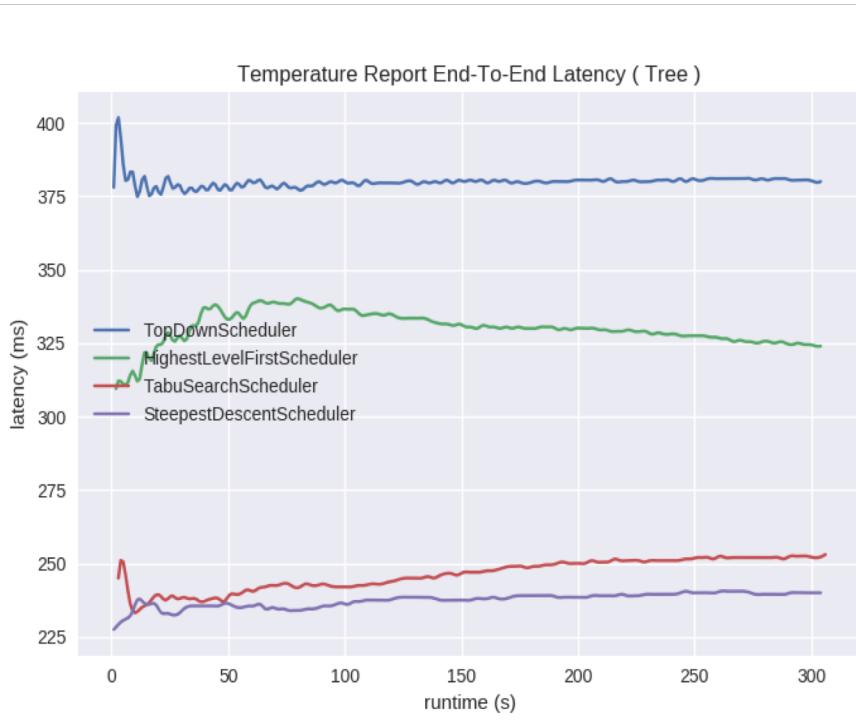
- Monitoring job for continuous temperature monitoring and fire detection



- Resource Classification: red -> high, yellow -> moderate, green -> low

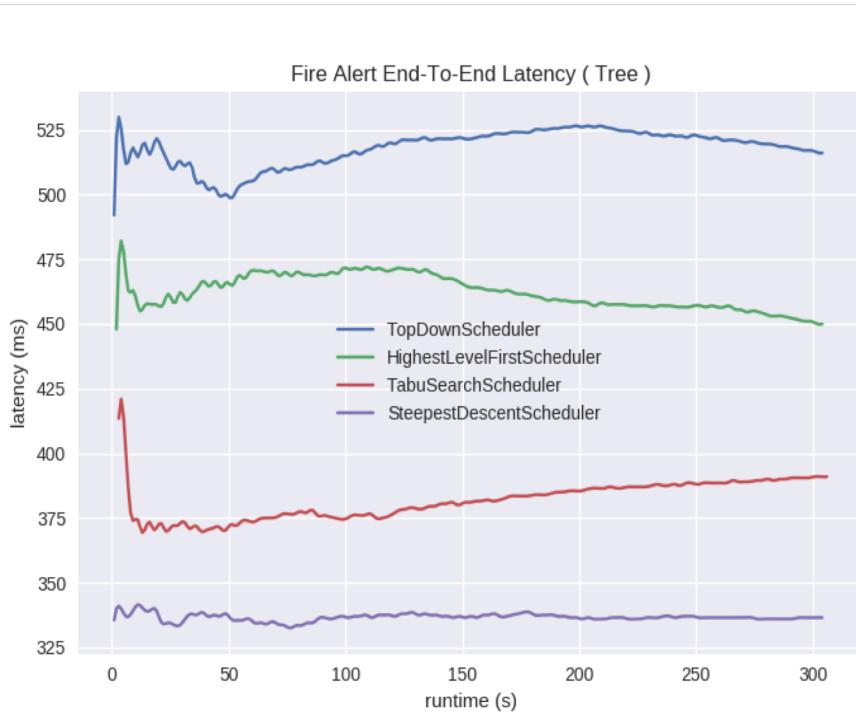
Evaluation: Results

- Ingestion rate: 100 TreeReports and 25 LocationReports per second



Evaluation: Results

- Ingestion rate: 100 TreeReports and 25 LocationReports per second



Conclusion

- Performance of the benchmark application varies drastically depending on the task placement
- Metrics-based methods yield up to 2x better schedules than the worst schedules in benchmark
- Yet:
 - Currently users need to provide additional inputs on their jobs
 - Environment and schedules are assumed static

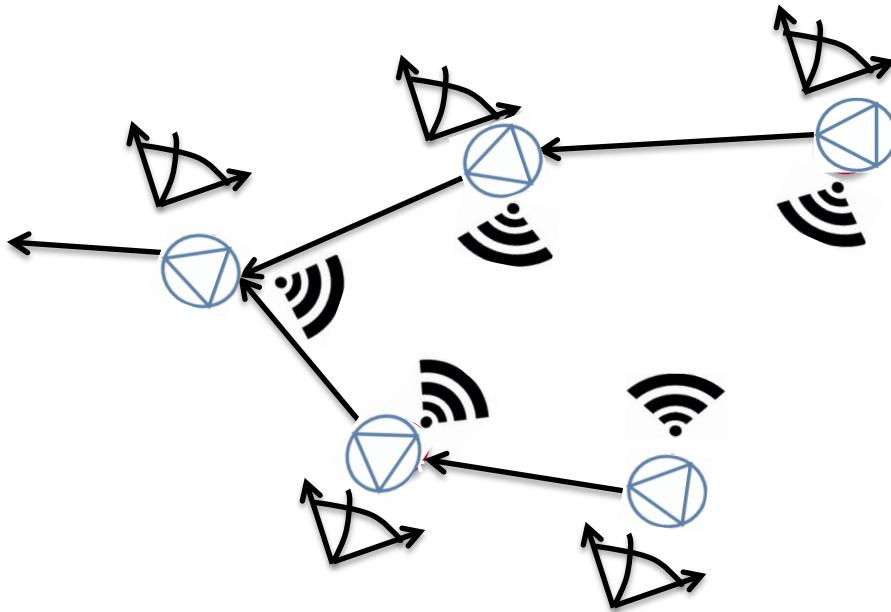
Future Work

- Towards adaptive and fully transparent methods
 - Automatic inference of resource requirements
 - Continuous monitoring of job and topology metrics
 - Live adaptation of schedules (either re-scheduling or migration of tasks)
- More comprehensive empirical evaluation
 - More realistic test environments
 - More benchmarks

Overview

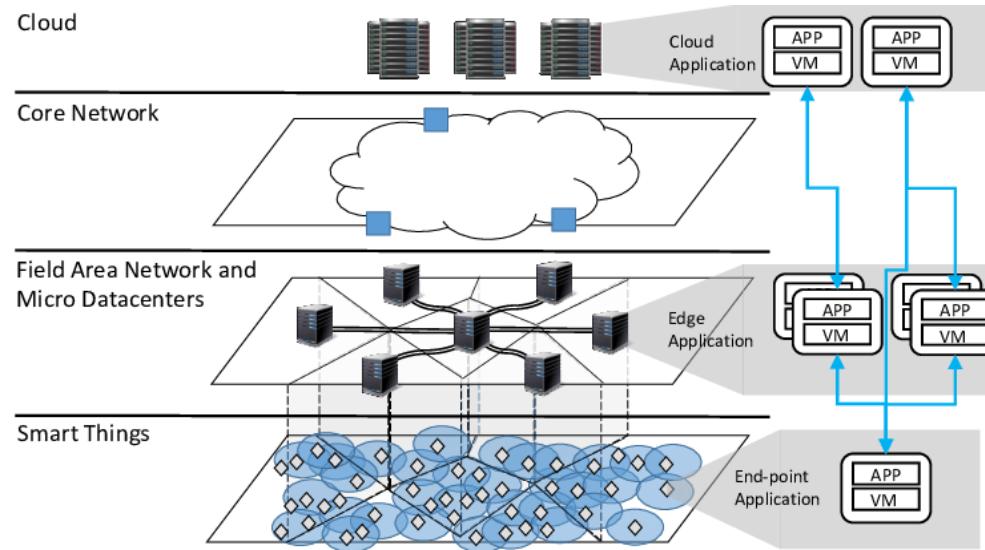
- Intro
- Federated, Hybrid and Multi-Clouds
 - Concepts
 - Policy enforcement
- Fog, Edge and IoT
 - Concepts
 - Placement in fog architectures
 - **Development practices: Testbed for fog computing**

Internet of Things Sensor Systems



- Remote monitoring applications, e.g.
 - rule-based with thresholds, aggregate statistics,
 - clustering, classification, anomaly detection, or
 - input for simulations and system optimization

Edge/Fog Architectures of the IoT



- Networks and nodes in IoT architectures:
 - Widely distributed and heterogeneous
 - Dynamically changing systems and workloads (e.g. events, failures, weather, attacks)

Critical Service Constraints



ENERGY



HEALTH



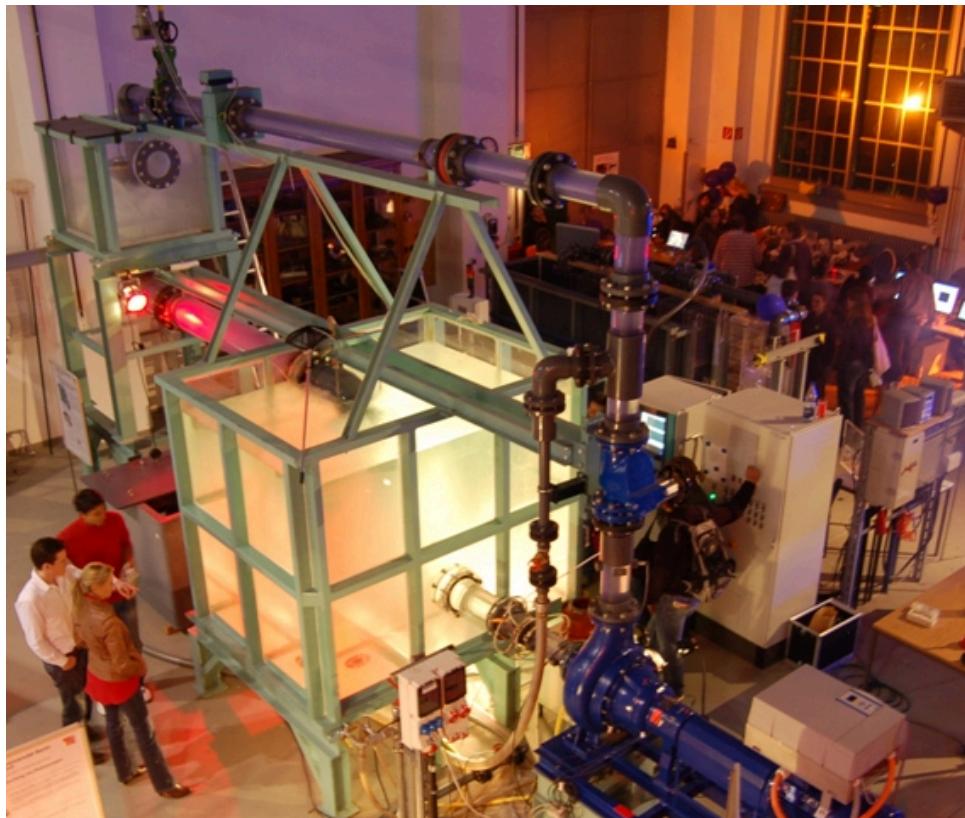
TRANSPORT



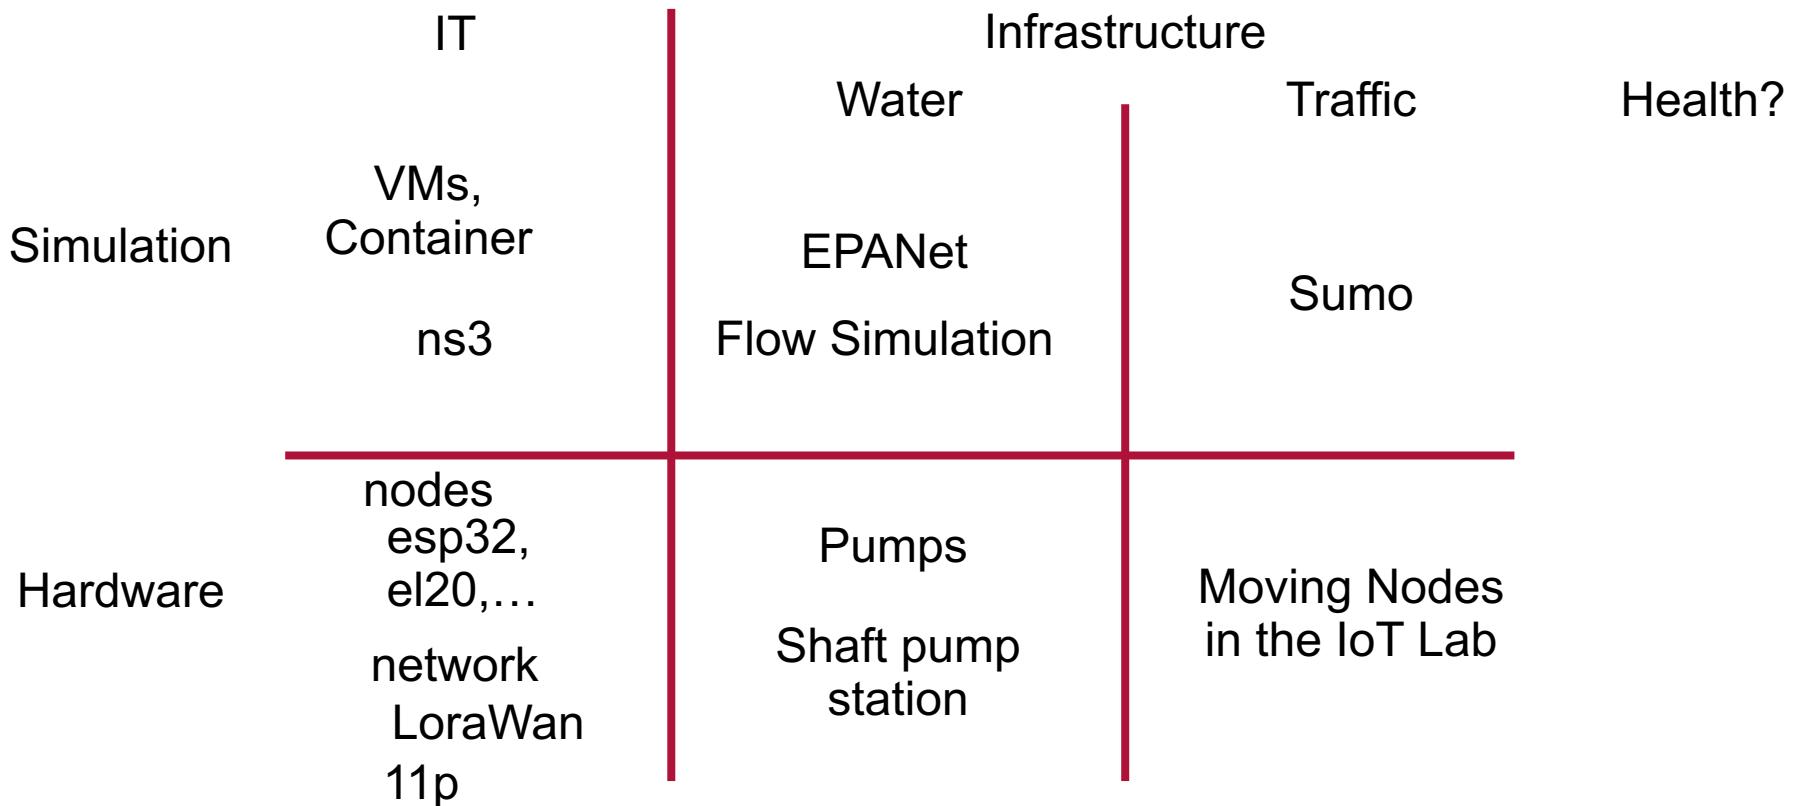
WATER

- Strong requirements in many critical IoT applications, e.g. in industrial and public infrastructures:
 - performance (latencies, resource usage)
 - reliability, availability, and security (exactly-once, max recovery times, uptime under stress/attack)

Hardware Testbeds



A Testbed for the Internet of Things



Integrated IoT Testbed (2/2)

- ns3 / NetEm
- VMs and containers
- qemu

e.g. water:
EPAnet

- esp32
- Raspberry Pis
- wally

- Pumps
- small
models

← Calibration (and continuous re-calibration)
← Virtual sensors in simulations