

Methods of Cloud Computing

Chapter 2: Virtual Resources



Complex and Distributed Systems
Faculty IV
Technische Universität Berlin



Operating Systems and Middleware
Hasso-Plattner-Institut
Universität Potsdam

Overview

- Virtual Resources and Infrastructure-as-a-Service
- Hardware Virtualization
 - Binary Translation, OS-Assisted Virtualization, Hardware-Assisted Virtualization
 - Virtual Machine Migration
 - Resource Isolation and Performance Implication
 - Case Study: Amazon EC2
- **OS-Level Virtualization**
 - Linux Containerization
 - LXC Containers and Docker
 - Comparison to Virtual Machines





OS Virtualization

- Also known as Container Virtualization
- Motivation: HW virtualization comes with too much overhead, large images, and long boot times
- Idea: do not virtualize entire machine, but...
 - Reuse operating system kernel and isolate applications
 - Virtualize access to resources used by processes
 - ◆ File system
 - ◆ Devices
 - ◆ Network
 - ◆ Other processes
 - ◆ ...

History of OS Virtualization

- Unix v7 **chroot** system call (1979)
 - Allows setting the file system root for processes
 - Unix philosophy: almost everything is accessed through the file system → almost everything can be virtualized though **chroot**
- Wave of container technologies in early 2000s
 - FreeBSD Jails, Linux VServer, Solaris Zones, ...
 - Different degrees of isolation, different tool chains
 - Mostly specific to Linux distribution
 - Never popular with big masses, mainly used by system admins and large companies

History of OS Virtualization

- LXC (Linux Containers, 2008) LXC
 - Easy-to-use user space tools for accessing Linux kernel process isolation features (cgroups, namespaces)
- Docker (2013) docker
 - Became the next big thing after the initial cloud hype
 - Most wide spread container technology and eco system
- Rocket (CoreOS, 2014) Rocket
 - Alternative to Docker
 - More focused on security and standardization
- LXD (Canonical) LXD
 - Focused on entire Linux distributions (not applications)

Overview


- Virtual Resources and Infrastructure-as-a-Service
- Hardware Virtualization
 - Binary Translation, OS-Assisted Virtualization, Hardware-Assisted Virtualization
 - Virtual Machine Migration
 - Resource Isolation and Performance Implication
 - Case Study: Amazon EC2
- OS-Level Virtualization
 - **Linux Containerization**
 - LXC Containers and Docker
 - Comparison to Virtual Machines

Linux Kernel Features for OS Virtualization

- Linux kernel contains large number of mechanisms for process isolation:
 - chroot system call
 - Namespaces
 - Capabilities
 - cgroups
 - SELinux
 - seccomp
 - ...
- Partially overlapping functionality, but different configuration approaches

chroot

- Chroot
 - Oldest mechanism for process isolation
 - System call, changes the root directory (“/”) of the calling process
 - Since “everything is a file” in Linux, many aspects of the underlying system can be virtualized
 - System calls, networking, etc. remain unchanged




Linux Namespaces













- Namespaces
 - Separate views on kernel resources for processes in different namespaces
 - Isolated resources:
 - ◆ IPC (Inter-process communication)
 - ◆ Network (devices, protocol stacks, firewalls, ports, ...)
 - ◆ Mount (mount points, file system structure)
 - ◆ PID (processes)
 - ◆ User (users and groups)
 - ◆ UTS (hostname and domain name)

Linux Namespaces

- Container 1: Share users and hostname, but isolate processes, network and mount points

 Namespaces used by processes running in Container 1

 Namespaces used by other processes


IPC	 Default	 Container 1
Network	 Default	 Container 1
Mount	 Default	 Container 1
PID	 Default	 Container 1
User	 Default 	
UTS	 Default 	

Linux Capabilities

- Capabilities
 - Traditionally, super user (root) is the only one to perform administrative tasks on a Linux system
 - New processes can be assigned selected capabilities
 - Long list of possible capabilities, including many system calls, device access, file system modifications, ...

Linux cgroups

- cgroups (control groups)
 - Allows definition of resource usage constraints for parts of the process tree
 - Used for limiting CPU, RAM, network, and disk usage for containers



Security Policies

- SELinux, AppArmor
 - Optional security kernel modules
 - Allow rule-based confinement of processes to limit their access to certain files and devices only
 - SELinux: complex and powerful rule definitions
 - AppArmor: more straight-forward security profiles
- Seccomp policies
 - seccomp system call puts the process in a secure computation mode, where only selected system calls are allowed

Overview

- Virtual Resources and Infrastructure-as-a-Service
- Hardware Virtualization
 - Binary Translation, OS-Assisted Virtualization, Hardware-Assisted Virtualization
 - Virtual Machine Migration
 - Resource Isolation and Performance Implication
 - Case Study: Amazon EC2
- OS-Level Virtualization
 - Linux Containerization
 - **LXC Containers and Docker**
 - Comparison to Virtual Machines

Examples for Container Technologies



LXC

- LXC: First Container technology widely adopted by end users



- Docker: Currently dominant container platform and eco system

Linux Containers

- LXC: Library and userspace tools (bash scripts) to access kernel process isolation features
- No daemon, containers and their file systems are stored in `/var/lib/lxc`

Language Bindings

(Python, Lua, Go, Python, Ruby, ...)

Userspace tools

`lxc-create`, `lxc-start`, `lxc-stop`, ...

liblxc

Kernel features:

Namespaces, capabilities, `chroot`, `cgroups`,
`AppArmor`, `SELinux`, `seccomp`

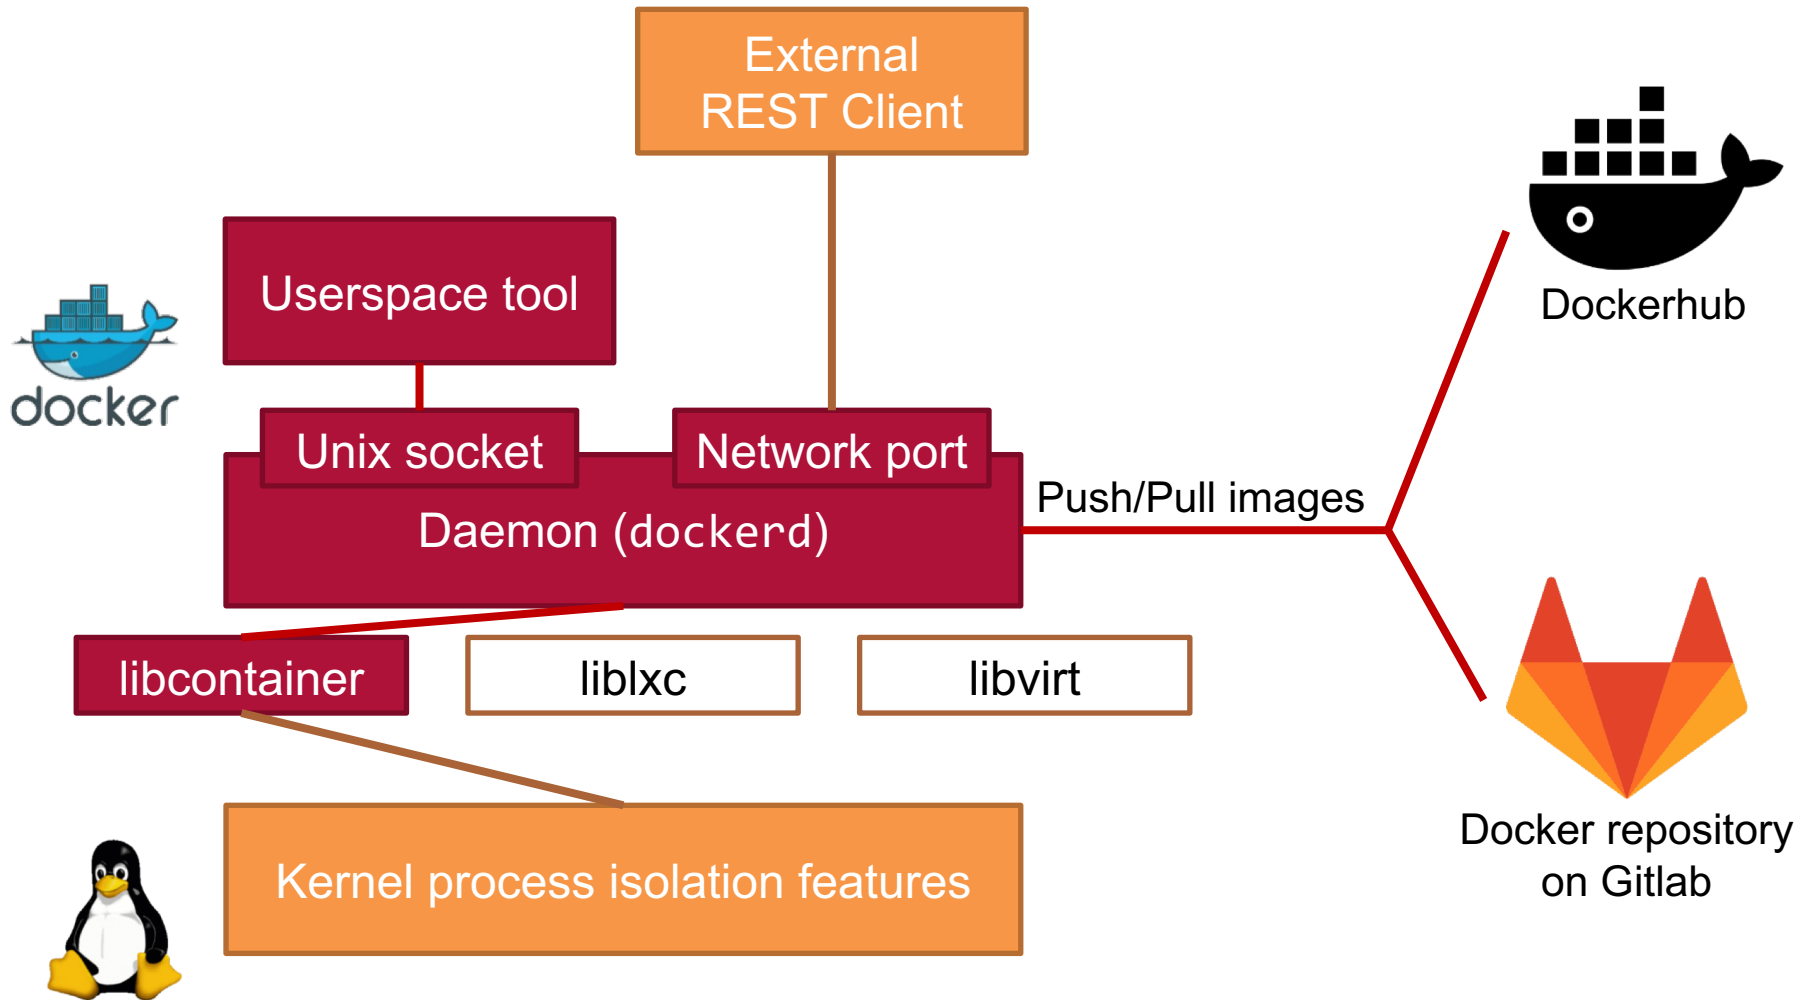
LXC Container Creation

- Container Creation
 1. Container directory allocated under `/var/lib/lxc`
 2. “Template script” executed to populate the file system of the container
 3. Process is spawned and chrooted into its file system
 4. Isolation properties are configured
 5. Main application process is spawned, all child processes will inherit isolation context
- LXC does not rely on an image format
- No public repository for sharing images (introduced later in LXD)

Docker

- Most popular container technology at the moment
- Docker includes:
 - Daemon and user space tools for managing local containers and images
 - Hierarchical image format
 - Public and private repositories for sharing images
 - Included and external capabilities for automatic orchestration of container infrastructures

Docker Components

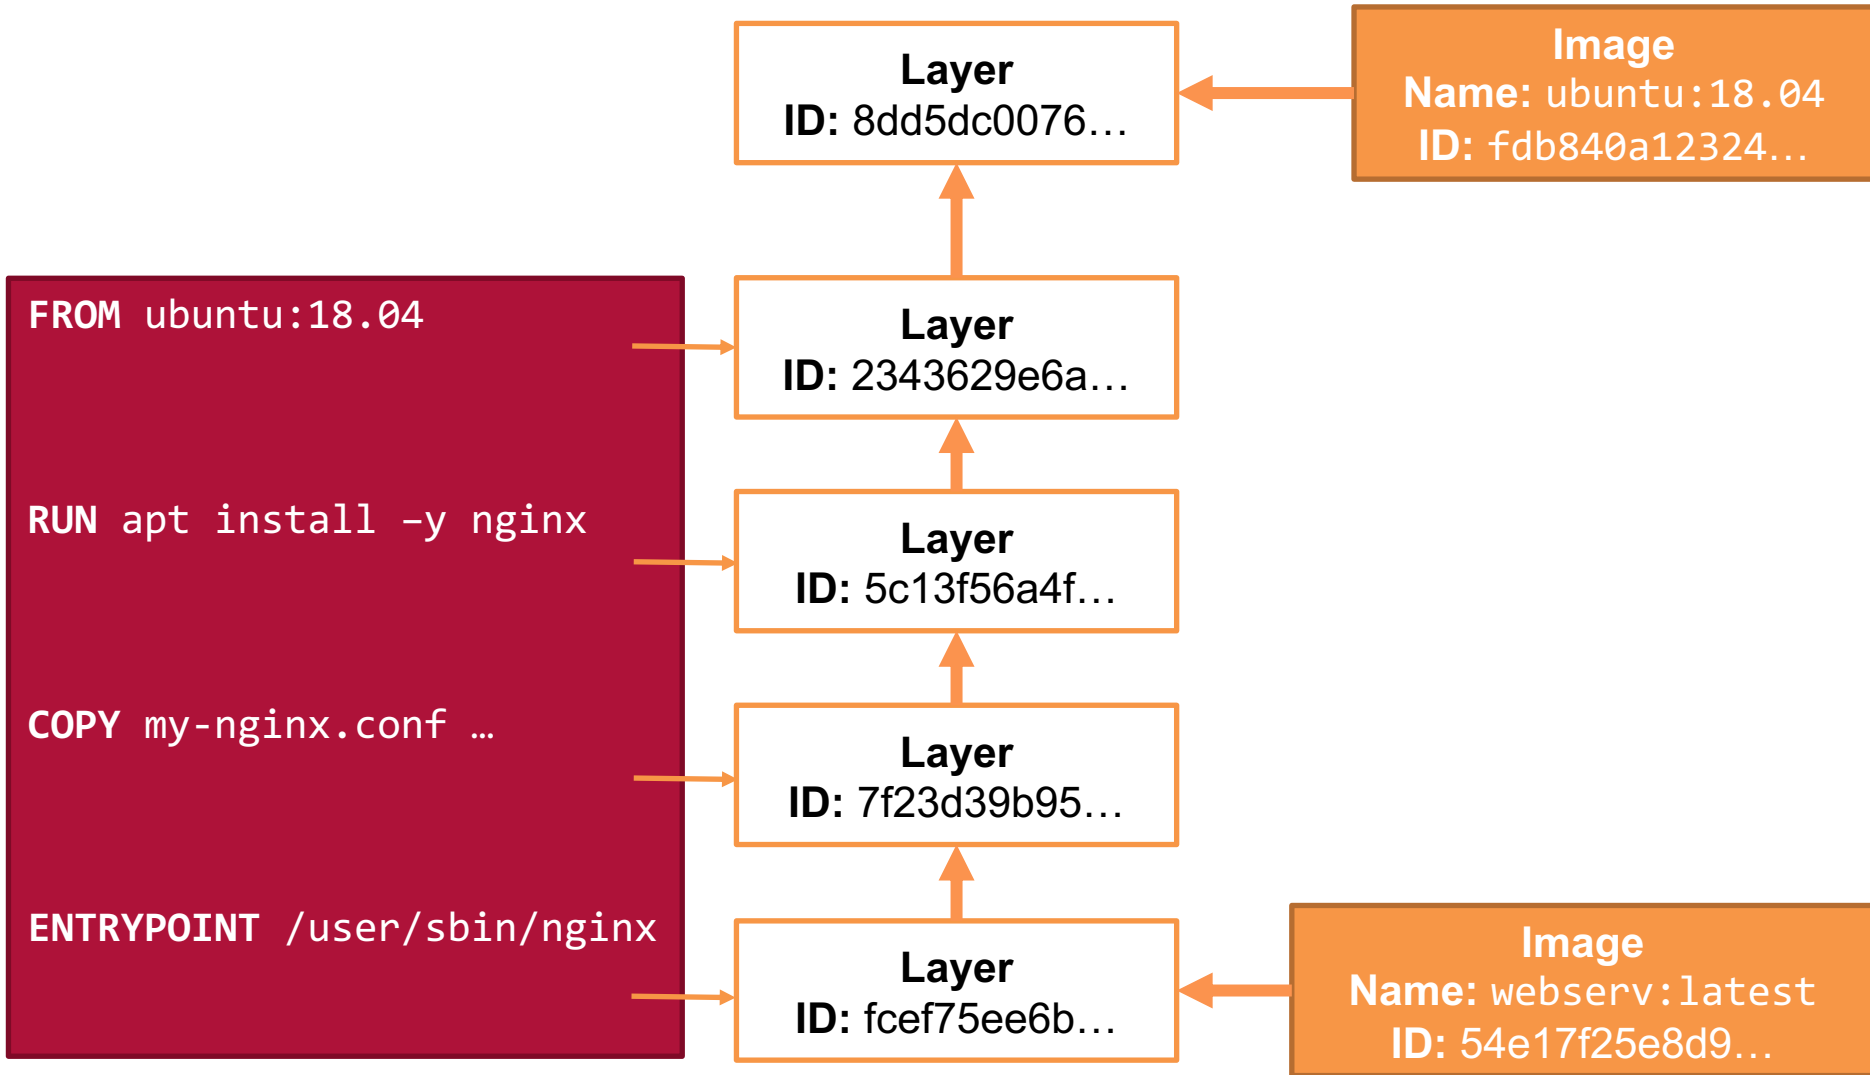


Dockerfiles

- Automatic way to create container images
- Text file with commands that modify files or change configuration values
- Intermediate steps can be cached using hashes, resulting in reduced image build times

```
FROM ubuntu:18.04
RUN apt install -y nginx
COPY my-nginx.conf /etc/nginx/nginx.conf
ENTRYPOINT /user/sbin/nginx
```

Docker Image Format



Docker Image Format

- Docker image contains:
 - List of layers that form the file system
 - ID (Hash of layer IDs and other config data)
 - Configuration data: ports, mounts, variables, ...
 - Other meta info (creation time, author, history, ...)
- Docker image **layer** contains:
 - Layer files (tarball): Files that were added/changed on this layer, relative to parent layer
 - ID of the layer (Hash of all files), ID of parent layer
 - Command that was used to create the layer

Overview

- Virtual Resources and Infrastructure-as-a-Service
- Hardware Virtualization
 - Binary Translation, OS-Assisted Virtualization, Hardware-Assisted Virtualization
 - Virtual Machine Migration
 - Resource Isolation and Performance Implication
 - Case Study: Amazon EC2
- OS-Level Virtualization
 - Linux Containerization
 - LXC Containers and Docker
 - **Comparison to Virtual Machines**



Containers vs VMs: Performance (CPU & RAM)

Workload	Docker	KVM
Linpack (GFLOPS)	290.9 [± 0.98]	284.2 [± 1.45]
Memory (Random Access, GLOps/s)	0.0124 [± 0.00044]	0.0125 [± 0.00032]
Memory (Sequential Access, GB/s)	45.6 [± 0.55]	45.0 [± 0.19]

- VMs introduce very low CPU and memory access overhead
 - Condition: exposing cache topology and CPU acceleration features (e.g. NUMA, FPU, SSE)



Source: IBM Research Report (An Updated Performance Comparison of Virtual Machines and Linux Containers, 2014)
[https://domino.research.ibm.com/library/cyberdig.nsf/papers/0929052195dd819c85257d2300681e7b/\\$file/rc25482.pdf](https://domino.research.ibm.com/library/cyberdig.nsf/papers/0929052195dd819c85257d2300681e7b/$file/rc25482.pdf)

Containers vs VMs: Performance (Network)




- Low CPU overhead per packet
- Reasons for latency increase:
 - NAT in Docker networking
 - Virtual network device in KVM (NAT might further increase latency)

Containers vs VMs: Performance (Disk IO)



- Similar throughput
- Penalty for latency and IOPs due to virtual IO device



Containers vs VMs: Image Size & Boot Time

- Virtual Machine images usually larger than container images (contain entire OS), but often image caching on execution host
- Boot time of VM can be orders of magnitude longer than container startup

Containers vs VMs: Isolation & Security

- Containers share the host OS kernel
 - Vulnerable to Linux kernel bugs
 - Denial-of-Service attacks: Resource usage, system calls, context switches of a container can starve others
 - Kernel parameters cannot be tuned to workload
- Virtualization technology older and VMM order of magnitude less code → more exploits discovered and fixed

Containers vs VMs: Summary

Virtual Machines

- Complete isolation
- Flexible OS
- Live migration

Containers

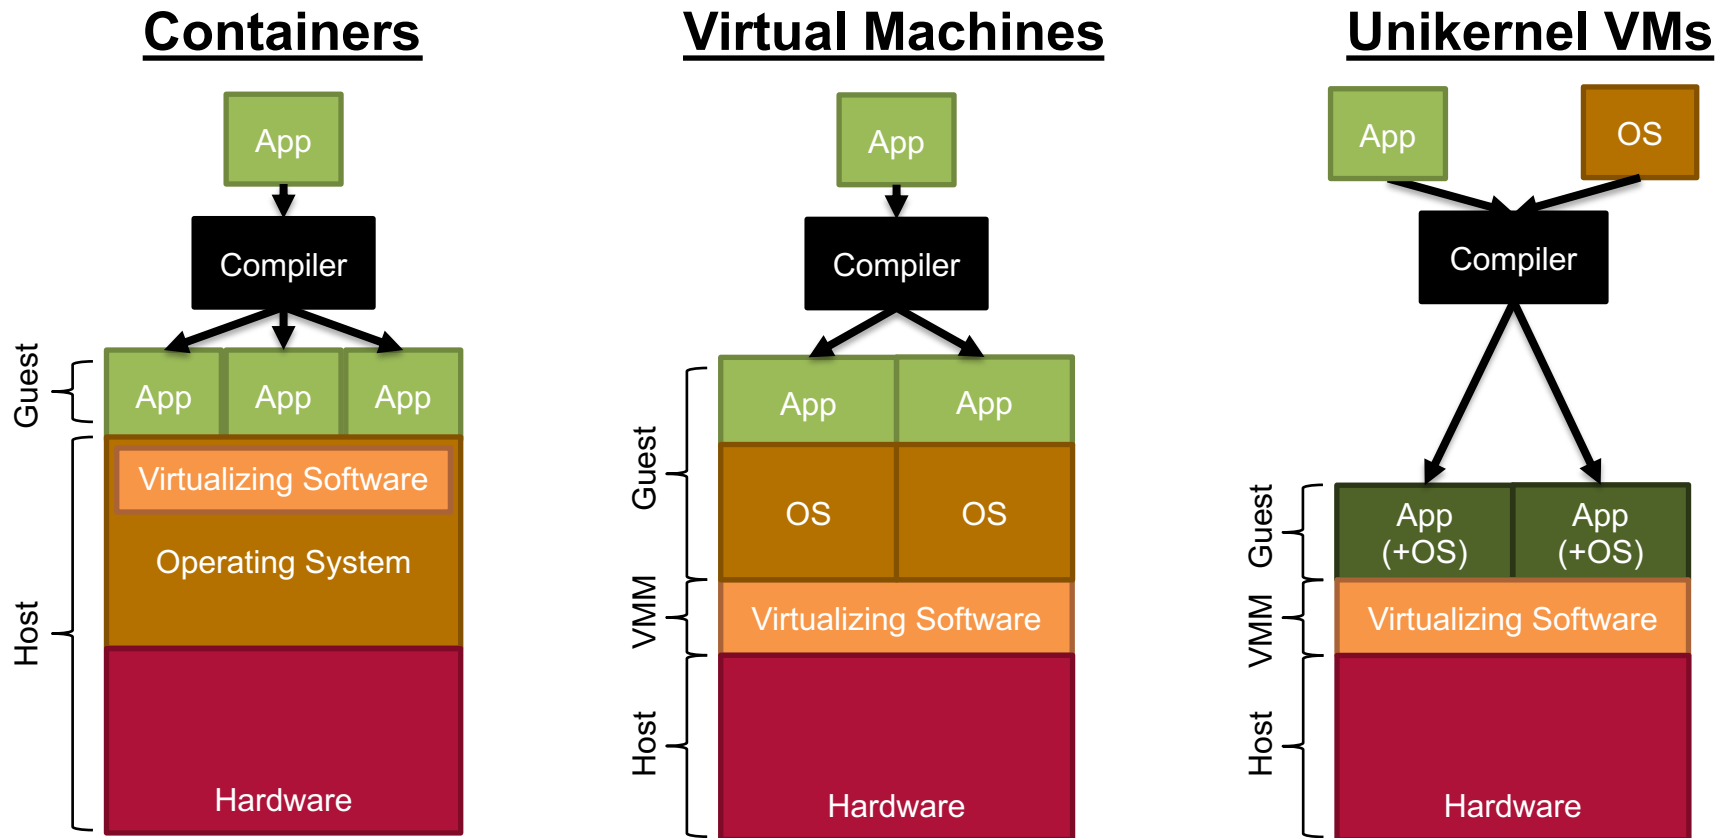
- Small images
- Quick startup
- Direct device access

→ Tradeoff between performance and isolation & security

Sidenote:

Unikernel VM Images

- Compile and link application code directly with all required OS functionality



Overview

- Virtual Resources and Infrastructure-as-a-Service
- Hardware Virtualization
 - Binary Translation, OS-Assisted Virtualization, Hardware-Assisted Virtualization
 - Virtual Machine Migration
 - Resource Isolation and Performance Implication
 - Case Study: Amazon EC2
- OS-Level Virtualization
 - Linux Containerization
 - LXC Containers and Docker
 - Comparison to Virtual Machines

Summary: Virtual Resources

- IaaS clouds let customers rent basic IT resources
 - Full control over OS and deployed applications in VMs
 - No long-term obligation or risk of over-/under-provisioning
- Virtualization as fundamental enabling technology
 - Several customers can share physical infrastructure
 - Different approaches to achieve virtualization
- Containers increasingly recognized as alternative or supplement to VMs

References

- Andrew S. Tanenbaum, Herbert Bos: Modern Operating Systems, Pearson, 2015
 - William Stallings: Operating Systems – Internals and Design Principles, 2015
- [3] G.J. Popek and R.P. Goldberg: “Formal Requirements for Virtualizable Third Generation Architectures”, Communications of the ACM, 17 (7), 1974
- [4] J.S. Robin, C.E. Irvine: “Analysis of the Intel Pentium’s Ability to Support a Secure Virtual Machine Monitor”, Proc. of the 9th Conference on USENIX Security Symposium, 2000
- [6] K. Adams, O. Agesen: “A Comparison of Software and Hardware Techniques for x86 Virtualization”, Proc. of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems, 2006
- [7] A. Whitaker , M. Shaw , S.D. Gribble: “Denali: Lightweight Virtual Machines for Distributed and Networked Applications”, Proc. of the 2002 USENIX Annual Technical Conference, 2002
- [8] P. Barham, B. Dragovic, K. Fraser, S. Hand, T. Harris, A. Ho, R. Neugebauer, I. Pratt, A. Warfield: “Xen and the Art of Virtualization”, Proc. of the 19th ACM Symposium on Operating Systems principles, 2003
- [11] O. Agesen: “Performance Aspects of x86 Virtualization”, VMWORLD 2007
- [13] G. Neiger, A. Santoni, F. Leung, D. Rodgers, R. Uhlig: “Intel Virtualization Technology: Hardware Support for Efficient Processor Virtualization”, Intel Technology Journal, 10 (3), 2006
- [15] C. Clark, K. Fraser, S. Hand, J.G. Hanseny, E. July, C. Limpach, I. Pratt, A. Wareld: “Live Migration of Virtual Machines”, Proc. of the 2nd conference on Symposium on Networked Systems Design & Implementation, 2005
- [18] J. Schad, J. Dittrich, J.-A. Quiané-Ruiz: “Runtime Measurements in the Cloud: Observing, Analyzing, and Reducing Variance”, Proc. of the VLDB Endowment, 3 (1-2), 2010